



Teradata Database

Introduction to Teradata

Release 13.10
B035-1091-109A
August 2010



The product or products described in this book are licensed products of Teradata Corporation or its affiliates.

Teradata, BYNET, DBC/1012, DecisionCast, DecisionFlow, DecisionPoint, Eye logo design, InfoWise, Meta Warehouse, MyCommerce, SeeChain, SeeCommerce, SeeRisk, Teradata Decision Experts, Teradata Source Experts, WebAnalyst, and You've Never Seen Your Business Like This Before are trademarks or registered trademarks of Teradata Corporation or its affiliates.

Adaptec and SCSISelect are trademarks or registered trademarks of Adaptec, Inc.

AMD Opteron and Opteron are trademarks of Advanced Micro Devices, Inc.

BakBone and NetVault are trademarks or registered trademarks of BakBone Software, Inc.

EMC, PowerPath, SRDF, and Symmetrix are registered trademarks of EMC Corporation.

GoldenGate is a trademark of GoldenGate Software, Inc.

Hewlett-Packard and HP are registered trademarks of Hewlett-Packard Company.

Intel, Pentium, and XEON are registered trademarks of Intel Corporation.

IBM, CICS, RACF, Tivoli, and z/OS are registered trademarks of International Business Machines Corporation.

Linux is a registered trademark of Linus Torvalds.

LSI and Engenio are registered trademarks of LSI Corporation.

Microsoft, Active Directory, Windows, Windows NT, and Windows Server are registered trademarks of Microsoft Corporation in the United States and other countries.

Novell and SUSE are registered trademarks of Novell, Inc., in the United States and other countries.

QLogic and SANbox are trademarks or registered trademarks of QLogic Corporation.

SAS and SAS/C are trademarks or registered trademarks of SAS Institute Inc.

SPARC is a registered trademark of SPARC International, Inc.

Sun Microsystems, Solaris, Sun, and Sun Java are trademarks or registered trademarks of Sun Microsystems, Inc., in the United States and other countries.

Symantec, NetBackup, and VERITAS are trademarks or registered trademarks of Symantec Corporation or its affiliates in the United States and other countries.

Unicode is a collective membership mark and a service mark of Unicode, Inc.

UNIX is a registered trademark of The Open Group in the United States and other countries.

Other product and company names mentioned herein may be the trademarks of their respective owners.

THE INFORMATION CONTAINED IN THIS DOCUMENT IS PROVIDED ON AN “AS-IS” BASIS, WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESS OR IMPLIED, INCLUDING THE IMPLIED WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, OR NON-INFRINGEMENT. SOME JURISDICTIONS DO NOT ALLOW THE EXCLUSION OF IMPLIED WARRANTIES, SO THE ABOVE EXCLUSION MAY NOT APPLY TO YOU. IN NO EVENT WILL TERADATA CORPORATION BE LIABLE FOR ANY INDIRECT, DIRECT, SPECIAL, INCIDENTAL, OR CONSEQUENTIAL DAMAGES, INCLUDING LOST PROFITS OR LOST SAVINGS, EVEN IF EXPRESSLY ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

The information contained in this document may contain references or cross-references to features, functions, products, or services that are not announced or available in your country. Such references do not imply that Teradata Corporation intends to announce such features, functions, products, or services in your country. Please consult your local Teradata Corporation representative for those features, functions, products, or services available in your country.

Information contained in this document may contain technical inaccuracies or typographical errors. Information may be changed or updated without notice. Teradata Corporation may also make improvements or changes in the products or services described in this information at any time without notice.

To maintain the quality of our products and services, we would like your comments on the accuracy, clarity, organization, and value of this document. Please e-mail: teradata-books@lists.teradata.com

Any comments or materials (collectively referred to as “Feedback”) sent to Teradata Corporation will be deemed non-confidential. Teradata Corporation will have no obligation of any kind with respect to Feedback and will be free to use, reproduce, disclose, exhibit, display, transform, create derivative works of, and distribute the Feedback and derivative works thereof without limitation on a royalty-free basis. Further, Teradata Corporation will be free to use any ideas, concepts, know-how, or techniques contained in such Feedback for any purpose whatsoever, including developing, manufacturing, or marketing products or services incorporating Feedback.

Copyright © 2000 – 2010 by Teradata Corporation. All Rights Reserved.

Purpose

This book provides an introduction to Teradata covering the following broad topics:

- The data warehouse and active Teradata
- The relational model and Teradata Database architecture
- Teradata Database hardware and software architecture
- Teradata Database RASUI (reliability, availability, serviceability, usability, and installability)
- Communication between the client and Teradata Database
- Data definitions and data manipulation using SQL
- SQL application development
- Data distribution and data access methods
- Concurrent control and transaction recovery
- The Data Dictionary
- International character support
- Query and database analysis tools
- Database security and system administration
- Managing and monitoring Teradata Database

Audience

This book is intended for users interested in a broad overview of Teradata. Such individuals may include database users or administrators.

Supported Software Releases and Operating Systems

This book supports Teradata® Database 13.10.

Teradata Database 13.10 supports:

- Microsoft Windows Server 2003 64-bit
- SUSE Linux Enterprise Server 10

Teradata Database client applications can support other operating systems.

Prerequisites

To gain an understanding of Teradata, you should be familiar with the following:

- Basic computer technology
- System hardware
- Teradata Tools and Utilities

Changes to This Book

Release	Description
Teradata Database 13.10 August 2010	Added the following: <ul style="list-style-type: none">• SQL UDFs• Utility management Updated the Viewpoint section.
Teradata Database 13.0 April 2009	<ul style="list-style-type: none">• Added the following:<ul style="list-style-type: none">• AWTs and performance groups to the resources you can monitor• DBC.AMPUsage view has statistics for a user/account string pair• DDL replication and replication rule sets• Down subtable recovery• Dynamic UDTs and Java UDFs• Ferret utility• Information on using DBS Control to change the value of the PrimaryIndexDefault General field• No Primary Index table• Period and Geospatial to SQL data types• Session statuses• Storage subpools• Teradata Viewpoint• Unicode and compatibility system views• Updated the following:<ul style="list-style-type: none">• DEPOT to be part of the file system• Workload class conditions

Additional Information

URL	Description
www.info.teradata.com/	<p>Use the Teradata Information Products Publishing Library site to:</p> <ul style="list-style-type: none"> View or download a manual: <ol style="list-style-type: none"> Under Online Publications, select General Search. Enter your search criteria and click Search. Download a documentation CD-ROM: <ol style="list-style-type: none"> Under Online Publications, select General Search. In the Title or Keyword field, enter <i>CD-ROM</i>, and click Search. Order printed manuals: <p>Under Print & CD Publications, select How to Order.</p>
www.teradata.com	<p>The Teradata home page provides links to numerous sources of information about Teradata. Links include:</p> <ul style="list-style-type: none"> Executive reports, case studies of customer experiences with Teradata, and thought leadership Technical information, solutions, and expert advice Press releases, mentions and media resources
www.teradata.com/t/TEN/	<p>Teradata Customer Education designs, develops and delivers education that builds skills and capabilities for our customers, enabling them to maximize their Teradata investment.</p>
www.teradataatyourservice.com	<p>Use Teradata @ Your Service to access Orange Books, technical alerts, and knowledge repositories, view and join forums, and download software patches.</p>
developer.teradata.com/	<p>Teradata Developer Exchange provides articles on using Teradata products, technical discussion forums, and code downloads.</p>

To maintain the quality of our products and services, we would like your comments on the accuracy, clarity, organization, and value of this document. Please e-mail: teradata-books@lists.teradata.com.

Table of Contents

Preface	3
Purpose	3
Audience	3
Supported Software Releases and Operating Systems	3
Prerequisites	4
Changes to This Book.....	4
Additional Information	5

Chapter 1: The Data Warehouse	19
About the Data Warehouse	19
About Active Data Warehouse.....	19
Strategic Queries	20
Tactical Queries	20
The Teradata Solution	20
Active Load	20
Active Access.....	21
Active Events	21
Active Workload Management.....	21
Active Enterprise Integration	21
Active Availability	22

SECTION 1 Teradata Overview

Chapter 2: Teradata	25
About Teradata Database.....	25
Designing Teradata Database	25
Attachment Methods.....	25
Teradata SQL	26

Character Support	26
Teradata Database as a “Single Data Store”	26
Teradata Database Capabilities	27
Teradata Database Software.....	27
Software Installation	28
Teradata Tools and Utilities.....	28
Supported Platforms	28
Installation Guides for Teradata Tools and Utilities.....	29
Application Programming Interfaces.....	29
Channel-Attached Connectivity Tools.....	29
Language Preprocessors.....	30
Load and Unload Utilities.....	30
Teradata Active System Management	32
Teradata Analyst Pack.....	32
Teradata Database Management Tools	33
Teradata Viewpoint.....	33
Teradata Meta Data Services.....	33
Storage Management Utilities.....	34

Chapter 3: Teradata Database Model

About the Relational Model.....	37
About the Relational Database.....	37
Set Theory and Relational Database Terminology	37
Tables, Rows, and Columns.....	38
Table Constraints.....	38
Rows and Columns	38

SECTION 2 Teradata Database Architecture

Chapter 4: Teradata Database Hardware and Software Architecture

SMP and MPP Platforms	41
The BYNET.....	42
Boardless BYNET.....	42
Disk Arrays	42

Logical Units	42
Vdisks	42
Cliques	43
Hot Standby Nodes	44
Virtual Processors	44
Access Module Processor	45
AMP Clusters	46
Parsing Engine	46
Request Processing	47
The Dispatcher	48
The AMPs	48
Example: SQL Request	49
Parallel Database Extensions	50
Teradata Database File System	51
Workstation Types and Available Platforms	51
System Console	51
Administration Workstation	51
Teradata Database Window	51
How Database Window Communicates with Teradata Database	52
Running DBW	52
Teradata Generic Security Service	52

Chapter 5: Teradata Database RASUI 55

Software Fault Tolerance	55
Vproc Migration	55
Fallback Tables	56
AMP Clusters	57
One-Cluster Configuration	57
Smaller Cluster Configuration	57
Journaling	58
Backup Archive and Restore	59
Table Rebuild Utility	60
Hardware Fault Tolerance	61
Teradata Replication Services Using Oracle GoldenGate	62

Chapter 6: Communication Between the Client and Teradata Database	65
Attachment Methods	65
Network Attachment Methods	65
.NET Data Provider	66
Java Database Connectivity	66
OLE DB Provider for Teradata	66
Open Database Connectivity	66
Teradata CLIv2 for Network-Attached Systems	66
Channel Attachment Method	68
Teradata CLIv2 for Channel-Attached Systems	68
Teradata Director Program	68
Teradata Database Server	68

SECTION 3 Using Teradata Database

Chapter 7: Database Objects, Databases, and Users	73
Tables	73
Table Types	73
Views	74
Creating Views	74
Benefits of Using Views	75
Restrictions on Using Views	75
SQL Stored Procedures	75
Using SQL Stored Procedures	75
Elements of an SQL Stored Procedure	76
External Stored Procedures	76
Usage	76
Macros	77
SQL Statements Related to Macros	77
Single-User and Multi-User Macros	77
Macro Processing	77
Triggers	78
Types of Triggers	78
When to Fire Triggers	78
ANSI-Specified Order	79

Using Triggers	79
User-Defined Functions	79
SQL UDFs	79
External UDFs	80
Usage	80
Related Topics	81
User-Defined Methods	81
Instance Methods	81
Constructor Methods	81
User-Defined Types	81
UDT Types	82
Databases and Users	82
Creating Databases and Users	82
Creating a Finance and Administration Database	83

Chapter 8: SQL

Using SQL	87
Types of SQL Statements	87
Data Definition Language Statements	87
Data Control Language Statements	88
Data Manipulation Language Statements	89
SQL Statement Syntax	90
Statement Execution	91
Statement Punctuation	91
The SELECT Statement	92
SELECT Statement and Set Operators	92
SELECT Statement and Joins	92
SQL Data Types	93
Data Types	93
Data Type Phrase	93
Data Type Attributes	94
Teradata Database Recursive Query	94
SQL Functions	95
Scalar Functions	95
Aggregate Functions	95
Ordered Analytical Functions	96
Cursors	96
Session Modes	97

Chapter 9: SQL Application Development99

SQL Applications	99
Client Applications	99
Embedded SQL Applications	99
Macros as SQL Applications	100
SQL Used to Create a Macro.....	100
Macro Usage.....	101
SQL Used to Modify a Macro.....	101
SQL Used to Delete a Macro.....	101
SQL Stored Procedures as SQL Applications.....	102
SQL Used to Create Stored Procedures.....	102
SQL Stored Procedure Example.....	102
SQL Used to Execute a Stored Procedure.....	103
DDL Statements with Stored Procedures.....	103
The EXPLAIN Request Modifier.....	104
How EXPLAIN Works.....	104
Benefits of Using EXPLAIN.....	104
Simple EXPLAIN Example.....	104
Third-Party Development	107
TS/API.....	107
Compatible Third-Party Software Products.....	107
Workload Management Application Programming Interface.....	107

Chapter 10: Data Distribution and Data Access Methods111

Teradata Database Indexes.....	111
Primary Indexes	111
Primary Indexes and Data Distribution.....	112
Primary Key	112
Foreign Key.....	112
Relationships Between Primary Indexes and Primary Keys.....	112
Partitioned Primary Indexes	113
Multilevel Partitioned Primary Index	114
Comparing Partitioned and Nonpartitioned Primary Indexes.....	114
Secondary Indexes.....	114
Secondary Index Subtables	114
Comparing Primary and Secondary Indexes.....	115
Join Indexes.....	115
Single-table Join Indexes.....	115

Multitable Join Indexes	116
Aggregate Join Indexes	116
Sparse Join Indexes	116
Hash Indexes	117
Index Specification	117
Creating Indexes	117
Strengths and Weaknesses of Various Types of Indexes	118
Hashing	120
Identity Column	120
Normalization	120
Normal Forms	120
First, Second, and Third Normal Forms	121
Referential Integrity	122
Referential Integrity Terminology	122
Referencing (Child) Table	122
Referenced (Parent) Table	122
Importance of Referential Integrity	122

Chapter 11: Concurrency Control and Transaction Recovery

About Concurrency Control	125
Transactions	125
Definition of a Transaction	125
Definition of Serializability	125
Transaction Semantics	126
ANSI Mode Transactions	126
Teradata Mode Transactions	127
Locks	127
Overview of Teradata Database Locking	127
Locking Database Management Systems	128
Lock Levels	128
Locking Severities	129
Automatic Database Lock Levels	129
Deadlocks and Deadlock Resolution	130
Host Utility Locks	130
HUT Lock Types	130
HUT Lock Characteristics	131
Recovery and Transactions	131
System and Media Recovery	131

System Restarts	132
Transaction Recovery	132
Down AMP Recovery	132
Down Subtable Recovery.....	133
Two-Phase Commit Protocol	133

Chapter 12: The Data Dictionary135

Data Dictionary.....	135
Data Dictionary Users	135
Data Dictionary Views	140
Users of Data Dictionary Views	140
SQL Access to the Data Dictionary	141

Chapter 13: International Language Support143

Character Set Overview	143
About Repertoire	143
Character Representation	144
External and Internal Character Sets.....	144
Character Data Translation.....	144
What Teradata Database Supports.....	144
Teradata Database Character Data Storage	145
Internal Server Character Sets.....	145
User Data	145
Object Names in the Data Dictionary	145
Language Support Modes.....	145
Overriding the Default Character Set for User Data.....	146
Standard Language Support Mode	146
LATIN Character Set.....	146
Compatible Languages	146
Japanese Language Support Mode	147
Advantages of Storing User Data Using UNICODE.....	147
User DBC Default Character Set	147
Extended Support.....	147

Chapter 14: Query and Database Analysis Tools 149

Teradata Visual Explain	149
Teradata System Emulation Tool	150
Teradata Index Wizard	150
Demographics	151
Teradata Statistics Wizard	151
Query Capture Facility	152
Target Level Emulation	152
Database Query Log	153
Database Object Use Count	153

Chapter 15: Teradata Database Security 157

Security Concepts	157
Users	157
Permanent Database Users	158
Directory-based Users	158
Proxy Users	158
Database Privileges	158
Directly Granted Privileges	159
Roles	159
External Roles	159
Profiles	160
User Authentication	160
Authentication Method	160
Logon Format	161
Logon Controls	162
Password Format Requirements	163
Password Controls	163
User Authorization	163
Authorization of Permanent Database Users	163
Authorization of Directory-Based Users	163
Authorization of Middle-Tier Application Users	164
Data Protection	164
Directory Management of Users	165
Supported Directories	165
Database Security Monitoring	165
Security Monitoring	166

Defining a Security Policy	167
Publishing a Security Policy	167

SECTION 4 Managing and Monitoring Teradata Database

Chapter 16: System Administration.....171

Session Management	171
Session Requests.....	171
Establishing a Session	171
Logon Operands.....	172
Administrative and Maintenance Utilities	172

Chapter 17: Database Management Tools and Utilities177

Data Archiving Utilities	177
Teradata Archive/Recovery Utility.....	178
Teradata Parallel Transporter	178
Teradata Parallel Transporter Application Programming Interface	179
Standalone Data Load and Unload Utilities.....	179
Teradata FastExport.....	179
Teradata FastLoad	180
Teradata MultiLoad.....	180
Teradata Parallel Data Pump	180
Access Modules	181
Basic Teradata Query	182
Session and Configuration Management Tools.....	182
System Resource and Workload Management Tools and Protocols	183
Write Ahead Logging.....	183
Ferret Utility.....	183
Priority Scheduler	184
Teradata Active System Management	184
Teradata SQL Assistant.....	188

Chapter 18: Aspects of System Monitoring 191

Teradata Viewpoint	191
QUERY STATE Command	192
Resource Usage Monitoring.....	192
Resource Usage Data Gathering.....	193
How to Control Collection and Logging of ResUsage Data.....	193
ResUsage Tables and Views.....	194
ResUsage Data Categories.....	194
ResUsage Macros	194
Summary Mode	194
Performance Monitoring.....	194
Account String Expansion.....	194
The TDPTMON.....	195
System Management Facility.....	195
Workload Management Application Programming Interface.....	195

Chapter 19: Teradata Meta Data Services 197

About Metadata.....	197
Types of Metadata.....	197
Teradata Meta Data Services	198
Creating Teradata Meta Data Repository	199
Connecting to Teradata Meta Data Repository.....	199

Glossary 201

Index..... 205

CHAPTER 1 The Data Warehouse

This chapter presents an overview of Teradata.

About the Data Warehouse

Initially, the data warehouse was a *historical* database, enterprise-wide and centralized, containing data derived from an operational database.

The data in the data warehouse was:

- Subject-oriented
- Integrated
- Usually identified by a timestamp
- Nonvolatile, that is, nothing was added or removed

Rows in the tables supporting the operational database were loaded into the data warehouse (the historical database) after they exceeded some well-defined date.

Data could be queried, but the responses returned only reflected historical information. In this sense, a data warehouse was initially *static*, and even if a historical data warehouse contained data that was being updated, it would still not be an *active* data warehouse.

About Active Data Warehouse

An active data warehouse:

- Provides a single up-to-date view of the enterprise on one platform.
- Represents a logically consistent store of detailed data available for strategic, tactical, and event driven business decision making.
- Relies on timely updates to the critical data (as close to real time as needed).
- Supports short, tactical queries that return in seconds, alongside of traditional decision support.

Strategic Queries

Strategic queries represent business questions that are intended to draw strategic advantage from large stores of data.

Strategic queries are often complex, involving aggregations and joins across multiple tables in the database. They are sometimes long-running and tend not to have a strict service level expectation.

Strategic queries are often ad hoc. They may require significant database resources to execute, they are often submitted from third-party tools, and they can take advantage of session pooling.

Tactical Queries

Tactical queries are short, highly tuned queries that facilitate action-taking or decision-making in a time-sensitive environment. They usually come with a clear service level expectation and consume a very small percentage of the overall system resources.

Tactical queries are usually repetitively executed and take advantage of techniques such as request (query plan) caching and session-pooling.

The Teradata Solution

In an active data warehouse, Teradata provides both strategic intelligence and operational intelligence.

- Strategic intelligence entails delivering intelligence through tools and utilities and query mechanisms that support strategic decision-making.

This includes, for example, providing users with simple as well as complex reports throughout the day which indicate the business trends that have occurred and that are occurring, which show why such trends occurred, and which predict if they will continue to occur.

- Operational intelligence entails delivering intelligence through tools and utilities and query mechanisms that support front-line or operational decision-making.

This includes, for example, ensuring aggressive Service Level Goals (SLGs) with respect to high performance, data freshness, and system availability.

Active Load

Teradata is able to load data actively and in a non-disruptive manner and, at the same time, process other workloads.

Teradata delivers Active Load through methods that support continuous data loading. These include streaming from a queue, more frequent batch updates, and moving changed data from another database platform to Teradata.

These methods exercise such Teradata Database features as queue tables and triggers, and use FastLoad, MultiLoad, TPump, standalone utilities, and Teradata Parallel Transporter.

Teradata can effectively manage a complex workload environment on a “single version of the business.”

Active Access

Teradata is able to access analytical intelligence quickly and consistently in support of operational business processes.

But the benefit of Active Access entails more than just speeding up user and customer requests. Active Access provides intelligence for operational and customer interactions consistently.

Active Access queries, also referred to as tactical queries, support tactical decision-making at the front-line. Such queries can be informational, such as simply retrieving a customer record or transaction, or they may include complex analytics.

Active Events

Teradata is able to detect a business event automatically, apply business rules against current and historical data, and initiate operational actions when appropriate. This enables enterprises to reduce the latency between the identification of an event and taking action with respect to it. Active Events entails more than event detection.

When notified of something important, Teradata presents users with recommendations for appropriate action. The analysis done for users may prescribe the best course of action or give them alternatives from which to choose.

Active Workload Management

Teradata is able to manage mixed workloads dynamically and to optimize system resource utilization to meet business goals.

Teradata Active System Management (ASM) is a portfolio of products that enables real-time system management.

Teradata ASM assists the database administrator in analyzing and establishing workloads and resource allocation to meet business needs. Teradata ASM facilitates monitoring workload requests to ensure that resources are used efficiently and that dynamic workloads are prioritized automatically.

Teradata ASM also provides state-of-the-art techniques to visualize the current operational environment and to analyze long-term trends. Teradata ASM enables database administrators to set SLGs, to monitor adherence to them, and to take any necessary steps to reallocate resources to meet business objectives.

Active Enterprise Integration

Teradata is able to integrate itself into enterprise business and technical architectures, especially those that support business users, partners, and customers. This simplifies the task of coordinating enterprise applications and business processes.

For example, a Teradata event, generated from a database trigger, calls a stored procedure, which inserts a row into a queue table and publishes a message via the Teradata JMS Provider. The message is delivered to a JMS queue on a WebLogic, SAP NetWeaver, or other JMS-compatible application server. SAP Customer Relationship Management receives the message, notifies the user, and takes an action.

Active Availability

Teradata is able to meet business objectives for its own availability. Moreover, it assists customers in identifying application-specific availability, recoverability, and performance requirements based on the impact of enterprise downtime. Teradata can also recommend strategies for achieving system availability goals.

SECTION 1 **Teradata Overview**

This chapter presents an overview of Teradata and its components.

About Teradata Database

Teradata Database is an information repository supported by tools and utilities that make it a complete and active relational database management system.

Designing Teradata Database

Teradata developers designed Teradata Database from mostly off-the-shelf hardware components. The result was an inexpensive, high-quality system that exceeded the performance of conventional relational database management systems. The hardware components of Teradata Database evolved from those of a simple parallel database computer into those of a general-purpose, massively parallel computer running the database.

The architecture supports both single-node, Symmetric Multiprocessing (SMP) systems and multinode, Massively Parallel Processing (MPP) systems in which the distributed functions communicate by means of a fast interconnect structure. The interconnect structure is the BYNET for MPP systems and the boardless BYNET for SMP systems.

Attachment Methods

To support its role in the active warehouse environment, Teradata Database can use either of two attachment methods to connect to other operational computer systems as illustrated in the following table.

This attachment method...	Allows the system to be attached...
network	to intelligent workstations and other computers and devices through a Local Area Network (LAN).
channel	directly to an I/O channel of a mainframe computer.

Teradata SQL

SQL is the language of relational database communication. Teradata SQL, which is broadly compatible with ANSI SQL, extends the capabilities of SQL by adding Teradata-specific extensions to the generic SQL statements.

You can run transactions in either Teradata or ANSI mode and these modes can be set or changed.

Character Support

Teradata Database has an international customer base. To accommodate communications in different languages, Teradata Database supports non-Latin character sets, including KANJI, KANJISJIS, and UNICODE.

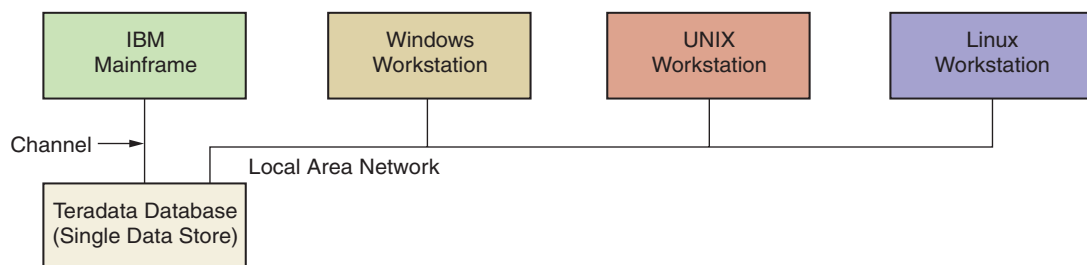
For detailed information about international character set support, see [Chapter 13: “International Language Support.”](#)

Teradata Database as a “Single Data Store”

A design goal of Teradata Database was to provide a *single data store* for a variety of client architectures. This approach greatly reduces data duplication and inaccuracies that can creep into data that is maintained in multiple stores.

This approach to data storage is known as the *single version of the business*, and Teradata Database implements this through heterogeneous client access. Clients can access a single copy of enterprise data and Teradata Database takes care of such things as data type translation, connections, concurrency, and workload management.

The following figure illustrates the idea of heterogeneous client access, where mainframe clients, network-attached workstations, and personal computers can access and manipulate the same database simultaneously. In this figure, the mainframe is attached via channel connections and other systems are attached via network connections.



1091H001

Teradata Database Capabilities

Teradata Database allows users to view and manage large amounts of data as a collection of related tables. Some of the capabilities of Teradata Database are listed in the following table.

Teradata Database provides...	That...
capacity	includes: <ul style="list-style-type: none"> Scaling from gigabytes to terabytes to petabytes of detailed data stored in billions of rows. Scaling to millions of millions of instructions per second (MIPS) to process data.
parallel processing	makes Teradata Database faster than other relational systems.
single data store	<ul style="list-style-type: none"> can be accessed by network-attached and channel-attached systems. supports the requirements of many diverse clients.
fault tolerance	automatically detects and recovers from hardware failures.
data integrity	ensures that transactions either complete or rollback to a consistent state if a fault occurs.
scalable growth	allows expansion without sacrificing performance.
SQL	serves as a standard access language that permits users to control data.

Teradata Database Software

Teradata Database software resides on the platform and implements the relational database environment. The platform software includes the following functional modules.

This module...	Provides...
Database Window	a tool that you can use to control the operation of Teradata Database.
Teradata Gateway	communications support. The Gateway software validates messages from clients that generate sessions over the network and it controls encryption.

This module...	Provides...
Parallel Data Extensions (PDE)	a software interface layer that lies between the operating system and database. PDE enables the database to operate in a parallel environment. For more information about PDE, see “ Parallel Database Extensions ” on page 50.
Teradata Database management software	the following modules: <ul style="list-style-type: none">• Parsing Engine (PE)• Access module processor (AMP)• Teradata Database file system For more information about Teradata Database file system, see “ Teradata Database File System ” on page 51.

Software Installation

The Parallel Upgrade Tool (PUT) automates much of the installation process for Teradata Database software. There are two major operational modes for PUT.

The operational mode...	Does the following...
Major upgrade	upgrades one or more software products to the next version.
Patch upgrade	applies patch packages to one or more software products.

Teradata Tools and Utilities

Teradata Tools and Utilities is a comprehensive suite of tools and utilities designed to operate in the client environment used to access Teradata Database.

Note: Teradata Database runs with or without channel- or network-attached client applications. The computer on which the utilities are installed can also run Teradata Database software.

Supported Platforms

Teradata Tools and Utilities applications are supported on a wide variety of operating platforms. For a list of supported platforms and product versions, see *Teradata Tools and Utilities 13.10 Supported Platforms and Product Versions*, which is available from <http://www.info.teradata.com/>.

Installation Guides for Teradata Tools and Utilities

To learn how to install Teradata Tools and Utilities, see the following installation guides:

- *Teradata Tools and Utilities Installation Guide for IBM z/OS*
- *Teradata Tools and Utilities Installation Guide for UNIX and Linux Systems*
- *Teradata Tools and Utilities Installation Guide for Microsoft Windows*

Application Programming Interfaces

Teradata provides a number of standardized interfaces to facilitate easy development of applications that access Teradata Database. The following table describes Teradata Application Programming Interfaces (APIs) and what each provides.

This utility for channel or network-attached clients or both ...	Provides...
.NET Data Provider for Teradata	an interface to Teradata Call-Level Interface version 2 (CLIV2) to connect, execute commands, and retrieve results from Teradata Database.
ODBC Driver for Teradata	access to Teradata Database from various tools, increasing the portability of access.
OLE DB Provider for Teradata	an interface for accessing and manipulating data stores that do not use SQL.
Teradata Call-Level Interface version 2 (CLIV2)	a Teradata proprietary API and library used by applications to communicate with Teradata Database.
Teradata Data Connector	a block-level I/O interface to one or more access modules that interface to hardware storage devices, software messaging systems, and OLE Database data sources.
Teradata JDBC Driver	platform-independent, Java-application access to Teradata Database from various tools increasing portability of data.

Channel-Attached Connectivity Tools

The following table describes tools and utilities on channel-attached clients and what each tool or utility provides.

This utility for channel-attached clients...	Provides...
Host Utility Consoles (HUTCNS)	access to a number of AMP-based utilities.
IBM® Customer Information Control System (CICS) Interface for Teradata	an interface that enables CICS macro or command-level application programs to access Teradata Database resources.

This utility for channel-attached clients...	Provides...
IBM IMS Interface for Teradata	an Information Management System (IMS) interface to Teradata Database.
Teradata CLIV2	a Teradata proprietary API and library used by applications to communicate with Teradata Database.
Teradata Director Program (TDP)	a high-performance interface for messages sent between the client and Teradata Database.
Teradata Transparency Series/Application Programming Interface (TS/API)	gateway services allowing products that access either DB2 or SQL/DS databases to access data stored on Teradata Database.

Language Preprocessors

Language preprocessors enable applications to access Teradata Database by interpreting SQL statements in C, COBOL, or Programming Language 1 (PL/I) programs. The following table describes the available Teradata preprocessors.

This preprocessor...	Provides a mechanism for...	For...
Teradata C Preprocessor	embedding SQL in C programs	channel- and network-attached clients.
Teradata COBOL Preprocessor	embedding SQL in COBOL programs	channel-attached and some network-attached clients.
Teradata PL/I Preprocessor	embedding SQL in PL/I programs	channel-attached clients.

Load and Unload Utilities

Teradata load and unload utilities offer a powerful solution for managing all your data load requirements from batch to real-time. Teradata load and unload utilities are fully parallel to provide optimal and scalable performance for getting data in and out of your Teradata Database. In addition, you can import and export data to and from host-based and client-resident data sources, including mainframe host databases, enterprise server databases or departmental data marts. The following table describes Teradata load and unload utilities.

This utility for channel- and network-attached clients...	Provides...
Basic Teradata Query (BTEQ)	a general-purpose, command-based tool that enables you to communicate with one or more Teradata Databases. BTEQ provides an interactive or batch interface that allows you to submit SQL statements, import and export data, and generate reports.
Teradata FastExport	a means of reading large volumes of data from Teradata Database.

This utility for channel- and network-attached clients...	Provides...
Teradata FastLoad	high-performance data loading from client files into empty tables.
Teradata MultiLoad	high-performance data maintenance, including inserts, updates, and deletions to existing tables.
Teradata Parallel Data Pump (TPump)	continuous update of tables; performs insert, update, and delete operations or a combination of those operations on multiple tables using the same source feed.
Teradata Parallel Transporter (Teradata PT)	a means to load data into and unload data from any accessible table in Teradata Database or from any other data stores for which an access operator or an access module exists.
Teradata Parallel Transporter Application Programming Interface (Teradata PT API)	a set of programming interfaces for loading and unloading data to and from Teradata Database systems. Teradata PT API enables an application to access the Teradata Database using proprietary Teradata load and export protocols. Because Teradata PT API is a functional library that is part of your client applications, it provides more control during the load and unload processes.

Teradata Access Modules

Access modules are adapters that allow all Teradata load and unload utilities to interface with a variety of data sources through standards-based interfaces. These standards-based modules let you read from a given data source as if you were reading from a flat file. The following table describes Teradata access modules.

This utility for channel- and network-attached clients...	Provides...
Named Pipes Access Module	an inter-process communication link between a writer process, such as Teradata FastExport, and a reader process, such as Teradata FastLoad.
Teradata Access Module for JMS	a fast, asynchronous method to import and export data between Teradata Database and any JMS-enabled messaging system. Teradata Access Module for JMS can be invoked by Teradata load and unload utilities.
Teradata OLE DB Access Module	a dynamic link library (DLL) that acts as an interface between Teradata load and export utilities and data sources for which an OLE DB provider is available. The access module quickly moves data between an OLE DB data source and Teradata Database without requiring intermediate storage.

This utility for channel- and network-attached clients...	Provides...
Teradata Data Connector	a software layer between a client utility and an access module. It establishes, maintains, and monitors the connection between a client utility and an access module by being statically linked to the client and dynamically linked to the access module.
Teradata WebSphere MQ Access Module	Teradata utilities to import data using IBM WebSphere MQ message queuing middleware.

Teradata Active System Management

Teradata Active System Management (ASM) is a portfolio of products designed for the analysis, organization and control of workloads inside the Teradata solution. These products help you keep pace with your increasingly complex production workloads, especially those dealing with critical business intelligence, analytics, or tactical queries. The following table describes Teradata Tools and Utilities products that are part of Teradata ASM and what each tool provides.

This tool for network-attached clients...	Provides...
Teradata Viewpoint	a means to set up rules (including workload limits on accounts, users, and objects, such as databases and tables) that manage database access, increase database efficiency, and enhance workload capacity.
Teradata Workload Analyzer (Teradata WA)	recommendations on workload definitions and operating rules to ensure that database performance meets Service Level Goals (SLGs).

Teradata Analyst Pack

As application environments expand to include mixed workloads for both decision support and near-real-time analytic processing, maximizing the performance of the Teradata Database becomes more challenging. Teradata provides the Teradata Analyst Pack to enable you to analyze and tune your database queries for better performance. The Teradata Analyst Pack simplifies your DBA and query planner's jobs by automating the steps required to analyze and optimize your Teradata Databases. The following table describes the Teradata Analyst Pack tools and what each tool provides.

This tool for network-attached clients...	Provides...
Teradata Index Wizard	analyses of various SQL query workloads and suggests candidate indexes to enhance performance of those queries.

This tool for network-attached clients...	Provides...
Teradata Statistics Wizard	automation for collecting workload statistics, or selecting recommended indexes or columns for statistics collection or re-collection.
Teradata System Emulation Tool (Teradata SET)	the capability to examine the query plans generated by the test system Optimizer as if the queries were processed on the production system.
Teradata Visual Explain (Teradata VE)	provides a graphical depiction of the execution plan of complex SQL statements.

Teradata Database Management Tools

You must ensure that your database can keep pace with changing requirements and the addition of new users and business applications. Teradata Database has a rich collection of tools and facilities to control database operation, administration, and maintenance.

The following table describes database management tools and what each tool provides.

This tool for network-attached clients...	Provides...
Teradata Administrator	an interface that you can use to perform database administration tasks.
Teradata Query Director (QD)	a program that routes sessions for high availability purposes.
Teradata Query Scheduler (QS)	a means to manage requests input to Teradata Database and keep the database running at optimum performance levels. The product consists of client, platform, and administrator components, plus a separate database within Teradata Database.
Teradata SQL Assistant for Microsoft Windows	a means of retrieving data from a Teradata Database or any ODBC-compliant database server. You can then view, manipulate, or store the extracted data on your PC. .NET Data Provider for Teradata can also be used to connect to a Teradata Database from Teradata SQL Assistant for Microsoft Windows.

Teradata Viewpoint

In addition to the Teradata Tools and Utilities database management tools, Teradata Viewpoint provides a suite of web-based applications that enable you to monitor and manage Teradata Database system performance from anywhere using a standard browser. For more information, see [“Teradata Viewpoint” on page 191](#).

Teradata Meta Data Services

Teradata Meta Data Services (MDS) is a comprehensive solution for managing metadata in data warehouse environments. This solution enables you to locate, consolidate, manage, and

navigate warehouse metadata. Teradata's single, unified approach makes sense to technical and business users alike, helping you reduce training and development costs while increasing end-user self-sufficiency, system utilization, and productivity.

The Teradata MDS utility for network-attached clients, provides an infrastructure for Teradata metadata and for creating tools to interchange metadata with external operational systems, Extraction Transformation and Load (ETL) tools, business intelligence tools, database modeling tools, and any other metadata sources.

Storage Management Utilities

The following table describes storage management utilities and what each utility provides.

This utility for channel- and network-attached clients...	Provides...
Archive/Recovery (Teradata ARC)	a means of archiving data to tape and disk and restoring data to Teradata Database.
Backup Application Software includes the following: <ul style="list-style-type: none">• NetVault® Plug-in for Teradata• NetBackup™ Teradata Extension• Tivoli® Storage Manager Teradata Extension	open architecture products for backup and restore functions for Windows and Linux clients.

For More Information

For more information on the topics presented in this chapter, see the following Teradata Database and Teradata Tools and Utilities books.

IF you want to learn more about...	See...
Archive/Recovery utility	<i>Teradata Archive/Recovery Utility Reference</i>
Backup/Archive/Recovery (BAR), including BAR framework products	<i>Teradata BAR Solutions Guide</i>
Teradata BTEQ	<i>Basic Teradata Query Reference</i>
CICS Interface for Teradata	<i>IBM CICS Interface for Teradata Reference</i>
Teradata CLIV2	<ul style="list-style-type: none">• <i>Teradata Call-Level Interface Version 2 Reference for Mainframe-Attached Systems</i>• <i>Teradata Call-Level Interface Version 2 Reference for Network-Attached Systems</i>
Host utility consoles	<ul style="list-style-type: none">• <i>Utilities</i>• <i>Database Administration</i>

IF you want to learn more about...	See...
IMS/DC Interface for Teradata	<i>IBM IMS/DC Interface for Teradata Reference</i>
Teradata JDBC Driver	<i>Teradata JDBC Driver User Guide</i>
Load and export utilities	<ul style="list-style-type: none"> • <i>Teradata FastExport Reference</i> • <i>Teradata FastLoad Reference</i> • <i>Teradata MultiLoad Reference</i> • <i>Teradata Parallel Data Pump Reference</i>
Teradata ODBC Driver	<i>ODBC Driver for Teradata User Guide</i>
OLE DB Provider for Teradata	<i>OLE DB Provider for Teradata User Guide</i>
SQL	<i>SQL Fundamentals</i>
Teradata Administrator	<i>Teradata Administrator User Guide</i>
Teradata Database architecture	<i>Database Design</i>
Teradata Data Connector	<i>Teradata Tools and Utilities Access Module Reference</i>
Teradata Director Program	<i>Teradata Director Program Reference</i>
Teradata Index Wizard	<i>Teradata Index Wizard User Guide</i>
Teradata Meta Data Services	<ul style="list-style-type: none"> • <i>Teradata Meta Data Services Administrator Guide</i> • <i>Teradata Meta Data Services Programmer Guide</i>
.NET Data Provider for Teradata	<i>.NET Data Provider for Teradata Release Definition</i>
Teradata Parallel Transporter, including Teradata Parallel Transport Application Programming Interface	<ul style="list-style-type: none"> • <i>Teradata Parallel Transporter Application Programming Interface Programmer Guide</i> • <i>Teradata Parallel Transporter Operator Programmer Guide</i> • <i>Teradata Parallel Transporter Reference</i> • <i>Teradata Parallel Transporter User Guide</i>
Teradata Preprocessor2 for Embedded SQL	<ul style="list-style-type: none"> • <i>Teradata Preprocessor2 for Embedded SQL Programmer Guide</i> • <i>SQL Stored Procedures and Embedded SQL</i>
Teradata Query Director	<i>Teradata Query Director User Guide</i>
Teradata Query Scheduler	<ul style="list-style-type: none"> • <i>Teradata Query Scheduler Administrator Guide</i> • <i>Teradata Query Scheduler User Guide</i>
Teradata SQL Assistant	<i>Teradata SQL Assistant for Microsoft Windows User Guide</i>
Teradata Statistics Wizard	<i>Teradata Statistics Wizard User Guide</i>
Teradata System Emulation Tool	<ul style="list-style-type: none"> • <i>SQL Request and Transaction Processing</i> • <i>Teradata System Emulation Tool User Guide</i>

IF you want to learn more about...	See...
Teradata Access Modules	<ul style="list-style-type: none">• <i>Teradata Tools and Utilities Access Module Programmer Guide</i>• <i>Teradata Tools and Utilities Access Module Reference</i>
Teradata Visual Explain	<i>Teradata Visual Explain User Guide</i>
Teradata Workload Analyzer	<i>Teradata Workload Analyzer User Guide</i>
TS/API products	<i>Teradata Transparency Series/ Application Programming Interface User Guide</i>

CHAPTER 3 Teradata Database Model

This chapter describes the concepts on which relational databases are modeled and discusses some of the objects that are part of a relational database.

About the Relational Model

The relational model for database management is based on concepts derived from the mathematical theory of sets. Basically, set theory defines a table as a relation. The number of rows is the *cardinality* of the relation, and the number of columns is the *degree*. Any manipulation of a table in a relational database has a consistent, predictable outcome, because the mathematical operations on relations are well-defined.

By way of comparison, database management products based on hierarchical, network, or object-oriented architectures are not built on rigorous theoretical foundations. Therefore, the behavior of such products is not as predictable or as flexible as that of relational products.

The SQL Optimizer in the database builds the most efficient access path to requested data. The Optimizer can readily adapt to changes in system variables by rebuilding access paths without programmer intervention. This adaptability is necessary because database definitions and data demographics can change.

About the Relational Database

Users perceive a relational database as a collection of objects, that is, as tables, views, macros, stored procedures, and triggers that are easily manipulated using SQL directly or specifically developed applications.

Set Theory and Relational Database Terminology

Relational databases are a generalization of the mathematics of set theory relations. Thus, the correspondences between set theory and relational databases are not always direct. The following table notes the correspondence between set theory and relational database terms.

Set Theory Term	Relational Database Term
Relation	Table
Tuple	Row
Attribute	Column

Tables, Rows, and Columns

Tables are two-dimensional objects consisting of rows and columns. Data is organized in tabular format and presented to the users of a relational database. References between tables define the relationships and constraints of data inside the tables themselves. For more information on different types of tables, see [“Tables” on page 73](#).

Table Constraints

You can define conditions that must be met before Teradata Database writes a given value to a column in a table. These conditions are called *constraints*. Constraints can include value ranges, equality or inequality conditions, and intercolumn dependencies. Teradata Database supports constraints at both the column and table levels.

During table creation and modification, you can specify constraints on single-column values as part of a column definition or on multiple columns using the CREATE and ALTER TABLE statements.

Rows and Columns

A column always contains the same kind of information. For example, a table that has information about employees would have a column for last name and nothing other than the employee last names should be placed in that column.

A row is one instance of all the columns in a table. For example, each row in the employee table would contain, among other things, the first name and the last name for that employee. The columns in a table represent entities, relationships, or attributes.

An *entity* is a person, place, or thing about which the table contains information. The table mentioned in the previous paragraphs contains information about the employee entity. Each table holds only one kind of row. The relational model requires that each row in a table be uniquely identified. To accomplish this, you define a uniqueness constraint to identify each row in the table. For more information about primary keys, see [“Relationships Between Primary Indexes and Primary Keys” on page 112](#).

For More Information

For more information on the topics presented in this chapter, see *Database Design*.

SECTION 2 **Teradata Database Architecture**

CHAPTER 4 Teradata Database Hardware and Software Architecture

This chapter briefly describes Teradata Database hardware components and software architecture.

The hardware that supports Teradata Database software is based on Symmetric Multiprocessing (SMP) technology.

The hardware can be combined with a communications network that connects the SMP systems to form Massively Parallel Processing (MPP) systems.

SMP and MPP Platforms

The components of the SMP and MPP hardware platforms include the following.

Component	Description	Function
Processor Node	<p>A hardware assembly containing several, tightly coupled, central processing units (CPUs) in an SMP configuration. An SMP node is connected to one or more disk arrays with the following installed on the node:</p> <ul style="list-style-type: none">• Teradata Database software• Client interface software• Operating system• Multiple processors with shared-memory• Failsafe power provisions <p>An MPP configuration is a configuration of two or more loosely coupled SMP nodes.</p>	Serves as the hardware platform upon which the database software operates.
BYNET	<p>Hardware interprocessor network to link nodes on an MPP system.</p> <p>Note: Single-node SMP systems use a software-configured virtual BYNET driver to implement BYNET services.</p>	Implements broadcast, multicast, or point-to-point communication between processors, depending on the situation.

These platforms use virtual processors (vprocs) that run a set of software processes on a node under the Parallel Database Extensions (PDE). For information about PDE, see [“Parallel Database Extensions” on page 50](#).

Vprocs provide the parallel environment that enables Teradata Database to run on SMP and MPP systems. For more information on vprocs, see [“Virtual Processors” on page 44](#).

The BYNET

At the most elementary level, you can look at the BYNET as a switched fabric that loosely couples all the SMP nodes in a multinode system. But the BYNET has capabilities that range far beyond those of a simple system bus.

The BYNET possesses high-speed logic that provides bi-directional broadcast, multicast, and point-to-point communication and merge functions.

A multinode system has at least two BYNETs. This creates a fault-tolerant environment and enhances interprocessor communication. Load-balancing software optimizes transmission of messages over the BYNETs. If one BYNET should fail, the second can handle the traffic.

Boardless BYNET

Single-node SMP systems use Boardless BYNET (or virtual BYNET) software to provide the same functions without the presence of BYNET hardware.

Disk Arrays

Teradata Database employs Redundant Array of Independent Disks (RAID) storage technology to provide data protection at the disk level. You use the RAID management software to group disk drives into RAID LUNS (Logical Units) to ensure that data is available in the event of a disk failure. Redundant implies that either data, functions, or components are duplicated in the architecture of the array.

Logical Units

The RAID Manager uses drive groups. A drive group is a set of drives that have been configured into one or more LUNs. Each LUN is uniquely identified.

The operating system recognizes a LUN as a disk and is not aware that it is actually writing to spaces on multiple disk drives. This technique allows RAID technology to provide data availability without affecting the operating system.

Vdisks

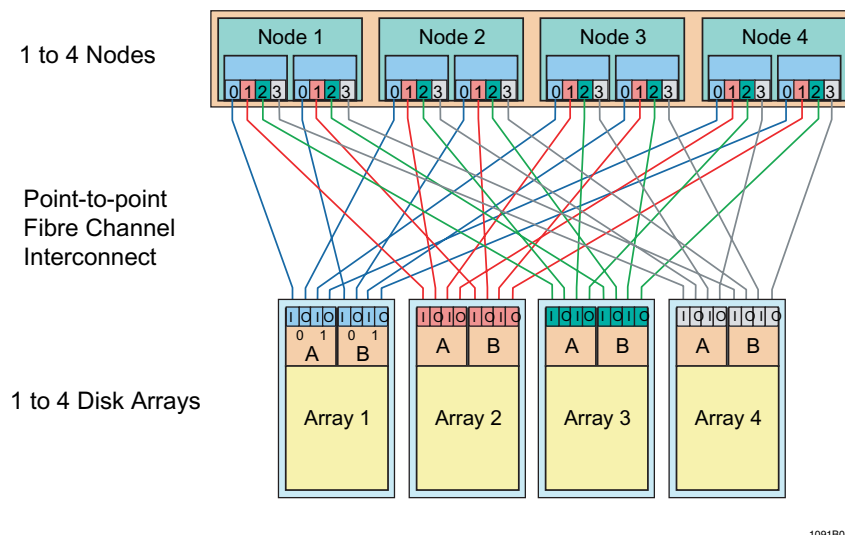
The group of cylinders currently assigned to an AMP is referred to as a vdisk, although the actual physical storage may derive from several different storage devices. For information about the role that AMPs play in Teradata Database architecture, see [“Virtual Processors” on page 44](#).

Cliques

The clique is a feature of some MPP systems that physically group nodes together by multiported access to common disk array units. Inter-node disk array connections are made using FibreChannel (FC) buses.

FC paths enable redundancy to ensure that loss of a processor node or disk controller does not limit data availability. The nodes do not share data. They only share access to the disk arrays. The following figure illustrates a four-node clique.

4-Node Clique



A clique is the mechanism that supports the migration of vprocs under PDE following a node failure. If a node in a clique fails, then vprocs migrate to other nodes in the clique and continue to operate while recovery occurs on their home node. For more detailed information on vprocs see [“Virtual Processors” on page 44](#).

PEs that manage physical channel connections *cannot* migrate because they are dependent on the hardware that is physically attached to the node to which they are assigned.

PEs for LAN-attached connections *do* migrate when a node failure occurs, as do all AMPs.

Hot Standby Nodes

Hot standby nodes allow spare nodes to be incorporated into the production environment. Teradata Database can use spare nodes to improve availability and maintain performance levels in the event of a node failure. A hot standby node is a node that:

- Is a member of a clique.
- Does not normally participate in Teradata Database operations.
- Can be brought in to participate in Teradata Database operations to compensate for the loss of a node in the clique.

Configuring a hot standby node can eliminate the system-wide performance degradation associated with the loss of a node. A hot standby node is added to each clique in the system. When a node fails, all AMPs and all LAN-attached PEs on the failed node migrate to the node designated as the hot standby. The hot standby node becomes a production node. When the failed node returns to service, it becomes the new hot standby node.

Configuring hot standby nodes eliminates:

- Restarts that are required to bring a failed node back into service.
- Degraded service when vprocs have migrated to other nodes in a clique.

Virtual Processors

The versatility of Teradata Database is based on virtual processors (vprocs) that eliminate dependency on specialized physical processors. Vprocs are a set of software processes that run on a node under Teradata Parallel Database Extensions (PDE) within the multitasking environment of the operating system.

The following table contains information about the different types of vprocs.

Vproc Type	Description
AMP	Access module processors perform database functions, such as executing database queries. Each AMP owns a portion of the overall database storage.
GTW	Gateway vprocs provide a socket interface to Teradata Database.
Node	The node vproc handles PDE and operating system functions not directly related to AMP and PE work. Node vprocs cannot be externally manipulated, and do not appear in the output of the Vproc Manager utility.
PE	Parsing engines perform session control, query parsing, security validation, query optimization, and query dispatch.
RSG	Relay Services Gateway provides a socket interface for the replication agent, and for relaying dictionary changes to the Teradata Meta Data Services utility.

Vproc Type	Description
VSS	Manages Teradata Database storage. AMPs acquire their portions of database storage through the TVS vproc.

A single system can support a maximum of 16,384 vprocs. The maximum number of vprocs per node can be as high as 128, but is typically between 6 and 12.

Each vproc is a separate, independent copy of the processor software, isolated from other vprocs, but sharing some of the physical resources of the node, such as memory and CPUs. Multiple vprocs can run on an SMP platform or a node.

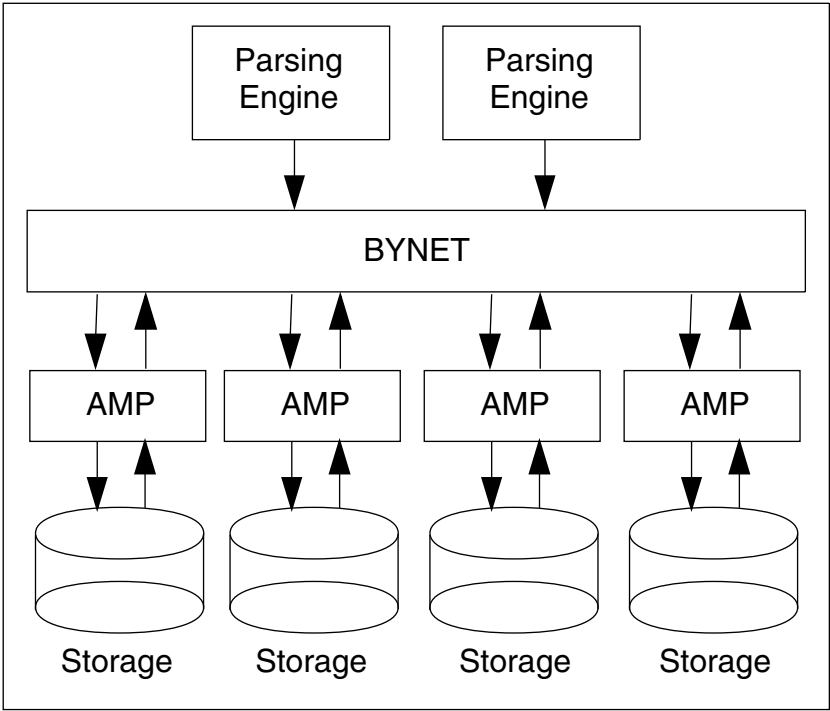
Vprocs and the tasks running under them communicate using unique-address messaging, as if they were physically isolated from one another. This message communication is done using the Boardless BYNET Driver software on single-node platforms or BYNET hardware and BYNET Driver software on multinode platforms.

Access Module Processor

The AMP vproc manages Teradata Database interactions with the disk subsystem. Each AMP manages a share of the disk storage.

AMP functions include...	For example...
database management tasks	<ul style="list-style-type: none"> Accounting Journaling Locking tables, rows, and databases Output data conversion
	During query processing: <ul style="list-style-type: none"> Sorting Joining data rows Aggregation
file system management	disk space management.

Each AMP, as represented in the following figure, manages a portion of the physical disk space. Each AMP stores its portion of each database table within that space.



1091C022

AMP Clusters

AMPs are grouped into logical clusters to enhance the fault-tolerant capabilities of Teradata Database. For more information on this method of creating additional fault tolerance in a system see [Chapter 5: “Teradata Database RASUL.”](#)

Parsing Engine

The PE is the vproc that communicates with the client system on one side and with the AMPs (via the BYNET) on the other side.

Each PE executes the database software that manages sessions, decomposes SQL statements into steps, possibly in parallel, and returns the answer rows to the requesting client.

The PE software consists of the following elements.

Parsing Engine Elements	Process
Parser	Decomposes SQL into relational data management processing steps.
Optimizer	Determines the most efficient path to access data.
Generator	Generates and packages steps.

Parsing Engine Elements	Process
Dispatcher	Receives processing steps from the parser and sends them to the appropriate AMPs via the BYNET.
	Monitors the completion of steps and handles errors encountered during processing.
Session Control	<ul style="list-style-type: none"> Manages session activities, such as logon, password validation, and logoff. Recovers sessions following client or server failures.

Request Processing

SQL is the language that you use to make requests of Teradata Database, that is, you use SQL to query Teradata Database.

The SQL parser handles all incoming SQL requests in the following sequence:

- 1 The Parser looks in the Request cache to determine if the request is already there.

IF the request is...	THEN the Parser...
in the Request cache	<p>reuses the plastic steps found in the cache and passes them to gncApply. Go to step 8 after checking privileges (step 4).</p> <p>Plastic steps are directives to the database management system that do not contain data values.</p>
not in the Request cache	begins processing the request with the Syntaxer.

- 2 The Syntaxer checks the syntax of an incoming request.

IF there are...	THEN the Syntaxer...
no errors	converts the request to a parse tree and passes it to the Resolver.
errors	passes an error message back to the requestor and stops.

- 3 The Resolver adds information from the Data Dictionary (or cached copy of the information) to convert database, table, view, stored procedure, and macro names to internal identifiers.

- 4 The security module checks privileges in the Data Dictionary.

IF the privileges are...	THEN the Security module...
valid	passes the request to the Optimizer.
not valid	aborts the request and passes an error message and stops.

- 5 The Optimizer determines the most effective way to implement the SQL request.
- 6 The Optimizer scans the request to determine where to place locks, then passes the optimized parse tree to the Generator.
- 7 The Generator transforms the optimized parse tree into plastic steps, caches the steps if appropriate, and passes them to gncApply.
- 8 gncApply takes the plastic steps produced by the Generator, binds in parameterized data if it exists, and transforms the plastic steps into concrete steps.

Concrete steps are directives to the AMPs that contain any needed user- or session-specific values and any needed data parcels.
- 9 gncApply passes the concrete steps to the Dispatcher.

The Dispatcher

The Dispatcher controls the sequence in which steps are executed. It also passes the steps to the BYNET to be distributed to the AMP database management software as follows:

- 1 The Dispatcher receives concrete steps from gncApply.
- 2 The Dispatcher places the first step on the BYNET; tells the BYNET whether the step is for one AMP, several AMPs, or all AMPs; and waits for a completion response.

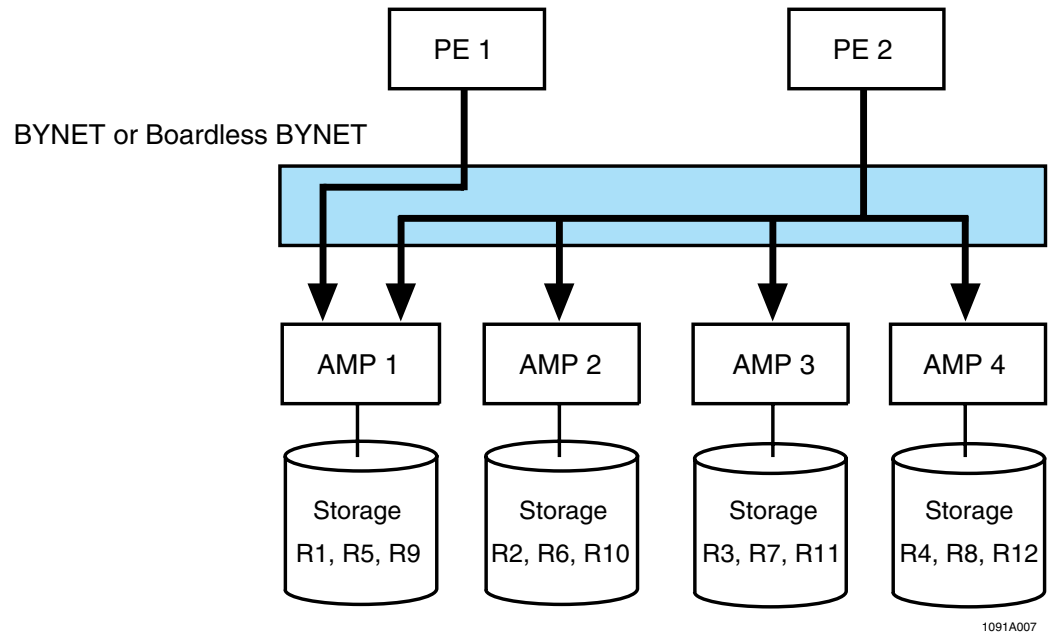
Whenever possible, Teradata Database performs steps in parallel to enhance performance. If there are no dependencies between a step and the following step, the following step can be dispatched before the first step completes, and the two execute in parallel. If there is a dependency, for example, the following step requires as input the data produced by the first step, then the following step cannot be dispatched until the first step completes.
- 3 The Dispatcher receives a completion response from all expected AMPs and places the next step on the BYNET. It continues to do this until all the AMP steps associated with a request are done.

The AMPs

AMPs obtain the rows required to process the requests (assuming that the AMPs are processing a SELECT statement). The BYNET transmits messages to and from the AMPs and PEs. An AMP step can be sent to one of the following:

- One AMP
- A selected set of AMPs, called a dynamic BYNET group
- All AMPs in the system

The following figure is based on the example in the next section. If access is through a primary index and a request is for a single row, the PE transmits steps to a single AMP, as shown at PE1. If the request is for many rows (an all-AMP request), the PE makes the BYNET broadcast the steps to all AMPs, as shown in PE2. To minimize system overhead, the PE can send a step to a subset of AMPs, when appropriate.



Example: SQL Request

As an example, consider the following Teradata SQL requests using a table containing checking account information. The example assumes that AcctNo column is the unique primary index for Table_01. For information about the types of indexes used by Teradata Database, see [Chapter 10: "Data Distribution and Data Access Methods."](#)

```
1. SELECT * FROM Table_01 WHERE AcctNo = 129317 ;
2. SELECT * FROM Table_01 WHERE AcctBal > 1000 ;
```

In this example:

- PEs 1 and 2 receive requests 1 and 2.
- The data for account 129317 is contained in table row R9 and stored on AMP1.
- Information about all account balances is distributed evenly among the disks of all four AMPs.

The sample Teradata SQL statement is processed in the following sequence:

- 1 PE 1 determines that the request is a primary index retrieval, which calls for the access and return of one specific row.
- 2 The Dispatcher in PE 1 issues a message to the BYNET containing an appropriate read step and R9/AMP 1 routing information. After AMP 1 returns the desired row, PE 1 transmits the data to the client.
- 3 The PE 2 Parser determines that this is an all-AMPs request, then issues a message to the BYNET containing the appropriate read step to be broadcast to all four AMPs.
- 4 After the AMPs return the results, PE 2 transmits the data to the client.

AMP steps are processed in the following sequence:

- 1 Lock—Serializes access in situations where concurrent access would compromise data consistency.
For some simple requests using Unique Primary Index (UPI), Nonunique Primary Index (NUPI), or Unique Secondary Index (USI) access, the lock step will be incorporated into step 2. For information about indexes and their uses, see [Chapter 10: “Data Distribution and Data Access Methods.”](#)
- 2 Operation—Performs the requested task. For complicated queries, there may be hundreds of operation steps.
- 3 End transaction—Causes the locks acquired in step 1 or 2 to be released.
The end transaction step tells all AMPs that worked on the request that processing is complete.

Parallel Database Extensions

Parallel Database Extensions (PDE) is a software interface layer that lies between the operating system and Teradata Database. PDE supports the parallelism that gives Teradata Database its speed and linear scalability. The operating system can be Linux or Microsoft Windows. PDE provides Teradata Database with the ability to:

- Run in a parallel environment
- Execute vprocs
- Apply a flexible priority scheduler to Teradata Database sessions
- Consistently manage memory, I/O, and messaging system interfaces across multiple OS platforms

PDE provides a series of parallel operating system services, which include:

- Facilities to manage parallel execution of database operations on multiple nodes.
- Dynamic distribution of database tasks.
- Coordination of task execution within and between nodes.

PDE enables MPP systems to take advantage of hardware features such as the BYNET and shared disk arrays.

Teradata Database File System

The file system is a layer of software between Teradata Database and PDE. File system service calls allow Teradata Database to store and retrieve data efficiently and with integrity without being concerned about the specific low-level operating system interfaces.

Workstation Types and Available Platforms

Workstations provide a window into the interworkings of Teradata Database. The following table shows the types of workstations available and their platforms.

Type of Workstation	Platform
System Console	SMP
AWS	MPP

System Console

The system console:

- Provides an input mechanism for the system and database administrators.
- Displays system status.
- Displays current system configuration.
- Displays performance statistics.
- Allows you to control various utilities.

Administration Workstation

The AWS can do everything a System Console can do, plus:

- Provide a single-system view in the multinode environment.
- Monitor system performance.

Teradata Database Window


Teradata DBW allows database administrators, system operators, and support personnel to control the operation of Teradata Database.

DBW is also the primary vehicle for starting and controlling the operation of Teradata Database utilities.

How Database Window Communicates with Teradata Database

DBW provides a graphical user interface to the Teradata Console Subsystem (CNS). Use DBW to issue database commands and run many of the database utilities. CNS is a part of the Parallel Database Extensions (PDE) software upon which the database runs.

The following figure illustrates the logical relationship among DBW, CNS, and Teradata Database.



1095D041

Running DBW

You can run DBW from the following locations:

- System Console
- Administration Workstation (AWS)
- Remote workstation or computer

To learn more about the DBW interface, see “Database Window (xdbw)” in *Utilities*.

Teradata Generic Security Service

Network security for Teradata is provided by Teradata Generic Security Service (TDGSS) software. It provides for secure communication between a client and Teradata Database.

For More Information

For more information on the topics presented in this chapter, see the following Teradata Database books.

IF you want to learn more about...	See...
SMP and MPP Platforms	<i>Performance Management</i>
Disk Arrays	<i>Database Administration</i>
Cliques	<i>Database Administration</i>
Hot Standby Nodes	<i>Database Administration</i>

IF you want to learn more about...	See...
Virtual Processors	<ul style="list-style-type: none"> • <i>Database Design</i> • <i>Database Administration</i> • <i>SQL Request and Transaction Processing</i>
Access Module Processor	<ul style="list-style-type: none"> • <i>Database Administration</i> • <i>Support Utilities</i>
Parsing Engine	<i>Database Administration</i>
Request Processing	<i>SQL Request and Transaction Processing</i>
Parallel Database Extensions	<i>Performance Management</i>
Teradata Database File System	<i>Utilities</i>
Workstation Types and Available Platforms	<ul style="list-style-type: none"> • AWS manuals (Visit www.info.teradata.com and search for the keyword “AWS.”) • <i>Performance Management</i>
Teradata Database Window	<i>Utilities</i>
Teradata General Security Service	<i>Security Administration</i>

CHAPTER 5 Teradata Database RASUI

Teradata Database addresses the critical requirements of reliability, availability, serviceability, usability, and installability (RASUI) by combining the following elements:

- Multiple microprocessors in a Symmetric Multi-Processing, (SMP) arrangement.
- RAID disk storage technology.
- Protection of Teradata Database from operating anomalies of the client platform.

Both hardware and software provide fault tolerance, some of which is mandatory and some of which is optional.

Software Fault Tolerance

This section explains the following Teradata Database facilities for software fault tolerance:


- Vproc migration
- Fallback tables
- AMP clusters
- Journaling
- Backup/Archive/Recovery
- Table Rebuild utility

Vproc Migration

Because the Parsing Engine (PE) and Access Module Processor (AMP) are vprocs and, therefore, software entities, they can migrate from their home node to another node within the same hardware clique if the home node fails for any reason. Although the system normally determines which vprocs migrate to which nodes, a user can configure preferred migratory destinations.

Vproc migration permits the system to function completely during a node failure, with some degradation of performance due to the non-functional hardware.

The following figure illustrates vproc migration, where the large X indicates a failed node, and arrows pointing to nodes still running indicate the migration of AMP3, AMP4, and PE2.



Note: PEs that manage physical channel connections cannot migrate during a node failure because they are dependent on the hardware that is physically attached to their assigned node.

Fallback Tables

A fallback table is a duplicate copy of a primary table. Each fallback row in a fallback table is stored on an AMP different from the one to which the primary row hashes. This storage technique maintains availability should the system lose an AMP and its associated disk storage in a cluster. In that event, the system would access data in the fallback rows.

The disadvantage of fallback is that this method doubles the storage space and the I/O (on INSERT, UPDATE, and DELETE statements) for tables. The advantage is that data is almost never unavailable because of one down AMP. Data is fully available during an AMP or disk outage, and recovery is automatic after repairs have been made.

Teradata Database permits the definition of fallback for individual tables. As a general rule, you should run all tables critical to your enterprise in fallback mode. You can run other, non-critical tables in nonfallback mode in order to maximize resource usage.

Even though RAID disk array technology may provide data access even when you have not specified fallback, neither RAID1 nor RAID5 provides the same level of protection as fallback.

You specify whether a table is fallback or not using the CREATE TABLE (or ALTER TABLE) statement. The default is *not* to create tables with fallback.

AMP Clusters

A cluster is a group of 2-16 AMPs that provide fallback capability for each other. A copy of each row is stored on a separate AMP *in the same* cluster. In a large system, you would probably create many AMP clusters. However, whether large or small, the concept of a cluster exists even if all the AMPs are in one cluster.

One-Cluster Configuration

Pictures best explain AMP clustering. The following figure illustrates a situation in which fallback is present with one cluster, which is essentially an unclustered system.

	AMP1	AMP2	AMP3	AMP4
Primary copy area	1,9,17	2,10,18	3,11,19	4,12,20
Fallback copy area	21,22,15	1,23,8	9,2,16	17,10,3

	AMP5	AMP6	AMP7	AMP8
Primary copy area	5,13,21	6,14,22	7,15,23	8,16,24
Fallback copy area	18,11,4	19,12,24	20,5,6	13,14,7

FG10A001

Note that the fallback copy of any row is always located on an AMP different from the AMP which holds the primary copy. This is an entry-level fault tolerance strategy. In this example which shows only a few rows, the data on AMP3 is fallback protected on AMPs 4, 5, and 6. However, in practice, some of the data on AMP3 would be fallback protected on each of the other AMPs in the system. The system becomes unavailable if two AMPs in a cluster go down.

Smaller Cluster Configuration

The following figure illustrates smaller clusters. Decreasing cluster size reduces the likelihood that two AMP failures will occur in the same cluster. The illustration shows the same 8-AMP configuration now partitioned into 2 AMP clusters of 4 AMPs each.

	AMP1	AMP2	AMP3	AMP4
Primary copy area	1,9,17	2,10,18	3,11,19	4,12,20
Fallback copy area	2,3,4	1,11,12	9,10,20	17,18,19
Cluster A				
Cluster B				
	AMP5	AMP6	AMP7	AMP8
Primary copy area	5,13,21	6,14,22	7,15,23	8,16,24
Fallback copy area	6,7,8	5,15,16	13,14,24	21,22,23

FG10A002

Compare this clustered configuration with the earlier illustration of an unclustered AMP configuration. In the example, the (primary) data on AMP3 is backed up on AMPs 1, 2, and 4 and the data on AMP6 is backed up on AMPs 5, 7, and 8.

If AMPs 3 and 6 fail at the same time, the system continues to function normally. Only if two failures occur within the same cluster does the system halt.

Subpools are logical groupings of AMPs and disks for fault-tolerance. In a single-clique system to ensure that a disk failure will not bring down both AMPs in a cluster, disks and AMPs are divided into two subpools, and clustering is done across the subpools.

Journaling

Teradata Database supports tables that are devoted to journaling. A journal is a record of some kind of activity. Teradata Database supports several kinds of journaling. The system does some journaling on its own, while you can specify whether to perform other journaling.

The following table explains the capabilities of the different Teradata Database journals.

This type of journal...	Does the following...	And Occurs ...
Down AMP recovery	<ul style="list-style-type: none"> Is active during an AMP failure only Journals fallback tables only Is used to recover the AMP after the AMP is repaired, then is discarded 	always.
Transient	<ul style="list-style-type: none"> Logs BEFORE images for transactions Is used by system to roll back failed transactions aborted either by the user or by the system Captures: <ul style="list-style-type: none"> Begin/End Transaction indicators Before row images for UPDATE and DELETE statements Row IDs for INSERT statements Control records for CREATE, DROP, DELETE, and ALTER statements Keeps each image on the same AMP as the row it describes Discards images when the transaction or rollback completes 	always.
Permanent	<ul style="list-style-type: none"> Is available for tables or databases Can contain before images, which permit rollback, or after images, which permit rollforward, or both before and after images Provides rollforward recovery Provides rollback recovery Provides full recovery of nonfallback tables Reduces need for frequent, full-table archives 	as specified by the user.

Backup Archive and Restore

Teradata Backup Archive and Restore (BAR) solutions use two different architectures:

- The advocated solution is the BAR Framework, which consists of Teradata Database nodes (Windows or Linux) connected to BAR servers using gigabit Ethernet LAN connections. The BAR servers initiate and run the backup, archive, and restore activities.
- The older solution is the direct-attached architecture where tape libraries and drives are directly connected to the Teradata Database nodes, which initiate and run the backup, archive, and restore processes.

The archive and restore product (ARCMAN) supports the following third party BAR backup application software products including:

- NetVault
- NetBackup

- Tivoli Storage Manager

Teradata Archive/Recovery Utility

Teradata Archive/Recovery (ARC) utility archives the following to client tape or client files and restores them to Teradata Database:

- Authorization objects
- Databases
- Data Dictionary tables
- External stored procedures
- Hash Indexes
- Join Indexes
- Methods
- Stored procedures
- Tables (or table partitions)
- Triggers
- UDFs
- UDTs
- Views

If your system is used only for decision support and is updated regularly with data loads, saving the data load media may be an acceptable backup/restore solution.

Table Rebuild Utility

Use the Table Rebuild utility to recreate a table, database, or entire disk on a single AMP under the following conditions:

- The table structure or data is damaged because of a software problem, head crash, power failure, or other malfunction.
- The affected tables are enabled for fallback protection.

Table rebuild can create all of the following on an AMP-by-AMP basis:

- Primary or fallback portions of a table.
- An entire table (both primary and fallback portions).
- All tables in a database.
- All tables on an individual AMP.

The Table Rebuild utility can also remove inconsistencies in stored procedure tables in a database. A Teradata Database system engineer, field engineer, or system support representative usually runs the Table Rebuild utility.

Hardware Fault Tolerance

Teradata Database provides the following facilities for hardware fault tolerance.

Facility	Description
Multiple BYNETs	Multinode Teradata Database servers are equipped with at least two BYNETs. Interprocessor traffic is never stopped unless all BYNETs fail. Within a BYNET, traffic can often be rerouted around failed components.
RAID disk units	<ul style="list-style-type: none"> Teradata Database servers use Redundant Arrays of Independent Disks (RAIDs) configured for use as RAID1, RAID5, or RAIDS. Non-array storage cannot use RAID technology. RAID1 arrays offer mirroring, the method of maintaining identical copies of data. RAID5 or RAIDS protects data from single-disk failures with a 25% increase in disk storage to provide parity. RAID1 provides better performance and data protection than RAID5/RAIDS, but is more expensive.
Multiple-channel and network connections	<p>In a client-server environment, multiple channel connections between mainframe and network-based clients ensure that most processing continues even if one or several connections between the clients and server are not working.</p> <p>Vproc migration is a software feature supporting this hardware issue.</p>
Isolation from client hardware defects	In a client-server environment, a server is isolated from many client hardware defects and can continue processing in spite of such defects.
Battery backup	All cabinets have battery backup in case of building power failures.
Power supplies and fans	Each cabinet in a configuration has redundant power supplies and fans to ensure fail-safe operation.
Hot swap capability for node components	<p>Teradata Database can allow some components to be removed and replaced while the system is running. This process is known as hot swap. Teradata Database offers hot swap capability for the following:</p> <ul style="list-style-type: none"> Disks within RAID arrays Fans Power supplies

Facility	Description
Cliques	<ul style="list-style-type: none">• A clique is a group of nodes sharing access to the same disk arrays. The nodes and disks are interconnected through FC buses and each node can communicate directly to all disks. This architecture provides and balances data availability in the case of a node failure.• A clique supports the migration of vprocs following a node failure. If a node in a clique fails, then its vprocs migrate to another node in the clique and continue to operate while recovery occurs on their home node. Migration minimizes the performance impact on the system.• PEs that manage physical channel connections <i>cannot</i> migrate because they depend on the hardware that is physically attached to the assigned node.• PEs for LAN-attached connections <i>do</i> migrate when a node failure occurs, as do all AMP vprocs.• To ensure maximum fault tolerance, no more than one node in a clique is placed in the same cabinet. Usually the battery backup feature makes this precaution unnecessary, but if you want maximum fault tolerance, then plan your cliques so the nodes are never in the same cabinet.

Teradata Replication Services Using Oracle GoldenGate

Teradata Replication Services Using Oracle GoldenGate preserves and synchronizes business-critical data by capturing row changes made to tables in one database and applying those changes to the same tables in other databases in near-real time. Replication ensures the safety and integrity of data by allowing one system to pass control to another during planned system downtime or an unplanned outage. This technology replicates data from active systems without interruption and limits the impact on system resources because only changes to user-selected tables are copied.

The changes captured and replicated can include inserts, updates, and deletes resulting from SQL DML statements or the execution of utilities, such as MultiLoad and Teradata TPump. The changes can be replicated between Teradata systems located across the street from each other or across the world.

Teradata Replication Services Using Oracle GoldenGate supports both an Active/Active configuration, in which both Teradata systems actively process transactions or an Active/Standby configuration, in which one system actively processes transactions and the other receives duplicate changes and stands by to become the active system if needed.

Hardware and software from both Oracle Corporation and Teradata are required to implement Teradata Replication Services Using Oracle GoldenGate. To purchase and set up Teradata Replication Services Using Oracle GoldenGate, contact your customer representative.

For More Information

For more information on the topics presented in this chapter, see the following Teradata Database and Teradata Tools and Utilities books.

IF you want to learn more about...	See...
Software Fault Tolerance, including: <ul style="list-style-type: none"> • Vproc Migration and Fallback Tables • Clusters (AMP clusters, one-cluster and small cluster configurations) • Journaling and Backup/Archive/Recovery (online archiving) • Table Rebuild Utility 	<ul style="list-style-type: none"> • <i>Database Administration</i> • <i>Teradata Archive/Recovery Utility Reference</i> • <i>SQL Data Definition Language</i> • <i>Utilities</i>
Hardware Fault Tolerance	<i>Database Design</i>
Teradata Replication Services Using Oracle GoldenGate	Teradata Replication Services Using Oracle GoldenGate

CHAPTER 6 **Communication Between the Client and Teradata Database**

This chapter describes various ways the client applications can communicate with Teradata Database.

Attachment Methods

Client applications can connect to Teradata Database using one of the following methods:

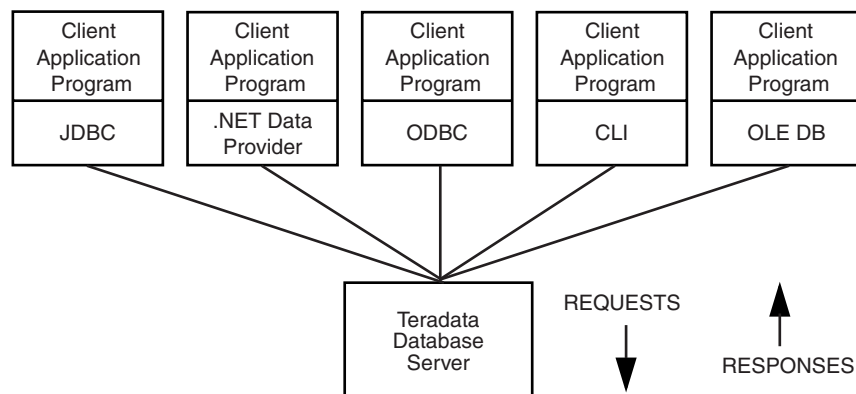
- Network-attached through a Local Area Network (LAN)
- Channel-attached through an IBM mainframe

Network Attachment Methods

Network-attached methods include:

- .NET Data Provider for Teradata
- Java Database Connectivity (JDBC)
- OLE DB Provider for Teradata
- Open Database Connectivity (ODBC)
- Teradata CLIv2 for NAS

The following figure illustrates the transparent connection between client applications and Teradata Database.



1091A006

.NET Data Provider

.NET Data Provider for Teradata uses CLIV2 to connect, execute commands, and retrieve results from Teradata Database. Results can be processed directly or cached in an ADO.NET DataSet. An ADO.NET DataSet can generate XML, and bridges relational and XML data sources.

Java Database Connectivity

JDBC is a specification for an API. The API allows platform-independent Java applications to access Teradata Database using SQL and, for Teradata Database running on 64-bit Windows and Linux, external stored procedures.

The JDBC API provides a standard set of interfaces for:

- Opening connections to databases
- Executing SQL statements
- Processing results

Teradata JDBC Driver provides access to Teradata Database using the Java language. Teradata JDBC Driver is a type 4 (pure Java) JDBC Driver. It is a set of Java classes that use TCP/IP to communicate directly with Teradata Database.

OLE DB Provider for Teradata

OLE DB Provider for Teradata allows programmers to design application programs that allow access between Teradata Database and data stores that do not use SQL. The application program (the consumer) requests database information from an intermediate program (the provider), which accesses Teradata Database. The intermediate program receives the response from Teradata Database and returns a copy of the desired data to the application program.

OLE DB Provider for Teradata also uses service providers. A service provider enhances the functionality of a provider; for example, the Microsoft Cursor Service for OLE DB adds client-side cursor support to any provider.

Open Database Connectivity

ODBC Driver for Teradata provides an interface to Teradata Databases using the industry standard ODBC API. ODBC Driver for Teradata provides Core-level SQL and Extension-level 1 (with some Extension-level 2) function call capability using the Windows Sockets (WinSock) Transmission Control Protocol/Internet Protocol (TCP/IP) communications software interface. ODBC operates independently of CLI.

Teradata CLIV2 for Network-Attached Systems

Teradata CLIV2 for network-attached systems (NAS) is a Teradata proprietary API and library providing an interface between applications on a LAN-connected client and Teradata Database server. Teradata CLIV2 for NAS can:

- Build parcels which are packaged by Micro Teradata Director Program (MTDP) and sent to Teradata Database using the Micro Operating System Interface (MOSI).

- Provide an application with a pointer to each of the parcels returned from Teradata Database.

MTDP


MTDP is the interface between Teradata CLIV2 for NAS and MOSI. Functions of MTDP include:

- Session initiation and termination
- Logging, verification, recovery, and restart
- Physical input to and output from the server

Note: MTDP does not control session balancing; session balancing on network-attached systems is controlled by Teradata Database Gateway on the server.

MOSI

MOSI is the interface between MTDP and Teradata Database. MOSI is a library of service routines providing operating system independence among clients that access Teradata Database. With MOSI, only one version of MTDP is required to run on all network-attached platforms.



2418B004

Channel Attachment Method

Channel attachment uses Teradata CLIV2 for channel-attached systems (CAS).

Teradata CLIV2 for Channel-Attached Systems

Teradata CLIV2 for CAS is a collection of callable service routines providing the interface between applications and the Teradata Director Program (TDP) on an IBM mainframe client. Teradata CLIV2 for CAS can operate with all versions of IBM operating systems, including Customer Information Control System (CICS), Information Management System (IMS), and IBM System z Operating System.

By way of TDP, Teradata CLIV2 for CAS sends requests to the server and provides client applications with responses from the server. Teradata CLIV2 for CAS provides support for:

- Managing multiple serially executed requests in a session
- Managing multiple simultaneous sessions to the same or different servers
- Using cooperative processing so an application can perform operations on the client and the server at the same time
- Communicating with Two-Phase Commit (2PC) coordinators for CICS and IMS transactions
- Generally insulating the application from the details of communicating with a server

Teradata Director Program

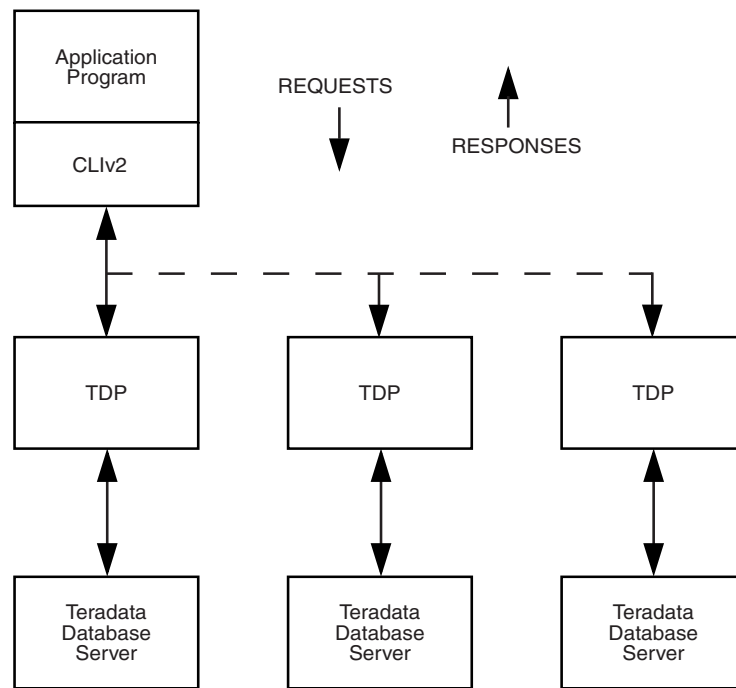
TDP manages communications between Teradata CLIV2 for CAS and the Teradata Database server. TDP executes on the same mainframe as Teradata CLIV2 for CAS, but runs as a different job or virtual machine. Although an individual TDP is associated with one logical server, any number of TDPs may operate and be simultaneously accessed by Teradata CLIV2 for CAS on the same mainframe. Each TDP is referred to by an application using an identifier called the TDPid that is unique in a mainframe; for example, TDP2.

Functions of TDP include:

- Session initiation and termination
- Logging, verification, recovery, and restart
- Physical input to and output from the server, including session balancing and queue maintenance
- Security

Teradata Database Server

A server implements the actual relational database that processes requests received from Teradata CLIV2 for CAS by way of TDP. The following figure illustrates the logical structure of the client-server interface on channel-attached systems.



1091B004

For More Information

For more information on the topics presented in this chapter, see the following Teradata Tools and Utilities books.

IF you want to learn more about...	See...
<p>Network Attachment Methods, including:</p> <ul style="list-style-type: none"> • .NET Data Provider for Teradata • Teradata CLIV2 for NAS • Java Database Connectivity, including Java language external stored procedures that use JDBC • OLE DB Provider for Teradata • Open Database Connectivity 	<ul style="list-style-type: none"> • .NET Data Provider for Teradata Release Definition • <i>Teradata Call-Level Interface Version 2 Reference for Network-Attached Systems</i> • <i>Teradata JDBC Driver User Guide</i> • <i>OLE DB Provider for Teradata User Guide</i> • <i>ODBC Driver for Teradata User Guide</i>
<p>Channel Attachment Method, including:</p> <ul style="list-style-type: none"> • Teradata CLIV2 for CAS • Teradata Director Program 	<ul style="list-style-type: none"> • <i>Teradata Call-Level Interface Version 2 Reference for Mainframe-Attached Systems</i> • <i>Teradata Director Program Reference</i>

SECTION 3 **Using Teradata Database**

CHAPTER 7 Database Objects, Databases, and Users

This chapter provides information about database objects stored in Teradata Database and space allocation for databases and users.

Tables

Tables are two-dimensional objects consisting of rows and columns. Data is organized in table format and presented to the users of a relational database.

Table Types

Table Type	Description
Permanent	Permanent tables allow different sessions and users to share table content.
Queue	Queue tables: <ul style="list-style-type: none">• Are permanent tables with a timestamp column. The timestamp indicates when each row was inserted into the table.• Establish first-in first-out (FIFO) ordering of table contents, which is needed for customer applications requiring event processing.
NoPI	NoPI tables are permanent tables that do not have primary indexes defined on them. They provide a performance advantage when used as staging tables to load data from FastLoad or TPump Array INSERT. They can have secondary indexes defined on them to avoid full-table scans during row access.
Error Logging	Error logging tables: <ul style="list-style-type: none">• Store information about errors on an associated permanent table.• Log information about insert and update errors.
Global Temporary	Global temporary tables: <ul style="list-style-type: none">• Are private to the session.• Are dropped automatically at the end of a session.• Have a persistent table definition stored in the Data Dictionary. The saved definition may be shared by multiple users and sessions with each session getting its own instance of the table.

Table Type	Description
Global Temporary Trace	Global temporary trace tables: <ul style="list-style-type: none">• Store trace output for the length of the session.• Have a persistent table definition stored in the Data Dictionary.• Are useful for debugging SQL stored procedures (via a call to an external stored procedure written to the trace output) and external routines (UDFs, UDMs, and external stored procedures).
Volatile	Volatile tables are used when: <ul style="list-style-type: none">• Only one session needs the table.• Only the creator needs to access the table.• You want better performance than a global temporary table.• You do not need the table definition after the session ends. Note: The definition of a volatile table can survive across a system restart if it is contained in a macro.
Derived	A derived table: <ul style="list-style-type: none">• Is a type of temporary table obtained from one or more other tables as the result of a subquery.• Is specified in an SQL SELECT statement.• Avoids the need to use the CREATE and DROP TABLE statements for storing retrieved information.• Is useful when you are coding more sophisticated, complex queries.

For more information about table types, see *SQL Fundamentals*.

Views

Database views are actually virtual tables that you can use *as if* they were physical tables to retrieve data defining columns from underlying views or tables, or from both.

A view does not contain data and is not materialized until a DML statement references it. View definitions are stored in the Data Dictionary.

Creating Views

A view is created from one or more base tables or from other views.

In fact, you can create hierarchies of views in which views are created from other views. This can be useful, but be aware that deleting any of the lower-level views invalidates dependencies of higher-level views in the hierarchy.

A view usually presents only a subset of the columns and rows in the base table or tables.

Moreover, some view columns do not exist in the underlying base tables. For example, it is possible to present data summaries in a view (for example, an average), which you cannot directly obtain from a base table.

Benefits of Using Views

There are at least four reasons to use views. Views provide:

- A user view of data in the database.
- Security for restricting table access and updates.
- Well-defined, well-tested, high-performance access to data.
- Logical data independence.

Restrictions on Using Views

You can use views as if they were tables in SELECT statements. Views are subject to some restrictions regarding INSERT, UPDATE, MERGE, and DELETE statements. For more information, see [“SQL Access to the Data Dictionary” on page 141](#).

SQL Stored Procedures

SQL stored procedures are executed on Teradata Database server space. It is a combination of procedural control statements, SQL statements, and control declarations that provide a procedural interface to Teradata Database.

Using SQL Stored Procedures

Using SQL stored procedures, you can build large and complex database applications. In addition to a set of SQL control statements and condition handling statements, an SQL stored procedure can contain the following:

- Multiple input and output parameters.
- Local variables and cursors.
- SQL DDL, DCL, and DML statements, including dynamic SQL, with a few exceptions.

Dynamic SQL is a method of invoking an SQL statement by creating and submitting it at runtime from within a stored procedure.

Applications based on SQL stored procedures provide the following benefits. They:

- Reduce network traffic in the client-server environment because stored procedures reside and execute on the server.
- Allow encapsulation and enforcement of business rules on the server, contributing to improved application maintenance.
- Provide better transaction control.
- Provide better security. The data access clause of an SQL stored procedure can restrict access to the database.
- Provide better security by granting the user access to the procedures rather than to the data tables.
- Provide an exception handling mechanism to handle the runtime conditions generated by the application.

- Allow all the SQL and SQL control statements embedded in an SQL stored procedure to be executed by submitting one CALL statement. Nested CALL statements further extend the versatility.

Elements of an SQL Stored Procedure

An SQL stored procedure contains some or all of the following elements.

This element...	Includes...
SQL control statements	nested or non-nested compound statements.
Control declarations	<ul style="list-style-type: none">• Condition handlers in DECLARE HANDLER statements for completion and exception conditions. Conditional handlers can be:<ul style="list-style-type: none">• CONTINUE or EXIT type.• Defined for a specific SQLSTATE code, the generic exception condition SQLEXCEPTION, or generic completion conditions NOT FOUND and SQLWARNING.• Cursor declarations in DECLARE CURSOR statements or in FOR iteration statements. Cursors can be either updatable or read only type. Cursors can also be result set cursors for returning the result of a SELECT statement executed in the stored procedure to the caller or client applications• Local variable declarations in DECLARE statements.
SQL transaction statements	DDL, DCL, DML, and SELECT statements, including dynamic SQL statements, with a few exceptions.
LOCKING modifiers	with all supported SQL statements except CALL.
Comments	bracketed and simple comments. Note: Nested bracketed comments are not allowed.

For more information, see [“SQL Stored Procedures as SQL Applications” on page 102](#).

External Stored Procedures

External stored procedures are written in the C, C++, or Java programming language, installed on the database, and then executed like stored procedures.

Usage

Here is a synopsis of the steps you take to develop, compile, install, and use external stored procedures:

- 1 Write, test, and debug the C, C++, or Java code for the procedure.
- 2 If you are using Java, place the class or classes for the external stored procedure in an archive file (JAR or ZIP) and call the SQLJ.INSTALL_JAR external stored procedure to register the archive file with the database.
- 3 Use CREATE PROCEDURE or REPLACE PROCEDURE for external stored procedures to create a database object for the external stored procedure.
- 4 Use GRANT to grant privileges to users who are authorized to use the external stored procedure.
- 5 Invoke the procedure using the CALL statement.

Macros

The macro database object consists of one or more SQL statements that can be executed by performing a single request. Each time the macro is performed, one or more rows of data may be returned.

SQL Statements Related to Macros

The following table lists the basic SQL statements that you can use with macros.

Use this statement...	To...
CREATE MACRO	incorporate a frequently used SQL statement or series of statements into a macro.
EXECUTE	run a macro. Note: A macro can also contain an EXECUTE statement that executes another macro.
DROP MACRO	delete a macro.

Single-User and Multi-User Macros

You can create a macro for your own use, or grant execution authorization to others. For example, your macro might enable a user in another department to perform operations on the data in Teradata Database. When executing the macro, the user need not be aware of the database access, the tables affected, or even the results.

Macro Processing

Regardless of the number of statements in a macro, Teradata Database treats it as a single request. When you execute a macro, the system processes either all of the SQL statements, or processes none of the statements. If a macro fails, the system aborts it, backs out any updates, and returns the database to its original state.

Triggers

The trigger defines events that happen when some other event, called a triggering event, occurs. This database object is essentially a stored SQL statement associated with a table called a *subject* table.

Teradata Database implementation of triggers complies with ANSI SQL specifications and provides extensions.

Triggers execute when any of the following modifies a specified column or columns in the subject table:

- DELETE
- INSERT
- UPDATE

Typically, the stored SQL statements perform a DELETE, INSERT, UPDATE, or MERGE on a table different from the subject table.

Types of Triggers

Teradata Database supports two types of triggers.

This type of trigger...	Fires for each...
statement	statement that modifies the subject table.
row	row modified in the subject table.

When to Fire Triggers

You can specify when triggers fire.

WHEN you specify...	THEN the triggered action...
BEFORE	<p>executes before the completion of the triggering event.</p> <p>As specified in the ANSI SQL standard, a BEFORE trigger cannot have data changing statements in the triggered action.</p>
AFTER	<p>executes after completion of the triggering event.</p> <p>Note: To support stored procedures the CALL statement is supported in the body of an AFTER trigger. Both row and statement triggers can call a stored procedure.</p>

Sometimes a request fires a trigger, which in turn, fires another trigger. Thus the outcome of one triggering event can itself become another trigger. Teradata Database processes and optimizes the triggered and triggering statements in parallel to maximize system performance.

ANSI-Specified Order

When you specify multiple triggers on a subject table, both BEFORE and AFTER triggers execute in the order in which they were created as determined by the timestamp of each trigger.

Triggers are sorted according to the preceding ANSI rule, unless you use the Teradata Database extension, ORDER. This extension allows you to specify the order in which the triggers execute, regardless of creation time stamp.

Using Triggers

You can use triggers to do various things:

- Define a trigger on the subject table to ensure that the UPDATE, INSERT, MERGE, and DELETE statements performed to the subject table are propagated to another table.
- Use triggers for auditing. For example, you can define a trigger which causes the INSERT statements in a log table when an employee receives a raise higher than 10%.
- Use a trigger to disallow massive UPDATE, INSERT, MERGE, or DELETE during business hours.
- Use a trigger to set a threshold. For example, you can use triggers to set thresholds for inventory of each item by store, to create a purchase order when the inventory drops below a threshold, or to change a price if the daily volume does not meet expectations.
- Use a trigger to call SQL stored procedures and external stored procedures.

User-Defined Functions

SQL provides a set of useful functions, but they might not satisfy all of the particular requirements you have to process your data.

Teradata Database supports two types of user-defined functions (UDFs) that allow you to extend SQL by writing your own functions:

- SQL UDFs
- External UDFs

SQL UDFs

SQL UDFs allow you to encapsulate regular SQL expressions in functions and then use them like standard SQL functions.

Rather than coding commonly used SQL expressions repeatedly in queries, you can objectize the SQL expressions through SQL UDFs.

Moving complex SQL expressions from queries to SQL UDFs makes the queries more readable and can reduce the client/server network traffic.

External UDFs

External UDFs allow you to write your own functions in the C, C++, or Java programming language, install them on the database, and then use them like standard SQL functions.

You can also install external UDF objects or packages from third-party vendors.

Teradata Database supports three types of external UDFs.

UDF Type	Description
Scalar	Scalar functions take input parameters and return a single value result. Examples of standard SQL scalar functions are CHARACTER_LENGTH, POSITION, and TRIM.
Aggregate	Aggregate functions produce summary results. They differ from scalar functions in that they take grouped sets of relational data, make a pass over each group, and return one result for the group. Some examples of standard SQL aggregate functions are AVG, SUM, MAX, and MIN.
Table	A table function is invoked in the FROM clause of a SELECT statement and returns a table to the statement.

Usage

To create and use an SQL UDF, follow these steps:

- 1 Use CREATE FUNCTION or REPLACE FUNCTION to define the UDF.
- 2 Use GRANT to grant privileges to users who are authorized to use the UDF.
- 3 Call the function.

Here is a synopsis of the steps you take to develop, compile, install, and use an external UDF:

- 1 Write, test, and debug the C, C++, or Java code for the UDF.
- 2 If you are using Java, place the class or classes for the UDF in an archive file (JAR or ZIP) and call the SQLJ.INSTALL_JAR external stored procedure to register the archive file with the database.
- 3 Use CREATE FUNCTION or REPLACE FUNCTION to create a database object for the UDF.
- 4 Use GRANT to grant privileges to users who are authorized to use the UDF.
- 5 Call the function.

Related Topics

For more information on ...	See ...
writing, testing, and debugging source code for an external UDF	<i>SQL External Routine Programming</i>
data definition statements related to UDFs, including CREATE FUNCTION and REPLACE FUNCTION	<i>SQL Data Definition Language</i>
invoking a table function in the FROM clause of a SELECT statement	<i>SQL Data Manipulation Language</i>
archiving and restoring UDFs	<i>Teradata Archive/Recovery Utility Reference</i>

User-Defined Methods

A User-Defined Method (UDM) is a special kind of UDF that is associated with a UDT. The term *method* and the acronym UDM are interchangeable.

Teradata Database supports two types of UDMs:

- Instance
- Constructor

Instance Methods

An instance method operates on a specific instance of a distinct or structured UDT. For example, an instance method named *area* might calculate and return the area of an instance of a structured UDT named *circle* that contains attributes *x*, *y*, and *radius*.

Instance methods can also provide transform, ordering, and cast functionality for a distinct or structured UDT. Teradata Database uses this functionality during certain operations involving the UDT.

Constructor Methods

A constructor method initializes an instance of a structured UDT.

A structured UDT can have more than one constructor method, each one providing different initialization options.

User-Defined Types

SQL provides a set of predefined data types, such as INTEGER and VARCHAR, that you can use to store the data that your application uses, but they might not satisfy all of the requirements you have to model your data.

User-defined types (UDTs) allow you to extend SQL by creating your own data types and then using them like predefined data types.

UDT Types

Teradata Database supports distinct and structured UDTs.

UDT Type	Description	Example
Distinct	A UDT that is based on a single predefined data type, such as INTEGER or VARCHAR.	A distinct UDT named euro that is based on a DECIMAL(8,2) data type can store monetary data.
Structured	A UDT that is a collection of one or more fields called attributes, each of which is defined as a predefined data type or other UDT (which allows nesting).	A structured UDT named circle can consist of x-coordinate, y-coordinate, and radius attributes.

Distinct and structured UDTs can define methods that operate on the UDT. For example, a distinct UDT named euro can define a method that converts the value to a US dollar amount. Similarly, a structured UDT named circle can define a method that computes the area of the circle using the radius attribute.

Teradata Database also supports a form of structured UDT called dynamic UDT. Instead of using a CREATE TYPE statement to define the UDT, like you use to define a distinct or structured type, you use the NEW VARIANT_TYPE expression to construct an instance of a dynamic UDT and define the attributes of the UDT at run time.

Unlike distinct and structured UDTs, which can appear almost anywhere that you can specify predefined types, you can only specify a dynamic UDT as the data type of (up to eight) input parameters to external UDFs. The benefit of dynamic UDTs is that they significantly increase the number of input arguments that you can pass in to external UDFs.

Databases and Users

While Teradata Database is a collection of related tables, views, stored procedures, macros, and so on, it also contains databases and users.

A database and a user are *almost* identical in Teradata Database. The major difference is that a user can log on to the system whereas the database cannot.

Creating Databases and Users

When Teradata Database is first installed on a server, one user exists on the system, that is, *User DBC* exists. *User DBC* owns all other databases and users in the system, and initially owns all the space in the entire system.

When Teradata Database is first installed on a server, User DBC is created. User DBC owns (with some exceptions):

- All free space on the system
- All databases and users created after installation

The database administrator manages this user and assigns space from *User DBC* to all other objects (or database and users).

To protect the security of system tables within Teradata Database, the database administrator typically creates *User System Administrator* from *User DBC*. The usual procedure is to assign all database disk space that system tables do not require to *User System Administrator*.

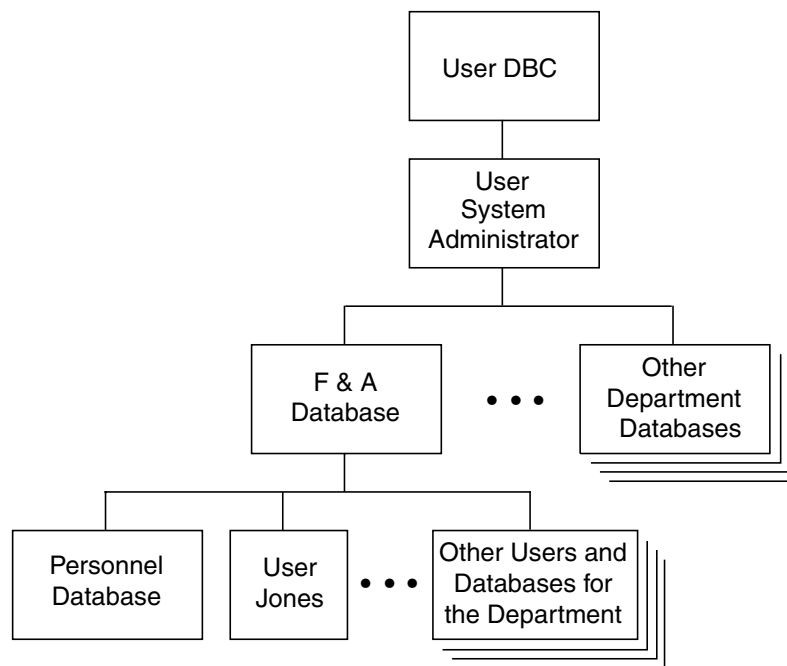
The database administrator then uses this user as a resource from which to allocate space to other databases and users of the system.

For information on how to create databases and users, see *Database Administration*.

Creating a Finance and Administration Database

Consider the following example: the database administrator needs to create a Finance and Administration (F&A) department database with *User Jones* as a supervisory user, that is, as database administrator within the F&A department.

The database administrator first creates *F&A Database* and then allocates space from it to *User Jones* to act as the F&A database administrator. The database administrator also allocates space from F&A to Jones for his personal use and for creating a *Personnel Database*, as well as other databases and other user space allocations. The following figure illustrates the hierarchy implicit in this relationship.



1091A004

F&A Database owns *Personnel Database* and all the other department databases.

F&A Database also owns *User Jones* and all other users within the department.

Because *User DBC* ultimately owns all databases and users, it is the final owner of all the databases and user space belonging to the organization.

This hierarchical ownership structure provides an owner of a database or user space with complete control over the security of owned data. An owner can archive the database or control access to it by granting or revoking privileges on it.

Note: There can only one immediate owner of a database or user.

For More Information

For more information on the topics presented in this chapter, see the following Teradata Database books.

If you want to learn more about...	See...
Tables	<ul style="list-style-type: none">• <i>SQL Fundamentals</i>• <i>SQL Data Definition Language</i>• <i>SQL Data Manipulation Language</i>
Queue tables	<ul style="list-style-type: none">• <i>SQL Fundamentals</i>• <i>SQL Data Definition Language</i>• <i>SQL Data Manipulation Language</i>
Views	<i>SQL Fundamentals</i>
SQL stored procedures	<ul style="list-style-type: none">• <i>SQL Data Definition Language</i>• <i>SQL Fundamentals</i>• <i>SQL Stored Procedures and Embedded SQL</i>
SQL external stored procedures	<ul style="list-style-type: none">• <i>SQL Data Definition Language</i>• <i>SQL Fundamentals</i>• <i>SQL Stored Procedures and Embedded SQL</i>
Macros	<i>SQL Fundamentals</i>
Triggers	<i>SQL Fundamentals</i>
User-Defined Functions	<ul style="list-style-type: none">• <i>SQL Data Definition Language</i>• <i>SQL Fundamentals</i>• <i>SQL External Routine Programming</i>
User-Defined Methods	<ul style="list-style-type: none">• <i>SQL Data Definition Language</i>• <i>SQL External Routine Programming</i>

If you want to learn more about...	See...
User-Defined Types	<ul style="list-style-type: none">• <i>SQL Data Definition Language</i>• <i>SQL Fundamentals</i>• <i>SQL External Routine Programming</i>
Databases and Users	<i>Database Administration</i>

This chapter describes SQL, the ANSI standard language for relational database management.

All application programming facilities ultimately make queries against Teradata Database using SQL because it is the only language Teradata Database understands.

To enhance the capabilities of SQL, Teradata Database has added extensions that are unique to Teradata Database. This comprehensive language is referred to as Teradata SQL.

Using SQL

SQL has the advantage of being the most commonly used language for relational database management systems. Because of this, both the data structures in Teradata Database and the commands for manipulating those structures are controlled using SQL. In addition, all applications, including those written in a client language with embedded SQL, macros, and ad hoc SQL queries, are written and executed using the same set of instructions and syntax.

Other database management systems use different languages for data definition and data manipulation and may not permit ad-hoc queries of the database. Teradata Database lets you use one language to define, query, and update your data.

Types of SQL Statements

The SQL language allows you, using SQL statements, to define database objects, to define user access to those objects, and to manipulate the data stored.

These functions form the principal functional families of SQL statements:

- Data Definition Language (DDL) statements
- Data Control Language (DCL) statements
- Data Manipulation Language (DML) statements

In addition, SQL provides HELP and SHOW statements that provide help about database object definitions, sessions and statistics, the EXPLAIN request modifier, SQL statement syntax, as well as displaying the SQL used to create tables.

The following sections contain information about the functional families of Teradata SQL.

Data Definition Language Statements

You use DDL statements to define the structure and instances of a database. DDL provides statements for the definition and description of database objects.

The following table lists some *basic* DDL statements. The list is not exhaustive.

Statement	Action
CREATE	Defines a new database object, such as a database, user, table, view, trigger, index, macro, stored procedure, user-defined type, user-defined function, or user-defined macro, depending on the object of the CREATE request.
DROP	Removes a database object, such as a database, user, table, view, trigger, index, macro, stored procedure, user-defined type, user-defined function, user-defined method, depending on the object of the DROP request.
ALTER	Changes, for example, a table, column, referential constraint, trigger, or index.
ALTER PROCEDURE	Recompiles an external stored procedure.
MODIFY	Changes a database or user definition.
RENAME	Changes, for example, the names of tables, triggers, views, stored procedures, and macros.
REPLACE	Replaces, for example, macros, triggers, stored procedures, and views
SET	Specifies, for example, time zones, the collation or character set for a session.
COLLECT	Collects optimizer or QCD statistics on, for example, a column, group of columns, index.
DATABASE	Specifies a default database.
COMMENT	Inserts or retrieves a text comment for a database object.

Successful execution of a DDL statement automatically creates, updates, or removes entries in the Data Dictionary. For information about the contents of the Data Dictionary, see [Chapter 12: “The Data Dictionary.”](#)

Data Control Language Statements

You use DCL statements to grant and revoke access to database objects and change ownership of those objects from one user or database to another. The results of DCL statement processing also are recorded in the Data Dictionary.

The following table lists some *basic* DCL statements. The list is not exhaustive.

Statement	Action
GRANT/REVOKE	Controls privileges of the users on an object.
GRANT LOGON/ REVOKE LOGON	Controls logon privileges to a host (client) or host group (if the special security user is enabled).

Statement	Action
GIVE	Gives a database object to another database object.

Data Manipulation Language Statements

You use DML statements to manipulate and process database values. You can insert new rows into a table, update one or more values in stored rows, or delete a row.

The following table list some *basic* DML statements. The list is not exhaustive.

Statement	Action
CHECKPOINT	Checkpoints a journal. CHECKPOINT is a statement that defines a recovery point in the journal that can later be used to restore the table contents to its state at a point in time. This can be useful if, for example, the table contents become incorrect due to hardware failure or an operational error.
DELETE	Removes a row (or rows) from a table.
ECHO	Echoes a string or command to a client.
INSERT	Inserts new rows into a table. For more information about a special case of INSERT, see Atomic Upsert later in this table.
MERGE	Combines both UPDATE and INSERT in a single SQL statement. Supports primary index operations only, similar to Atomic Upsert but with fewer constraints.
These statements: <ul style="list-style-type: none"> • ABORT • ROLLBACK • COMMIT • BEGIN TRANSACTION • END TRANSACTION 	Allows you to manage transactions.
SELECT	Returns specific row data in the form of a result table.

Statement	Action
UPDATE	Modifies data in one or more rows of a table. For more information about a special case of UPDATE, see Atomic Upsert later in this table.
	Atomic Upsert The upsert form of the UPDATE DML statement is a Teradata Database extension of the ANSI SQL standard designed to enhance the performance of TPump utility by allowing the statement to support atomic upsert. For more information about how TPump operates, see “Teradata Parallel Data Pump” on page 180 . This feature allows Teradata TPump and all other CLIV2-, ODBC-, and JDBC-based applications to perform single-row upsert operations using an optimally efficient single-pass strategy. This single-pass upsert is called atomic to emphasize that its component UPDATE and INSERT SQL statements are grouped together and performed as a single, or atomic, SQL statement.

SQL Statement Syntax

A typical SQL request consists of the following:

- A statement keyword
- One or more column names
- A database name
- A table name
- One or more optional clauses introduced by keywords

For example, in the following single-statement request, the statement keyword is SELECT:

```
SELECT deptno, name, salary
FROM personnel.employee
WHERE deptno IN(100, 500)
ORDER BY deptno, name
;
```

The select list and FROM clause for this statement is made up of the following names:

- Deptno, name, and salary (the column names)
- Personnel (the database name)
- Employee (the table name)

The search condition, or WHERE clause, is introduced by the keyword WHERE, as in:

```
WHERE deptno IN(100, 500)
```

The sort ordering, or ORDER BY clause, is introduced by the keywords ORDER BY, as in:

```
ORDER BY deptno, name
```

Statement Execution

Teradata Database offers the following ways to invoke an executable statement:

- Interactively from a terminal
- Embedded within an application program
- Dynamically created within an embedded application
- Embedded within a stored procedure or external stored procedure
- Dynamically created within an SQL stored procedure
- Via a trigger
- Embedded within a macro

Statement Punctuation

You can use punctuation to separate or identify the parts of an SQL statement.

This syntax element...	Named...	Performs this function in a SQL statement...
.	period	separates database names from table names and table names from a particular column name (for example, personnel.employee.deptno).
,	comma	separates and distinguishes column names in the select list, or column names or parameters in an optional clause.
'	apostrophe	delimits the boundaries of character string constants.
()	left and right parentheses	groups expressions or defines the limits of a phrase.
;	semicolon	separates statements in multi-statement requests and terminates requests submitted via certain utilities such as BTEQ.
"	double quotation marks	identifies user names that might otherwise conflict with SQL reserved words or that would not be valid names in the absence of the double quotation marks.
:	colon	prefixes reference parameters or client system variables.

To include an apostrophe or show possession in a title, double the apostrophes.

The SELECT Statement

SELECT is probably the most frequently used SQL statement. It specifies the table columns from which to obtain the data you want, the corresponding database (if different from the current default database), and the table or tables that you need to reference within that database.

The SELECT statement further specifies how, in what format, and in what order the system returns the set of result data.

You can use the following options, lists, and clauses with the SELECT statement to request data from Teradata Database. The list is not exhaustive.

- DISTINCT option
- FROM clause
- WHERE clause, including subqueries
- GROUP BY clause
- HAVING clause
- QUALIFY clause
- ORDER BY clause
 - CASESPECIFIC option
 - International sort orders
- WITH clause
- Query expressions and set operators

Another variation is the SELECT INTO statement, which is used in embedded SQL and stored procedures. This statement selects at most one row from a table and assigns the values in that row to host variables in embedded SQL or to local variables or parameters in Teradata Database stored procedures.

SELECT Statement and Set Operators

The SELECT statement can use the set operators UNION, INTERSECT, and MINUS/EXCEPT. These set operators allow you to manipulate the answers to two or more queries by combining the results of each query into a single result set.

You can use the set operators within, for example, the following operations:

- View definitions
- Derived tables
- Subqueries

SELECT Statement and Joins

A SELECT statement can reference data in two or more tables and the relational join combines the data from the referenced tables.

In this way, the SELECT statement defines a join of specified tables to retrieve data more efficiently than without defining a join of tables.

You can specify both inner joins and outer joins:

- An inner join selects data from two or more tables that meet specific join conditions. Each source must be named and the join condition, that is the common relationship among the tables to be joined, can be on an ON clause or a WHERE clause.
- The outer join is an extension of the inner join that includes rows that qualify for a simple inner join, as well as a specified set of rows that do not match the join conditions expressed by the query.

SQL Data Types

You must specify a data type for each column when you use SQL to create a table because Teradata Database does not provide a default data type. You can include a data type to specify data conversions in expressions.

Data Types

The following table describes the SQL data types Teradata Database supports.

Teradata Database supports ...	Including, for example...
Teradata Database data types	<ul style="list-style-type: none">• Byte• Graphic• Period• Geospatial <p>For more information on geospatial types, see <i>SQL Geospatial Types</i>.</p>
ANSI-compliant data types	<ul style="list-style-type: none">• Large Objects (LOBs):<ul style="list-style-type: none">• Binary Large Objects (BLOBs)• Character Large Objects (CLOBs)• Character• DateTime• Interval• Numeric
User-defined Types (UDTs)	<ul style="list-style-type: none">• Distinct• Structured

Data Type Phrase

A data type phrase does the following:

- Determines how data is stored on Teradata Database.

- Specifies how data is presented to the user.

Data Type Attributes

You can use Teradata SQL to define the attributes of a data value. Data type attributes control, among other things:

- Import format (internal representation of stored data).
- Export format (how data is presented for a column or an expression result).

You must define data type attributes when you define a column.

You can override the default values of data type attributes. For example, when you create a table, you can use a FORMAT phrase to override the output format of a data type.

The following table summarizes data type attributes.

Data Type Attribute	ANSI	Teradata Database Extension to ANSI
NOT NULL	X	
UPPERCASE		X
[NOT] CASESPECIFIC		X
FORMAT <i>string_literal</i>		X
TITLE <i>string_literal</i>		X
NAMED <i>name</i>		X
DEFAULT <i>value</i>	X	
DEFAULT USER	X	
DEFAULT DATE		X
DEFAULT TIME		X
DEFAULT NULL	X	
WITH DEFAULT		X
CHARACTER SET	X	

Teradata Database Recursive Query

A recursive query is a named query expression that references itself in its definition. The self-referencing capability gives the user a simple way to search a table using iterative self-join and set operations.

The recursive query feature benefits the user by reducing the complexity of the queries and allowing a certain class of queries to execute more efficiently.

Recursive queries are implemented using the WITH RECURSIVE clause in the statement and the RECURSIVE clause in the CREATE VIEW statement.

SQL Functions

The control statements of SQL stored procedures make SQL a computationally complete (that is, procedural) language. You can write your own UDFs and external stored procedures in C, C++, or Java to define what you want.

Procedural languages contain functions that perform complex operations. The usual SQL statements do not support many functions. However, to reduce the reliance on ancillary application code, SQL does support the following standard functions:

- Scalar
- Aggregate
- Ordered analytical

In addition, you can create scalar, aggregate, and table functions to meet specific needs.

Scalar Functions

A scalar function works on input parameters to create a result.

When it is part of an expression, the function is invoked as needed whenever expressions are evaluated for an SQL statement. When a function completes, its result is used by the expression in which the function was referenced.

For example, the following request returns the current date plus 13 years.

```
SELECT ADD_MONTHS (CURRENT_DATE, 12*13);
```

The following request returns the date 6 months ago.

```
SELECT ADD_MONTHS (CURRENT_DATE, -6);
```

Aggregate Functions

Sometimes the information you want can only be derived from data in a set of rows, instead of individual rows.

Aggregate functions produce results from sets of relational data that you have grouped (optionally) using a GROUP BY or ORDER BY clause. Aggregate functions process each set and produce one result for each set.

The following table lists a few examples of aggregate functions.

The function...	Returns the...
AVG	arithmetic average of the values in a specified column.

The function...	Returns the...
COUNT	number of qualified rows.
MAX	maximum column value for the specified column.
MIN	minimum column value for the specified column.
SUM	arithmetic sum of a specified column.

Ordered Analytical Functions

Ordered analytical functions work over a range of data for a particular set of rows in some specific order to produce a result for each row in the set.

Like aggregate functions, ordered analytical functions are called for each item in a set. But unlike an aggregate function, an ordered analytical function produces a result for each detail item.

Ordered analytical functions allow you to perform sophisticated data mining on the information in your databases to get the answers to questions that SQL otherwise cannot provide.

The following table lists two examples of ordered analytical functions.

The following function...	Returns the...
AVG	arithmetic average of all values in the specified expression for each row in the group. The OVER() phrase must be specified to make AVG an ordered analytical function.
RANK	ordered ranking of rows based on the value of the column being ranked.

Cursors

Traditional application development languages cannot deal with result tables without some kind of intermediary mechanism because SQL is a set-oriented language. The intermediary mechanism is the cursor.

A cursor is a pointer that the application program uses to move through a result table.

You declare a cursor for a SELECT request, and then open the named cursor. The act of opening the cursor executes the SQL request.

You use the FETCH... INTO... statement to individually fetch and write the rows into host variables. The application can then use the host variables to do computations.

Teradata Preprocessor2 uses cursors to mark or tag the first row accessed by an SQL query. Preprocessor2 then increments the cursor as needed.

SQL stored procedures use:

- Cursors to fetch one result row at a time and then to execute SQL and SQL control statements as required for each row. Local variables or parameters from the stored procedure can be used for computations.
- Result set cursors to return the result of a SELECT statement executed in the stored procedure to the caller of the stored procedure or the client application.

Session Modes

Teradata Database supports two session modes:

- ANSI
- Teradata

The semantics of ANSI session mode conform to those of the ANSI SQL:2008 standard. The semantics of Teradata session mode, though largely conforming with ANSI semantics, have several defaults that differ in important ways from the ANSI semantics.

For More Information

For more information on the topics presented in this chapter, see the following Teradata Database books.

If you want to learn more about...	See...
SQL	<i>SQL Fundamentals</i>
Types of SQL Statements	<ul style="list-style-type: none">• <i>SQL Data Control Language</i>• <i>SQL Data Definition Language</i>• <i>SQL Data Manipulation Language</i>
The SELECT Statement	<ul style="list-style-type: none">• <i>SQL Data Manipulation Language</i>• <i>SQL Functions, Operators, Expressions, and Predicates</i>• <i>SQL Fundamentals</i>
SQL Data Types	<i>SQL Data Types and Literals</i>
Teradata Database Recursive Query	<i>SQL Fundamentals</i>
SQL Functions	<i>SQL Functions, Operators, Expressions, and Predicates</i>

If you want to learn more about...	See...
Cursors	<ul style="list-style-type: none">• <i>SQL Data Definition Language</i>• <i>SQL Stored Procedures and Embedded SQL</i>
Session Modes	<ul style="list-style-type: none">• <i>SQL Fundamentals</i>• <i>SQL Request and Transaction Processing</i>

CHAPTER 9 SQL Application Development

This chapter describes the tools used to develop applications for Teradata Database and the interfaces used to establish communications between the applications and Teradata Database.

SQL Applications

Client Applications

Client applications can use the following APIs to communicate with Teradata Database:

- .NET Data Provider
- Java Database Connectivity (JDBC)
- Open database Connectivity (ODBC)

For information on these APIs, see [“Network Attachment Methods” on page 65](#).

Embedded SQL Applications

When you write applications using embedded SQL, you insert SQL requests into your application program, which must be written in one of the supported programming languages shown in [“Supported Languages and Platforms” on page 100](#).

Because third-generation application development languages do not have facilities for dealing with results sets, embedded SQL contains extensions to executable SQL that permit declarations.

Embedded SQL declarations include:

- Code to encapsulate the SQL from the application language
- Cursor definition and manipulation

A cursor is a pointer device you use to read through a results table one record/row at a time. For more information about cursors, see [“Cursors” on page 96](#).

Using Embedded SQL

The client application languages that support embedded SQL are all compiled languages. SQL is not defined for any of them. For this reason, you must precompile your embedded SQL code to translate the SQL into native code before you can compile the source using a native compiler.

The precompiler tool is called Preprocessor2 (PP2), and you use it to:

- Read your application source code to look for the defined SQL code fragments.

- Interpret the intent of the code after it isolates all the SQL code in the application and translates it into Call-Level Interface (CLI) calls.
- Comment out all the SQL source.

The output of the precompiler is native language source code with CLI calls substituted for the SQL source. After the precompiler generates the output, you can process the converted source code with the native language compiler. For information about Call-Level Interface communications interface, see [Chapter 6: “Communication Between the Client and Teradata Database.”](#)

Supported Languages and Platforms

Preprocessor2 supports the following application development languages on the specified platforms.

Application Development Language	Platform
C	<ul style="list-style-type: none">• IBM mainframe clients• UNIX® operating system clients and some other workstation clients
COBOL	<ul style="list-style-type: none">• IBM mainframe clients• Some workstation clients
PL/I	IBM mainframes

Macros as SQL Applications

Teradata Database macros are SQL statements that the server stores and executes. Macros provide an easy way to execute frequently used SQL operations. Macros are particularly useful for enforcing data integrity rules, providing data security, and improving performance.

SQL Used to Create a Macro

You use the CREATE MACRO statement to create Teradata Database macros.

For example, suppose you want to define a macro for adding new employees to the Employee table and incrementing the EmpCount field in the Department table. The CREATE MACRO statement looks like this:

```
CREATE MACRO NewEmp (name VARCHAR(12),
                    number INTEGER NOT NULL,
                    dept INTEGER DEFAULT 100
                    )
AS (INSERT INTO Employee (Name,
                        EmpNo,
                        DeptNo
                        )
VALUES (name,
        number,
```

```

        dept
    )
    ;
    UPDATE Department
    SET EmpCount=EmpCount+1
    WHERE DeptNo=dept
    ;
)
;

```

This macro defines parameters that users must fill in each time they execute the macro. A leading colon (:) indicates a reference to a parameter within the macro.

Macro Usage

The following example shows how to use the NewEmp macro to insert data into the Employee and Department tables.

The information to be inserted is the name, employee number, and department number for employee H. Goldsmith. The EXECUTE macro statement looks like this:

```
EXECUTE NewEmp ('Goldsmith H', 10015, 600);
```

SQL Used to Modify a Macro

The following example shows how to modify a macro. Suppose you want to change the NewEmp macro so that the default department number is 300 instead of 100. The REPLACE MACRO statement looks like this:

```

REPLACE MACRO NewEmp (name VARCHAR(12),
                      number INTEGER NOT NULL,
                      dept INTEGER DEFAULT 300)
AS (INSERT INTO Employee (Name,
                          EmpNo,
                          DeptNo)
VALUES (name,
        number,
        dept)
;
    UPDATE Department
    SET EmpCount=EmpCount+1
    WHERE DeptNo=dept
    ;
)
;

```

SQL Used to Delete a Macro

The example that follows shows how to delete a macro. Suppose you want to drop the NewEmp macro from the database. The DROP MACRO statement looks like this:

```
DROP MACRO NewEmp;
```

SQL Stored Procedures as SQL Applications

SQL stored procedures are database applications created by combining SQL control statements with other SQL elements and condition handlers. They provide a procedural interface to Teradata Database and many of the same benefits as embedded SQL.

SQL stored procedures conform to ANSI SQL standard with some exceptions.

SQL Used to Create Stored Procedures

Teradata SQL supports creating, modifying, dropping, renaming, and controlling privileges of stored procedures through DDL and DCL statements.

You can create or replace an external stored procedure through the COMPILE command in Basic Teradata Query (BTEQ), BTEQ for Microsoft Windows systems (BTEQWIN), and SQL Assistant. You must specify a source file as input for the COMPILE command.

Stored procedures do not need to be compiled, but external stored procedures do.

You can also create or modify a stored procedures using the CREATE PROCEDURE or REPLACE PROCEDURE statement from CLIV2, ODBC, and JDBC applications.

SQL Stored Procedure Example

Assume you want to create an SQL stored procedure named NewProc that you can use to add new employees to the Employee table and retrieve the department name of the department to which the employee belongs. You can also report an error, in case the row that you are trying to insert already exists, and handle that error condition.

The following SQL stored procedure definition includes nested, labeled compound statements. The compound statement labeled L3 is nested within the outer compound statement L1. Note that the compound statement labeled L2 is the handler action clause of the condition handler.

This SQL stored procedure defines parameters that must be filled in each time it is called (executed).

```
CREATE PROCEDURE NewProc (IN name CHAR(12),
                          IN num INTEGER,
                          IN dept INTEGER,
                          OUT dname CHAR(10),
                          INOUT p1 VARCHAR(30))
L1: BEGIN
  DECLARE CONTINUE HANDLER FOR SQLSTATE value '23505'
  L2: BEGIN
    SET p1='Duplicate Row'
    ;
  END L2;
  L3: BEGIN
    INSERT INTO Employee (EmpName, EmpNo, DeptNo)
    VALUES (name, num, dept)
    ;

    SELECT DeptName
```

```

        INTO dname FROM Department
        WHERE DeptNo = dept;
        IF SQLCODE <> 0 THEN LEAVE L3;
        ...
        END L3
    ;
END L1
;

```

SQL Used to Execute a Stored Procedure

After compiling an external stored procedure, procedures are stored as objects in Teradata Database. You can execute stored procedures from Teradata Database client products using the SQL CALL statement. Arguments for all input (IN or INOUT) parameters of the stored procedure must be submitted with the CALL statement.

BTEQ and other Teradata Database client products support stored procedure execution and DDL. These include:

- JDBC
- ODBC
- CLIV2
- PP2
- Teradata SQL Assistant
- BTEQWIN (BTEQ for Windows)

DDL Statements with Stored Procedures

You can use the following DDL statements with stored procedures. The list is not exhaustive.

Use This Statement...	To...
CREATE PROCEDURE	direct the stored procedure compiler to create a procedure from the SQL statements in the remainder of the statement text.
ALTER PROCEDURE	direct the stored procedure compiler to recompile a stored procedure created in an earlier version of Teradata Database without executing SHOW PROCEDURE and REPLACE PROCEDURE statements.
DROP PROCEDURE	drop a stored procedure.
RENAME PROCEDURE	rename a procedure.
REPLACE PROCEDURE	direct the stored procedure compiler to replace the definition of an existing stored procedure. If the specified stored procedure does not exist, create a new procedure by that name from the SQL statements in the remainder of the source text.
HELP PROCEDURE ... ATTRIBUTES	view all the parameters and parameter attributes of a procedure, or the creation time attributes of a procedure.

Use This Statement...	To...
HELP 'SPL'	display a list of all DDL and control statements associated with stored procedures.
HELP 'SPL' <i>command_name</i>	display help about the command you have named.
SHOW PROCEDURE	view the current definition (source text) of a procedure. The text is returned in the same format as defined by the creator.

The EXPLAIN Request Modifier

Teradata SQL supplies a very powerful EXPLAIN request modifier that allows you to see the execution plan of a query.

The EXPLAIN request modifier not only explains how a request will be processed, but provides an estimate of the number of rows involved as well as the performance impact of the request. Teradata Database supports EXPLAIN request modifiers with detailed Optimizer information including, for example, cost estimates for Insert, Update, Upsert, Merge, and Delete steps, as well as spool size estimates.

How EXPLAIN Works

The EXPLAIN request modifier that precedes any SQL request causes Teradata Database to display the execution plan for that request. The request itself is not submitted for execution.

When you perform an EXPLAIN against any SQL request, that request is parsed and optimized. The access and join plans generated by the Optimizer are returned in the form of a text file that explains the (possibly parallel) steps used in the execution of the request. Also included is the relative cost required to complete the request given the statistics with which the Optimizer had to work. If the statistics are not reasonably accurate, the cost estimate may not be accurate.

Benefits of Using EXPLAIN

EXPLAIN helps you to evaluate complex queries and to develop alternative, more efficient, processing strategies. You may be able to get a better plan by collecting more statistics on more columns, or by defining additional indexes. Your knowledge of the actual demographics information may allow you to identify row count estimates that seem badly wrong, and help to pinpoint areas where additional statistics would be helpful.

Simple EXPLAIN Example

The EXPLAIN example shown below results from joining tables with the following table definitions:

```
CREATE TABLE customer  
(c_custkey INTEGER,
```



```

c_name CHAR(26),
c_address VARCHAR(41),
c_nationkey INTEGER,
c_phone CHAR(16),
c_acctbal DECIMAL(13,2),
c_mktsegment CHAR(21),
c_comment VARCHAR(127))
UNIQUE PRIMARY INDEX( c_custkey );

*** Table has been created.
*** Total elapsed time was 1 second.

+-----+-----+-----+-----+-----+-----+-----+
CREATE TABLE orders
(o_orderkey INTEGER NOT NULL,
o_custkey INTEGER,
o_orderstatus CHAR(1),
o_totalprice DECIMAL(13,2) NOT NULL,
o_orderdate DATE FORMAT 'yyyy-mm-dd' NOT NULL,
o_orderpriority CHAR(21),
o_clerk CHAR(16),
o_shippriority INTEGER,
o_comment VARCHAR(79))
UNIQUE PRIMARY INDEX(o_orderkey);

*** Table has been created.
*** Total elapsed time was 1 second.

+-----+-----+-----+-----+-----+-----+-----+
CREATE TABLE lineitem
(l_orderkey INTEGER NOT NULL,
l_partkey INTEGER NOT NULL,
l_suppkey INTEGER,
l_linenummer INTEGER,
l_quantity INTEGER NOT NULL,
l_extendedprice DECIMAL(13,2) NOT NULL,
l_discount DECIMAL(13,2),
l_tax DECIMAL(13,2),
l_returnflag CHAR(1),
l_linestatus CHAR(1),
l_shipdate DATE FORMAT 'yyyy-mm-dd',
l_commitdate DATE FORMAT 'yyyy-mm-dd',
l_receiptdate DATE FORMAT 'yyyy-mm-dd',
l_shipinstruct VARCHAR(25),
l_shipmode VARCHAR(10),
l_comment VARCHAR(44))
PRIMARY INDEX( l_orderkey );

*** Table has been created.
*** Total elapsed time was 1 second.

+-----+-----+-----+-----+-----+-----+-----+
collect stats orders index (o_orderkey) values
(0,0,1,10,1,1000000,1000000)
;

```

```
*** Update completed. One row changed.
*** Total elapsed time was 1 second.
```

```
+-----+-----+-----+-----+-----+-----+-----+
collect stats lineitem index (l_orderkey) values
(0,0,1,10,1,500000,1000000
);
```

```
*** Update completed. One row changed.
*** Total elapsed time was 1 second.
```

The following statement defines a join index on these tables.

```
CREATE JOIN INDEX order_join_line AS
SELECT ( l_orderkey, o_orderdate, o_custkey, o_totalprice ),
( l_partkey, l_quantity, l_extendedprice, l_shipdate )
FROM lineitem
LEFT JOIN orders ON l_orderkey = o_orderkey
ORDER BY o_orderdate
PRIMARY INDEX (l_orderkey);
```

```
*** Index has been created.
*** Total elapsed time was 1 second
```

The following EXPLAIN shows that the Optimizer used the newly created join index, `order_join_line`.

```
EXPLAIN
SELECT o_orderdate, o_custkey, l_partkey, l_quantity,
l_extendedprice
FROM lineitem , orders
WHERE l_orderkey = o_orderkey;
```

```
*** Help information returned. 14 rows.
*** Total elapsed time was 1 second.
```

Explanation

- ```

1) First, we lock a distinct EXPLAINSAMPLE."pseudo table" for read on
a RowHash to prevent global deadlock for
EXPLAINSAMPLE.ORDER_JOIN_LINE.
2) Next, we lock EXPLAINSAMPLE.ORDER_JOIN_LINE for read.
3) We do an all-AMPs RETRIEVE step from EXPLAINSAMPLE.ORDER_JOIN_LINE
by way of an all-rows scan with a condition of ("NOT
(EXPLAINSAMPLE.ORDER_JOIN_LINE.o_orderdate IS NULL)") into Spool 1
(group_amps), which is built locally on the AMPs. The size of
Spool 1 is estimated with high confidence to be 36 rows. The
estimated time for this step is 0.01 seconds.
4) Finally, we send out an END TRANSACTION step to all AMPs involved
in processing the request.
-> The contents of Spool 1 are sent back to the user as the result of
statement 1. The total estimated time is 0.01 seconds.
```

The following statement drops the join index named `order_join_line`.

```
drop join index order_join_line;
```

```
*** Index has been dropped.
*** Total elapsed time was 1 second.
```

# Third-Party Development

Teradata Database supports many third-party software products. The two general components of supported products include those of the transparency series and the native interface products.

## TS/API

The Transparency Series/Application Program Interface (TS/API) provides a gateway between the IBM mainframe relational database product DB2 (IBM System z Operating System version) and Teradata Database.

TS/API enables an SQL statement formulated for DB2 to be translated into Teradata SQL, permitting DB2 applications to access data stored in Teradata Database.

## Compatible Third-Party Software Products

Many third-party, interactive query products operate in conjunction with Teradata Database, permitting queries formulated in a native query language to access Teradata Database.

The list of supported third-party products changes frequently. For a current list, contact your Teradata sales office.

## Workload Management Application Programming Interface

Workload Management API consists of interfaces to PM/APIs and open APIs. These interfaces are used to:

- Monitor system and session-level activities.
- Manage Teradata Active System Management (ASM) rules.
- Update components stored in a custom database called TDWM.
- Track system usage and manage task priorities.

For more information about the APIs, see *Workload Management API: PM/API and Open API*.

### PM/API

PM/APIs provide access to PMPC routines resident in Teradata Database. The PMPC subsystem is available through a logon partition called MONITOR, using a specialized PM/API subset of CLIV2.

PM/APIs have the following features:

- CLIV2 data is acquired in near real time, with less overhead and minimal possibility of being blocked. These capabilities allow frequent in-process performance analysis.
- A CLIV2 request saves the raw data in an in-memory buffer where a client application program can easily retrieve the data for real-time analysis or importing into custom reports.

- A CLIV2 request provides access to data that the resource usage does not. For example, session-level resource usage data, and data on application locks and which application is being blocked.

Using PM/APIs may not be the right choice for all performance monitoring requirements. Standard performance monitoring tools and reports (such as resource usage macros and tables) may be sufficient.

For details, see *Workload Management API: PM/API and Open API*.

## Open API

The open API provides an SQL interface to System PMPC through user-defined functions (UDFs) and external stored procedures.

The following are examples of some of the UDFs and external stored procedures used:

| Use the ...                 | To ...                                                                                                                                       |
|-----------------------------|----------------------------------------------------------------------------------------------------------------------------------------------|
| AbortSessions function      | abort queries submitted by a set of users that have been running longer than 10 minutes and have been skewed by more than 10% for 5 minutes. |
| TDWMRuleControl function    | temporarily enable a rule to block an application from accessing a database while it is synchronized between two active systems.             |
| GetQueryBandValue procedure | query the DBQLogTbl based on specified names and values in the QueryBand field.                                                              |

Most of the SQL interfaces available to System PMPC provide similar functionality to the CLIV2 interfaces.

**Note:** Most open APIs do not follow transaction rules. If, within a transaction, a UDF or external stored procedure is called that performs an action (for example, setting a session account string) and the transaction rolls back, the action of the UDF or external stored procedure is not rolled back.

However, those interfaces that update the Teradata Dynamic Workload Management database, such as the TDWMRuleControl, TDWMObjectAssn, and TDWMSetLimits procedures, must follow transaction rules. If one of these interfaces is called within a transaction, the update will be rolled back if the transaction is aborted.

For more information on the SQL interfaces described in this section and the differences between the open API and PM/API, see *Workload Management API: PM/API and Open API*.

## For More Information

For more information on the topics presented in this chapter, see the following Teradata Database and Teradata Tools and Utilities books.

| IF you want to learn more about...                                                                                                        | See...                                                                                                                                                                                               |
|-------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| SQL Applications, including: <ul style="list-style-type: none"> <li>• Client Applications</li> <li>• Embedded SQL Applications</li> </ul> | <ul style="list-style-type: none"> <li>• <i>SQL Stored Procedures and Embedded SQL</i></li> <li>• <i>Teradata Preprocessor2 for Embedded SQL Programmer Guide</i></li> </ul>                         |
| Macros as SQL Applications                                                                                                                | <ul style="list-style-type: none"> <li>• <i>SQL Fundamentals</i></li> <li>• <i>SQL Data Definition Language</i></li> <li>• <i>SQL Data Manipulation Language</i></li> </ul>                          |
| Teradata Database SQL stored procedures as SQL applications                                                                               | <ul style="list-style-type: none"> <li>• <i>SQL Fundamentals</i></li> <li>• <i>SQL Stored Procedures and Embedded SQL</i></li> </ul>                                                                 |
| The EXPLAIN Request Modifier                                                                                                              | <ul style="list-style-type: none"> <li>• <i>SQL Data Manipulation Language</i></li> <li>• <i>SQL Request and Transaction Processing</i></li> </ul>                                                   |
| Third-Party Development, including: <ul style="list-style-type: none"> <li>• TS/API</li> <li>• Workload Management APIs</li> </ul>        | <ul style="list-style-type: none"> <li>• <i>Teradata Transparency Series/ Application Programming Interface User Guide</i></li> <li>• <i>Workload Management API: PM/API and Open API</i></li> </ul> |



## CHAPTER 10 **Data Distribution and Data Access Methods**

---

This chapter describes how Teradata Database handles data distribution, including the normalization of data and referential integrity, and data access methods.

### **Teradata Database Indexes**

An index is a physical mechanism used to store and access the rows of a table. Indexes on tables in a relational database function much like indexes in books, they speed up information retrieval.

In general, Teradata Database uses indexes to:

- Distribute data rows.
- Locate data rows.
- Improve performance.

Indexed access is usually more efficient than searching all rows of a table.

- Ensure uniqueness of the index values.

Only one row of a table can have a particular value in the column or columns defined as a unique index.

Teradata Database supports the following types of indexes:

- Primary
- Partitioned Primary
- Secondary
- Join
- Hash
- Special indexes for referential integrity

These indexes are discussed in the following sections.

### **Primary Indexes**

You can create a table with a Unique Primary Index (UPI), a Non-Unique Primary Index (NUPI), or No Primary Index (NoPI).

| IF you create a table with ... | THEN ...                                                                  |
|--------------------------------|---------------------------------------------------------------------------|
| a UPI                          | the PI is a column, or columns, that has no duplicate values.             |
| a NUPI                         | the PI is a column, or columns, that may have duplicate values.           |
| NoPI                           | there is no PI column and rows are not hashed based on any column values. |

## Primary Indexes and Data Distribution

Unique Primary Indexes (UPIs) guarantee uniform distribution of table rows.

Nonunique Primary Indexes (NUPIs) can cause skewed data. While not a guarantor of uniform row distribution, the degree of uniqueness of the index will determine the degree of uniformity of the distribution. Because all rows with the same PI value are distributed to the same AMP, columns with a small number of distinct values that are repeated frequently do not make good PI candidates.

The most efficient access methods to get data in a table is through the PI. For this reason, choosing a PI should take the following design goal into consideration: choosing a PI that gives good distribution of data across the AMPs must be balanced against choosing a PI that reflects the most common usage pattern of the table.

## Primary Key

A Primary Key (PK), a term that comes from data modeling, defines a column, or columns, that uniquely identify a row in a table. Because it is used for identification, a PK cannot be null. There must be something in that column, or columns, that uniquely identify it. Moreover, PK values should not be changed. Historical information, as well as relationships with others tables, may be lost if a PK is changed or re-used.

A PK is a logical relational database concept. It may or may not be the best column, or columns, to choose as a PI for a table.

## Foreign Key

A Foreign Key (FK) identifies table relationships. They model the relationship between data values across tables. Relational databases, like Teradata Database, permit data values to associate across more than one table.

Thus each FK a table may have must exist somewhere as a PK. That is, there must be referential integrity between FKs and PKs.

## Relationships Between Primary Indexes and Primary Keys

The following table describes some of the relationships between PKs and PIs.



| Primary Key                                                                                                                                                                                                                                                     | Primary Index                                                                                      |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------|
| Identifies a row uniquely.                                                                                                                                                                                                                                      | Distributes rows.                                                                                  |
| Does not imply access path.                                                                                                                                                                                                                                     | Defines most common access path.                                                                   |
| Must be unique.                                                                                                                                                                                                                                                 | May be unique or nonunique.                                                                        |
| May not be null.                                                                                                                                                                                                                                                | May be null.                                                                                       |
| Causes a Unique Primary Index (UPI) or Unique Secondary Index (USI) to be created.                                                                                                                                                                              | N/A                                                                                                |
| Constraint used to ensure referential integrity.                                                                                                                                                                                                                | Physical access mechanism.                                                                         |
| Required by Teradata Database only if referential integrity checks are to be performed.                                                                                                                                                                         | Defined for most production tables. Some staging tables may not have a primary index (NoPI table). |
| <ul style="list-style-type: none"> <li>• If Teradata Database performs referential integrity checks, then the column limit is 64.</li> <li>• If Teradata Database performs no referential integrity checks, then there is no arbitrary column limit.</li> </ul> | 64-column limit.                                                                                   |
| Values should not be changed if you want to maintain data integrity and preserve historical relations among tables.                                                                                                                                             | Values can be changed.                                                                             |

The columns chosen for the UPI of a table are frequently the same columns identified as the PK during the data modeling process, but no hard-and-fast rule makes this so. In fact, physical database design considerations often lead to a choice of columns other than those of the primary key for the PI of a table.

## Partitioned Primary Indexes

Both Unique Primary Indexes (UPIs) and Nonunique Primary Indexes (NUPIs) can be partitioned, though a non-partitioned PI is the default Teradata Database PI.

A Partitioned Primary Index (PPI), like a nonpartitioned PI, provides an access path to the rows in the base table, as well as global temporary tables, volatile tables, and noncompressed join indexes via the PI values.

When a table or join index is created with a PPI, the rows are hashed to the appropriate AMPs based on the PI columns and assigned to an appropriate partition based on the value of a partitioning expression that you define when you create or alter the table. Once assigned to a partition, the rows are stored in row hash order.

## Multilevel Partitioned Primary Index

Teradata Database supports a Multilevel Partitioned Primary Index (MLPPI) wherever a Partitioned Primary Index (PPI) is supported. Multilevel partitioning allows each partition to be subpartitioned. Each level must define at least two partitions. The number of levels of partitioning cannot exceed 15. The product of the number of partitions defined cannot exceed 65, 535. The number of levels of partitioning may be restricted further by other limits such as, for example, the maximum size of the table header or Data Dictionary entry sizes.

## Comparing Partitioned and Nonpartitioned Primary Indexes

PPIs are designed to optimize range queries while also providing efficient PI join strategies. A range query requests data that falls within specified boundaries.

The following table provides a comparison of PPI and PI index capabilities.

| Capabilities                                                | Partitioned                 | Nonpartitioned |
|-------------------------------------------------------------|-----------------------------|----------------|
| Hash distributed to the AMPs by the hash of the PI columns. | Yes                         | Yes            |
| Partitioned on each AMP on some set of columns.             | Yes                         | No             |
| Ordered by hash of the PI columns on each AMP.              | Yes (within each partition) | Yes            |

## Secondary Indexes

Secondary Indexes (SIs) allow access to information in a table by alternate, less frequently used paths and improve performance by avoiding full table scans.

Although SIs add to table overhead, in terms of disk space and maintenance, you can drop and recreate SIs as needed.

SIs:

- Do not affect the distribution of rows across AMPs.
- Can be unique or nonunique.
- Are used by the Optimizer when the indexes can improve query performance.
- Can be useful for NoPI tables.

## Secondary Index Subtables

The system builds subtables for all SIs. The subtable contains the index rows that associate the SI value with one or more rows in the base table.

When column values change or when new rows are added to the base table, the system updates the rows in the subtable. When you drop the SI, the system removes the subtable and the disk space it used becomes available.

## Comparing Primary and Secondary Indexes

The following table provides a brief comparison of PI and SI features.

| Feature                              | Primary | Secondary       |
|--------------------------------------|---------|-----------------|
| Can be unique or nonunique           | Both    | Both            |
| Affects row distribution             | Yes     | No              |
| Create and drop dynamically          | No      | Yes             |
| Improves access                      | Yes     | Yes             |
| Create using multiple data types     | Yes     | Yes             |
| Requires separate physical structure | No      | Yes, a subtable |
| Requires extra processing overhead   | No      | Yes             |

## Join Indexes

A Join Index (JI) is an indexing structure containing columns from one or more base tables.

Some queries can be satisfied by examining only the JI when all referenced columns are stored in the index. Such queries are said to be *covered* by the JI. Other queries may use the JI to qualify a few rows, then refer to the base tables to obtain requested columns that are not stored in the JI. Such queries are said to be *partially-covered* by the index.

Because Teradata Database supports multitable, partially-covering JIs, all types of JIs, except the aggregate JI, can be joined to their base tables to retrieve columns that are referenced by a query but are not stored in the JI. Aggregate JIs can be defined for commonly-used aggregation queries.

Much like SIs, JIs impose additional processing on insert and delete operations and update operations which change the value of columns stored in the JI. The performance trade-off considerations are similar to those for SIs.

### Single-table Join Indexes

Join indexes are similar to base tables in that they support a primary index, which can be used for direct access to one or a few rows.

A single-table JI is an index structure that contains rows from only a single-table. This type of structure has been found to be very useful by Teradata Database users because it provides an alternative approach (primary index) to directly accessing data.

## Multitable Join Indexes

When queries frequently request a particular join, it may be beneficial to predefine the join with a multitable JI. The Optimizer can use the predefined join instead of performing the same join repetitively.

## Aggregate Join Indexes

When query performance is of utmost importance, aggregate JIs offer an extremely efficient, cost-effective method of resolving queries that frequently specify the same aggregate operations on the same column or columns. When aggregate JIs are available, the system does not have to repeat aggregate calculations for every query.

You can define an aggregate JI on two or more tables, or on a single-table. A single-table aggregate JI includes a summary table with:

- A subset of columns from a base table.
- Additional columns for the aggregate summaries of the base table columns.

You can create an aggregate JI using:

- SUM function  
A SUM aggregate JI contains a hidden column containing the row count, so that AVERAGE can be calculated from the JI.
- COUNT function
- GROUP BY clause

## Sparse Join Indexes

Another capability of the JI allows you to index a portion of the table using the WHERE clause in the CREATE JOIN INDEX statement to limit the rows indexed. You can limit the rows that are included in the JI to a subset of the rows in the table based on an SQL query result.

Any JI, whether simple or aggregate, multitable or single-table, can be sparse. For example, the following DDL creates J1, which is an aggregate JI containing only the sales records from 2007:

```
CREATE JOIN INDEX J1 AS
SELECT storeid, deptid, SUM(sales_dollars)
FROM sales
WHERE EXTRACT(year FROM sales_date) = 2007
GROUP BY storeid, deptid;
```

When you enter a query, the Optimizer determines whether accessing J1 gives the correct answer and is more efficient than accessing the base tables. This sparse JI would be selected by the Optimizer only for queries that restricted themselves to data from the year 2007.

## Hash Indexes

The hash index provides an index structure that can be hash-distributed to AMPs in various ways. The index has characteristics similar to a single-table JI with a row identifier that provides transparent access to the base table. A hash index may be simpler to create than a corresponding JI.

The hash index has been designed to improve query performance in a manner similar to a single-table JI. In particular, you can specify a hash index to:

- Cover columns in a query so that the base table does not need to be accessed.
- Serve as an alternate access method to the base table in a join or retrieval operation.

## Index Specification

User tables can be created with or without a primary index. Tables without a primary index (NoPI) are typically used as staging tables. If you do not explicitly specify a PRIMARY INDEX clause, NO PRIMARY INDEX clause, or PRIMARY KEY or UNIQUE constraint in the CREATE TABLE statement, Teradata Database creates a table with the first column as the nonunique primary index by default. (You can modify this behavior by using the DBS Control utility and changing the value of the PrimaryIndexDefault General field.)

### Creating Indexes

The following table provides general information about creating indexes.

| To specify a...                  | Use the following statement...                                                      | And the following clause... |
|----------------------------------|-------------------------------------------------------------------------------------|-----------------------------|
| Unique Primary Index (UPI)       | CREATE TABLE                                                                        | UNIQUE PRIMARY INDEX        |
| Nonunique Primary Index (NUPI)   | CREATE TABLE                                                                        | PRIMARY INDEX               |
| No Primary Index (NoPI) Table    | CREATE TABLE                                                                        | NO PRIMARY INDEX            |
| Unique Secondary Index (USI)     | CREATE TABLE                                                                        | UNIQUE INDEX                |
|                                  | CREATE INDEX                                                                        | UNIQUE                      |
| Nonunique Secondary Index (NUSI) | CREATE TABLE                                                                        | INDEX                       |
|                                  | CREATE INDEX                                                                        | N/A                         |
| Join Index (JI)                  | CREATE JOIN INDEX<br><b>Note:</b> A JI can provide an index across multiple tables. | N/A                         |
| Hash Index                       | CREATE HASH INDEX                                                                   | N/A                         |

Indexes are also created when the PRIMARY KEY and UNIQUE constraints are specified.

## Strengths and Weaknesses of Various Types of Indexes

Teradata Database does not require or allow users to explicitly dictate how indexes should be used for a particular query. The Optimizer costs all of the reasonable alternatives and selects the one that is estimated to be the least expensive.

The object of any query plan is to return accurate results as quickly as possible. Therefore, the Optimizer uses an index or indexes only if the index speeds up query processing. In some cases, the Optimizer processes the query without using any index.

Optimizer index selection for a query plan:

- Can have a direct impact on overall Teradata Database performance.
- Is not always a straightforward process.
- Is based partly on usage expectations.

The following table assumes execution of a simple SELECT statement and explains the strengths and weaknesses of some of the various indexing methods.

| This access method...          | Has the following strengths...                                                                                                                                                                                                                                                 | And the following possible drawbacks...                                                                                                                                                                                    |
|--------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Unique Primary Index (UPI)     | <ul style="list-style-type: none"><li>• is the most efficient access method when the SQL statement contains the PI value</li><li>• involves one AMP and one row</li><li>• requires no spool file (for a simple SELECT)</li><li>• can obtain the most granular locks</li></ul>  | none, in the context of a SELECT statement specifying a PI value. However, a poorly chosen PI can cause poor overall performance in a large workload.                                                                      |
| Nonunique Primary Index (NUPI) | <ul style="list-style-type: none"><li>• provides efficient access when the SQL statement contains the PI value</li><li>• involves one AMP</li><li>• can obtain granular locks</li><li>• may not require a spool file as long as the number of rows returned is small</li></ul> | <ul style="list-style-type: none"><li>• may slow down INSERTs for a SET table with no USIs.</li><li>• may decrease the efficiency of SELECTs containing the PI value when some values are repeated in many rows.</li></ul> |
| Unique Secondary Index (USI)   | <ul style="list-style-type: none"><li>• provides efficient access when the SQL statement contains the USI values, and you do not specify PI values</li><li>• involves two AMPs and one row</li><li>• requires no spool file (for a simple SELECT)</li></ul>                    | requires additional overhead for INSERT, UPDATE, MERGE, and DELETE statements.                                                                                                                                             |

| This access method...                            | Has the following strengths...                                                                                                                                                                                                                                                                                                                                                                              | And the following possible drawbacks...                                                                                                                                                                                                                                                                                                                                                                               |
|--------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Nonunique Secondary Index (NUSI)                 | <ul style="list-style-type: none"> <li>provides efficient access when the number of rows per value in the table is relatively small</li> <li>involves all AMPS and probably multiple rows</li> <li>provides access using information that may be more readily available than a UPI value, such as employee last name, compared to an employee number</li> <li>may require a spool file</li> </ul>           | <ul style="list-style-type: none"> <li>requires additional overhead for INSERT, UPDATE, MERGE, and DELETE statements.</li> <li>will not be used by the Optimizer if the number of data blocks accessed is a significant percentage of the data blocks in the table because the Optimizer will determine that a full table scan is cheaper.</li> </ul>                                                                 |
| Full table scan                                  | <ul style="list-style-type: none"> <li>accesses each row only once</li> <li>provides access using any arbitrary set of column conditions</li> </ul>                                                                                                                                                                                                                                                         | <ul style="list-style-type: none"> <li>examines every row.</li> <li>usually requires a spool file possibly as large as the base table.</li> </ul>                                                                                                                                                                                                                                                                     |
| Multitable join index (JI)                       | <ul style="list-style-type: none"> <li>can eliminate the need to perform certain joins and aggregates repetitively</li> <li>may be able to satisfy a query without referencing the base tables</li> <li>can have a different PI from that of the base table</li> <li>can replace an NUSI or a USI</li> </ul>                                                                                                | <ul style="list-style-type: none"> <li>requires additional overhead for INSERT, UPDATE, MERGE, and DELETE statements for any of the base tables that contribute to the multitable JI.</li> <li>usually is not suitable for data in tables subjected to a large number of daily INSERT, UPDATE, MERGE, and DELETE statements.</li> <li>imposes some restrictions on operations performed on the base table.</li> </ul> |
| Single-table join index (JI)<br>or<br>hash index | <ul style="list-style-type: none"> <li>can isolate frequently used columns (or their aggregates for JIs only) from those that are seldom used</li> <li>can reduce number of physical I/Os when only commonly used columns are referenced</li> <li>can have a different PI from that of the base table</li> </ul>                                                                                            | <ul style="list-style-type: none"> <li>requires additional overhead for INSERT, UPDATE, MERGE, and DELETE statements.</li> <li>imposes some restrictions on operations performed on the base table.</li> </ul>                                                                                                                                                                                                        |
| Sparse join index (JI)                           | <ul style="list-style-type: none"> <li>can be stored in less space than an ordinary JI</li> <li>reduces the additional overhead associated with INSERT, UPDATE, MERGE, and DELETE statements to the base table when compared with an ordinary JI</li> <li>can exclude common values that occur in many rows to help ensure that the Optimizer chooses to use the JI to access less common values</li> </ul> | <ul style="list-style-type: none"> <li>requires additional overhead for INSERT, UPDATE, MERGE, and DELETE statements to the base table.</li> <li>imposes some restrictions on operations performed on the base table.</li> </ul>                                                                                                                                                                                      |

## Hashing

Teradata Database uses hashing to distribute data for tables with a PI to disk storage and uses indexes to access the data.

Because the architecture of Teradata Database is massively parallel, it requires an efficient means of distributing and retrieving its data. That efficient method is hashing. Virtually all Teradata Database indexes are based on (or partially based on) row hash values rather than table column values.

For PIs, Teradata Database obtains a row hash by hashing the values of the PI columns. The row hash and a sequence number, which is assigned to distinguish between rows with the same row hash within a table, are collectively called a row identifier and uniquely identify each row in a table. A partition identifier is also part of the row identifier in the case of PPI tables. For more information on PPI, see [“Partitioned Primary Indexes” on page 113](#).

For SIs, Teradata Database computes a hash value using the hash of the values of the SI columns. This value is used for access when an SI value is specified in the SQL. The SI subtable records the hash value for the SI, the actual value of the index columns (for synonym resolution), and a list of primary index row identifiers for the table being indexed.

## Identity Column

Identity Column is a column attribute option defined in the ANSI standard. When associated with a column, this attribute causes the system to generate a unique, table-level number for every row that is inserted into the table.

Identity columns have many applications, including the automatic generation of UPIs, USIs, primary keys, and surrogate keys. For example, an identity column can serve as a UPI to ensure even data distribution when you import data from a system that does not have a PI.

For more information about indexes, see [“Teradata Database Indexes” on page 111](#).

## Normalization

Normalization is the process of reducing a complex database schema into a simple, stable one. Generally this process involves removing redundant attributes, keys, and relationships from the conceptual data model.

### Normal Forms

Normalization theory is constructed around the concept of *normal forms* that define a system of constraints. If a relation meets the constraints of a particular normal form, we say that relation is in normal form.



By definition, a relational database is always normalized to first normal form, because the column values are always *atomic*. That is, a column can contain one and only one value. Null is considered a value in this context.

But to simply leave it at that invites a number of problems including redundancy and potential update anomalies. The higher normal forms were developed to correct those problems.

## First, Second, and Third Normal Forms

First, second, and third normal forms are stepping stones to the Boyce-Codd normal form and, when appropriate, the higher normal forms.

### First Normal Form

First normal form (1NF) is definitive of a relational database. If we are to consider a database relational, then all relations in the database are in 1NF.

We say a relation is in 1NF if all fields within that relation are atomic. We sometimes refer to this concept as the elimination of repeating groups from a relation. Furthermore, first normal form allows no hierarchies of data values.

### Second Normal Form

Second normal form (2NF) deals with the elimination of circular dependencies from a relation. We say a relation is in 2NF if it is in 1NF and if every non-key attribute is fully dependent on the entire Primary Key.

A non-key attribute is any attribute that is not part of the Primary Key for the relation.

### Third Normal Form

Third normal form (3NF) deals with the elimination of non-key attributes that do not describe the Primary Key.

For a relation to be in 3NF, the relationship between any two non-Primary Key columns, or groups of columns, in a relation must *not* be one-to-one in either direction.

We say attributes are mutually independent if none of them is functionally dependent on any combination of the others. This mutual independence ensures that we can update individual attributes without any danger of affecting any other attribute in a row.

The following list of benefits summarizes the advantages of implementing a normalized logical model in 3NF.

- Greater number of relations
- More PI choices
- Optimal distribution of data
- Fewer full table scans

## Referential Integrity

Traditional referential integrity is the concept of relationships between tables, based on the definition of a primary key and a foreign key. The concept states that a row cannot exist in a table with a non-null value for a *referencing* column if an equal value does not exist in a *referenced* column.

Using referential integrity, you can specify columns within a *referencing* table that are foreign keys for columns in some other *referenced* table. You must define referenced columns as either primary key columns or unique columns.

Referential integrity is a reliable mechanism that prevents accidental database inconsistencies when you perform inserts, merges, updates, and deletes.

### Referential Integrity Terminology

We use the following terms to explain the referential integrity concept.

| Term         | Definition                                                                                                                                          |
|--------------|-----------------------------------------------------------------------------------------------------------------------------------------------------|
| Parent Table | The table referred to by a Child table. Also called the “referenced table.”                                                                         |
| Child Table  | A table in which the referential constraints are defined. Also called the “referencing table.”                                                      |
| Parent Key   | A candidate key in the parent table.                                                                                                                |
| Primary Key  | With respect to referential integrity, a primary key is a parent table column set that is referred to by a foreign key column set in a child table. |
| Foreign Key  | With respect to referential integrity, a foreign key is a child table column set that refers to a primary key column set in a parent table.         |

### Referencing (Child) Table

We call the referencing table the Child table, and we call the specified Child table columns the referencing columns. Referencing columns should be of the same number and have the same data type as the referenced table key.

### Referenced (Parent) Table

A Child table must have a parent table, and the referenced table is referred to as the Parent table. The parent key columns are the referenced columns.

### Importance of Referential Integrity

Referential integrity is important, because it keeps you from introducing errors into your database. Suppose you have an Order Parts table like the following.

| Order Number | Part Number | Quantity |
|--------------|-------------|----------|
| PK           |             | Not Null |
| FK           | FK          |          |
| 1            | 1           | 110      |
| 1            | 2           | 275      |
| 2            | 1           | 152      |

Part number and order number, each foreign keys in this relation, also form the composite primary key.

Suppose you were to delete the row defined by the primary key value 1 in the PART NUMBER table. The foreign key for the first and third rows in the ORDER PART table would now be inconsistent, because there would be no row in the PART NUMBER table with a primary key of 1 to support it. Such a situation shows a loss of referential integrity.

Teradata Database provides referential integrity to prevent this from happening. If you try to delete a row from the PART NUMBER table for which you have specified referential integrity, the database management system will not allow you to remove the row if the part number is referenced in child tables.

Besides data integrity and data consistency, referential integrity provides these benefits.

| Benefit                            | Description                                                                                                                                                                                                                 |
|------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Increases development productivity | You do not need to code SQL statements to enforce referential integrity constraints because Teradata Database automatically enforces referential integrity.                                                                 |
| Requires fewer written programs    | All update activities are programmed to ensure that referential integrity constraints are not violated, because Teradata Database enforces referential integrity in all environments. Additional programs are not required. |
| Allows optimizations               | Referential integrity allows optimizations to occur, such as join elimination.                                                                                                                                              |

## For More Information

For more information on the topics presented in this chapter, see the following Teradata Database books.

| If you want to learn more about... | See...                                                                                                                                                                                                                                                                                                                                         |
|------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Teradata Database Indexes          | <ul style="list-style-type: none"> <li>• <i>Database Design</i></li> <li>• <i>SQL Fundamentals</i></li> </ul>                                                                                                                                                                                                                                  |
| Primary Indexes                    | <ul style="list-style-type: none"> <li>• <i>Database Administration</i></li> <li>• <i>Database Design</i></li> <li>• <i>SQL Data Definition Language</i></li> <li>• <i>SQL Fundamentals</i></li> </ul>                                                                                                                                         |
| No Primary Index (NoPI) Tables     | <ul style="list-style-type: none"> <li>• <i>Database Design</i></li> <li>• <i>SQL Data Definition Language</i></li> <li>• <i>SQL Data Manipulation Language</i></li> <li>• <i>SQL Request and Transaction Processing</i></li> </ul>                                                                                                            |
| Partitioned Primary Indexes        | <ul style="list-style-type: none"> <li>• <i>Database Design</i></li> <li>• <i>SQL Data Definition Language</i></li> <li>• <i>SQL Functions, Operators, Expressions, and Predicates</i></li> <li>• <i>SQL Fundamentals</i></li> <li>• <i>SQL Data Manipulation Language</i></li> <li>• <i>SQL Request and Transaction Processing</i></li> </ul> |
| Secondary Indexes                  | <ul style="list-style-type: none"> <li>• <i>Database Administration</i></li> </ul>                                                                                                                                                                                                                                                             |
| Join Indexes                       | <ul style="list-style-type: none"> <li>• <i>Database Design</i></li> <li>• <i>SQL Data Definition Language</i></li> <li>• <i>SQL Fundamentals</i></li> <li>• <i>SQL Request and Transaction Processing</i></li> </ul>                                                                                                                          |
| Hash Indexes                       | <ul style="list-style-type: none"> <li>• <i>Database Design</i></li> <li>• <i>SQL Data Definition Language</i></li> </ul>                                                                                                                                                                                                                      |
| Index Specification                | <i>SQL Data Manipulation Language</i>                                                                                                                                                                                                                                                                                                          |
| Hashing                            | <ul style="list-style-type: none"> <li>• <i>Database Design</i></li> <li>• <i>SQL Data Definition Language</i></li> <li>• <i>SQL Request and Transaction Processing</i></li> </ul>                                                                                                                                                             |
| Identity column                    | <i>SQL Data Definition Language</i>                                                                                                                                                                                                                                                                                                            |
| Normalization                      | <i>Database Design</i>                                                                                                                                                                                                                                                                                                                         |
| Referential integrity              |                                                                                                                                                                                                                                                                                                                                                |

## CHAPTER 11 **Concurrency Control and Transaction Recovery**

---

This chapter describes the concurrency control in relational database management systems and how to use transaction journaling (permanent journaling) to recover lost data or restore an inconsistent database to a consistent state.

### **About Concurrency Control**

Concurrency control prevents concurrently running processes from improperly inserting, deleting, or updating the same data. A system maintains concurrency control through two mechanisms:

- Transactions
- Locks

### **Transactions**

Transactions are a mandatory facility for maintaining the integrity of a database while running multiple, concurrent operations.

#### **Definition of a Transaction**

A transaction is a logical unit of work and the unit of recovery. The requests nested within a transaction must either all happen or not happen at all. Transactions are atomic. A partial transaction cannot exist.

#### **Definition of Serializability**

A set of transactions is serializable if the set produces the same result as some arbitrary serial execution of those same transactions for arbitrary input.

A set of transactions is correct only if it is serializable. The Two-Phase Locking (2PL) protocol ensures the serializability of transactions.

The two phases are the “growing phase” and the “shrinking phase”.

| In the...       | A transaction must...                                                                                                                                                                                               |
|-----------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| growing phase   | acquire a lock on an object before operating on it.<br><br>The Teradata Optimizer requests locks as close to the beginning of the transaction as possible.                                                          |
| shrinking phase | never acquire any more locks after it has released a lock.<br><br>Lock release is an all-or-none operation. Once acquired, locks are not released until the transaction has committed or is completely rolled back. |

## Transaction Semantics

Teradata Database supports both ANSI transaction semantics and Teradata transaction semantics. A system parameter specifies the default transaction mode for a site. However, you can override the default for a session.

Teradata Database returns a failure when a transaction operating in Teradata semantics mode issues a COMMIT statement. Teradata Database supports the ANSI COMMIT statement in ANSI transaction mode.

## ANSI Mode Transactions

All ANSI transactions are implicitly opened. Either of the following events opens an ANSI transaction:

- Execution of the first SQL request in a session.
- Execution of the first request following the close of a previous transaction.

Transactions close when the application performs a COMMIT, ROLLBACK, or ABORT request.

When the transaction contains a DDL statement, including DATABASE and SET SESSION, which are considered DDL statements in this context, the statement must be the last request in the transaction other than the transaction closing statement.

A session executing under ANSI transaction semantics allows neither the BEGIN TRANSACTION statement, the END TRANSACTION statement, nor the two-phase commit protocol. When an application submits these statements in ANSI mode, the database software generates an error.

In ANSI mode, the system rolls back the entire transaction if the current request:

- Results in a deadlock.
- Performs a DDL statement that aborts.
- Executes an explicit ROLLBACK or ABORT statement.

Teradata Database accepts the ABORT and ROLLBACK statements in ANSI mode, including conditional forms of those statements. If the system detects an error for either a single or multistatement request, it only rolls back that request, and the transaction remains open, except in special circumstances.

Application-initiated, asynchronous aborts also cause full transaction rollback in ANSI mode.

## Teradata Mode Transactions

Teradata mode transactions can be either implicit or explicit.

An explicit, or user-generated, transaction is a single set of BEGIN TRANSACTION/END TRANSACTION statements surrounding one or more requests.

All other requests are implicit transactions.

Consider the following transaction:

```
BEGIN TRANSACTION;
DELETE FROM Employee
WHERE Name = 'Smith T';
UPDATE Department
SET EmpCount=EmpCount-1
WHERE DeptNo=500;
END TRANSACTION;
```

If an error occurs during the processing of either the DELETE or UPDATE statement within the BEGIN TRANSACTION and END TRANSACTION statements, the system restores both Employee and Department tables to the states at which they were before the transaction began.

If an error occurs during a Teradata transaction, then the system rolls back the *entire* transaction.

## Locks

A lock is a means of controlling access to some resource. Teradata Database can lock several different types of resources in several different ways.

### Overview of Teradata Database Locking

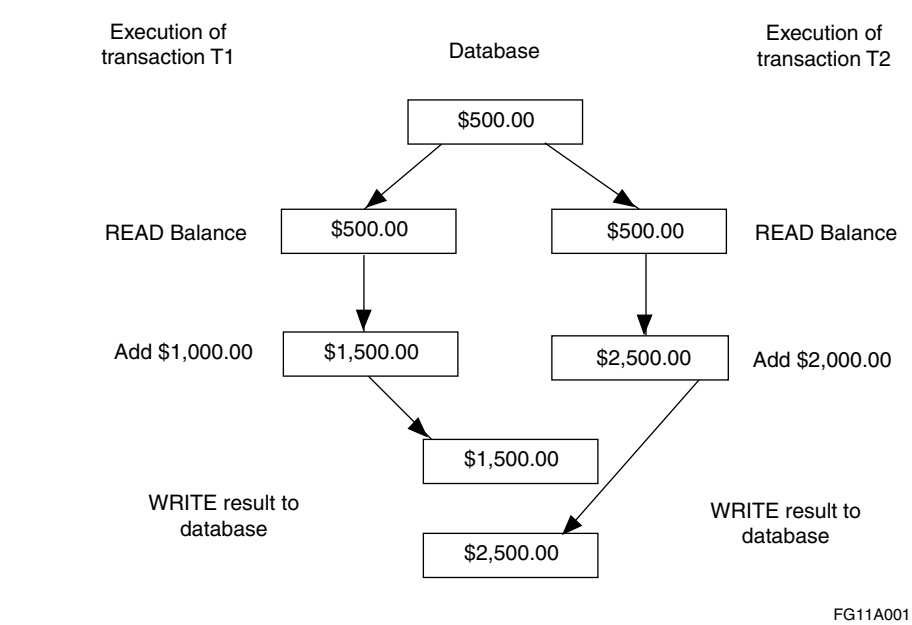
Most locks used on Teradata Database resources are obtained automatically. Users can upgrade the severity of a lock with the LOCKING request modifier but never downgrade the severity of a lock. The data integrity requirement of a request decides the type of lock that the system uses.

A request for a resource locked by another user is queued (in the case of a conflicting lock level) until the process using the resource releases its lock on that resource. A user can specify that the request be aborted if the lock cannot be obtained immediately.

Locking Database Management Systems

The lost update anomaly best explains why database management systems, in which multiple processes are accessing the same database, require locks.

The following figure provides an example of this anomaly.



This example shows a nonserialized set of transactions. If locking had been in effect, the database would not have been able to add \$3000.00 to \$500.00 and get two different and wrong results.

The example demonstrates the most common problem encountered in a transaction processing system without locks. Although several other problems arise when locking is not in effect, the lost update problem sufficiently illustrates the need for locking.

Lock Levels

Teradata Database lock manager can lock the following objects.

| Object Locked | Description                                                                                       |
|---------------|---------------------------------------------------------------------------------------------------|
| Database      | Locks rows of all tables in the database.                                                         |
| Table         | Locks all rows in the table and any index and fallback subtables.                                 |
| Row hash      | Locks the primary copy of a row and all rows that share the same hash code within the same table. |



## Locking Severities

Users can apply four different types of locking severities on Teradata Database resources. The following table explains these types.

| Lock Type | Description                                                                                                                                                                                                                                        |
|-----------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Exclusive | The requester has exclusive privileges to the locked resource. No other process can read from, write to, or access the locked resource in any way.                                                                                                 |
| Write     | The requester has exclusive privileges to the locked resource except for readers not concerned with data consistency.                                                                                                                              |
| Read      | Several users can hold Read locks on a resource, during which the system permits no modification of that resource.<br><br>Read locks ensure consistency during read operations such as those that occur during a SELECT statement.                 |
| Access    | The requestor is willing to accept minor inconsistencies of the data while accessing the database (an approximation is good enough).<br><br>An access lock permits modifications on the underlying data while the SELECT operation is in progress. |

This same information is illustrated in the following table.

| Lock Request | Lock Type Held |         |         |         |           |
|--------------|----------------|---------|---------|---------|-----------|
|              | None           | Access  | Read    | Write   | Exclusive |
| Access       | Granted        | Granted | Granted | Granted | Queued    |
| Read         | Granted        | Granted | Granted | Queued  | Queued    |
| Write        | Granted        | Granted | Queued  | Queued  | Queued    |
| Exclusive    | Granted        | Queued  | Queued  | Queued  | Queued    |

## Automatic Database Lock Levels

Teradata Database applies most of its locks automatically. The following table illustrates how Teradata Database applies different locks for various types of SQL statements.

| Type of SQL Statement | Locking Level by Access Type |                      | Locking Mode |
|-----------------------|------------------------------|----------------------|--------------|
|                       | UPI/NUPI/USI                 | NUSI/Full Table Scan |              |
| SELECT                | Row Hash                     | Table                | Read         |
| UPDATE                | Row Hash                     | Table                | Write        |
| DELETE                | Row Hash                     | Table                | Write        |

| Type of SQL Statement                               | Locking Level by Access Type |                      | Locking Mode |
|-----------------------------------------------------|------------------------------|----------------------|--------------|
|                                                     | UPI/NUPI/USI                 | NUSI/Full Table Scan |              |
| INSERT                                              | Row Hash                     | Not applicable       | Write        |
| CREATE DATABASE<br>DROP DATABASE<br>MODIFY DATABASE | Not applicable               | Database             | Exclusive    |
| CREATE TABLE<br>DROP TABLE<br>ALTER TABLE           | Not applicable               | Table                | Exclusive    |

## Deadlocks and Deadlock Resolution

A deadlock occurs when transaction 1 places a lock on resource A, and then needs to lock resource B. But resource B has already been locked by transaction 2, which in turn needs to place a lock on resource A. This state of affairs is called a deadlock. To resolve a deadlock, Teradata Database aborts the younger transaction, the one that has held the resource the least amount of time, and performs a rollback.

If you used BTEQ to submit the transaction, the database reports the deadlock abort to BTEQ. BTEQ resubmits only the request that caused the error (the default behavior), not the complete transaction. Because this can result in partially committed transactions, you must take care when writing a BTEQ script to ensure that the transaction is one request. For example, a statement in BTEQ ends with a semicolon (;) as the last non-blank character in the line. Thus, BTEQ sees the following example as two requests:

```
sel * from x;
sel * from y;
```

However, if you write these same statements in the following way, BTEQ sees them as only one request:

```
sel * from x
; sel * from y;
```

## Host Utility Locks

The locking operation that the client-resident Teradata Archive/Recovery utility uses is different from the locking operation that Teradata Database performs. Teradata Database documentation and utilities frequently refer to archive locks as HUT (Host Utility) locks.

### HUT Lock Types

Teradata Database places HUT locks as follows.

| Lock Type  | Object Locked                                                                                                                                                            |
|------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Read       | Any object being archived.                                                                                                                                               |
| Group Read | Rows of a table being archived if and only if the table is defined for an after-image permanent journal and if you select the appropriate option on the ARCHIVE command. |
| Write      | Permanent journal table being restored.                                                                                                                                  |
| Write      | All tables in a ROLLFORWARD or ROLLBACKWARD during recovery operations.                                                                                                  |
| Write      | Journal table being deleted.                                                                                                                                             |
| Exclusive  | Any object being restored.                                                                                                                                               |

## HUT Lock Characteristics

HUT locks have the following characteristics:

- Associated with the currently logged-on user who entered the statement rather than with a job or transaction.
- Placed only on objects on the AMPs that are participating in a utility operation.
- Placed at the cluster level during a CLUSTER dump.
- Placed for a user on an object at one level never conflicts with another level of lock on the same object for the same user.
- Remain active until they are released either by the RELEASE LOCK option of an ARC command or by the execution of a Teradata SQL RELEASE LOCK statement after a utility operation completes.
- Automatically reinstated following Teradata Database restart if they had not been released.

## Recovery and Transactions

Recovery is a process by which an inconsistent database is brought back to a consistent state.

Transactions play the critical role in this process because they are used to “play back” (using the term in its most general sense) a series of updates to the database, either taking it back to some earlier state or bringing it forward to a current state.

## System and Media Recovery

The following sections describe the behavior of Teradata Database when it encounters different types of errors or failures.

## System Restarts

Unscheduled restarts occur for one of the following reasons:

- AMP or disk failure
- Software failure
- Disk parity error

Failures and errors affect all software recovery in the same way. Hardware failures take the affected component offline and it remains offline until repaired or replaced.

## Transaction Recovery

Two types of automatic transaction recovery can occur:

- Single transaction recovery
- Database recovery

The following table details what happens when the two automatic recovery mechanisms take place.

| This recovery type... | Happens when Teradata Database...                                                                                                                                                                                                                                                                      |
|-----------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| single transaction    | aborts a single transaction because of: <ul style="list-style-type: none"><li>• Transaction deadlock</li><li>• User error</li><li>• User-initiated abort command</li><li>• An inconsistent data table</li></ul> Single transaction recovery uses the transient journal to effect its data restoration. |
| database              | performs a restart for one of the following reasons: <ul style="list-style-type: none"><li>• Hardware failure</li><li>• Software failure</li><li>• User command</li></ul>                                                                                                                              |

## Down AMP Recovery

When an AMP fails to come online during system recovery, which is done using the down AMP Recovery Journal, Teradata Database continues to process requests using fallback data. When the down AMP comes back online, down AMP recovery procedures begin to bring the data for the AMP up-to-date as follows.

| IF there are...                        | THEN the AMP recovers... |
|----------------------------------------|--------------------------|
| a large number of rows to be processed | offline.                 |
| only a few rows to be processed        | online.                  |

After all updates are made, we consider the AMP to be fully recovered.

## Down Subtable Recovery

Teradata Database can isolate errors to a data or index subtable, or to a range of rows (region) in a data or index subtable. When these errors are found, Teradata Database marks only the affected subtable or region down and takes a snapshot dump.

In-progress transactions that require the down subtable or region will be aborted. Subsequent transactions that require access to the down subtable or region will not be allowed until the problem is fixed. Transactions that do not require access to the affected subtable or region will continue to run.

## Two-Phase Commit Protocol

Two-phase commit (2PC) is a protocol for assuring update consistency across distributed databases in which each participant in the transaction commit operation votes to either commit or abort the changes. Participants wait before committing a change until they know that all participants can commit.

A participant is a “database manager” that performs some work on behalf of the transaction and that commits or aborts changes to the database. A participant can also be a coordinator of participants at a lower level.

By voting to commit, a participant guarantees that it can either commit or roll back its part of the transaction, even if it crashes before receiving the result of the vote.

The 2PC protocol allows the development of Customer Information Control System (CICS) and Information Management System (IMS) applications that can update one or more Teradata Database databases or databases, or both under some other DBMS in a synchronized manner. The result is that all updates requested in a defined unit of work will either succeed or fail.

## For More Information

For more information on the topics presented in this chapter, see the following Teradata Database and Teradata Tools and Utilities books.

| IF you want to learn more about... | See...                                        |
|------------------------------------|-----------------------------------------------|
| Transactions                       | <i>SQL Request and Transaction Processing</i> |
| ANSI Mode Transactions             |                                               |
| Teradata Mode Transactions         |                                               |

| IF you want to learn more about... | See...                                                                                                                                                                                                               |
|------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Host Utility Locks                 | <ul style="list-style-type: none"><li>• <i>SQL Request and Transaction Processing</i></li><li>• <i>Database Administration</i></li></ul>                                                                             |
| Locks                              | <ul style="list-style-type: none"><li>• <i>SQL Request and Transaction Processing</i></li><li>• <i>Database Administration</i></li></ul>                                                                             |
| System and Media Recovery          | <ul style="list-style-type: none"><li>• <i>Database Administration</i></li><li>• <i>Utilities</i></li></ul>                                                                                                          |
| Two-phase Commit Protocol          | <ul style="list-style-type: none"><li>• <i>Teradata Director Program Reference</i></li><li>• <i>IBM CICS Interface for Teradata Reference</i></li><li>• <i>IBM IMS/DC Interface for Teradata Reference</i></li></ul> |

## CHAPTER 12 The Data Dictionary

---

This chapter provides information about the Data Dictionary.

### Data Dictionary

The Teradata Database Data Dictionary is composed of tables and views that reside in the system user named DBC. The tables are reserved for use by the system and contain metadata about the objects in the system, privileges, system events, and system usage. The views provide access to the information in the tables. The tables contain current definitions, control information, and general information about the following:

- Authorization
- Accounts
- Character sets
- Columns
- Constraints
- Databases
- Disk Space
- End Users
- Events
- External stored procedures
- Indexes
- JAR and ZIP archive files
- Logs
- Macros
- Privileges
- Profiles
- Resource usage
- Roles
- Rules
- Sessions and session attributes
- Statistics
- Stored procedures
- Tables
- Translations
- Triggers
- User-defined functions
- User-defined methods
- User-defined types
- Views

### Data Dictionary Users

Users typically use views to obtain information on objects, and database administrators generally create and update tables that system views reference. The following table contains information about what is stored in the Data Dictionary when you create some of the most important objects.

| WHEN you create...           | THEN the definition of the object is stored along with the following details...                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
|------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| a table                      | <ul style="list-style-type: none"> <li>• Table location, identification, database name, table name, creator name, version, database name, and user names of all owners in the hierarchy.</li> <li>• Each column in the table, including column name, data type, length, and phrases.</li> <li>• User/creator access privileges on the table.</li> <li>• Indexes defined for the table.</li> <li>• Constraints defined for the table.</li> <li>• Table backup and protection (including fallback and permanent journaling status).</li> <li>• Date and time the object was created.</li> </ul> |
| a database                   | <ul style="list-style-type: none"> <li>• Database name, creator name, owner name, and account name</li> <li>• Space allocation (if any) including: <ul style="list-style-type: none"> <li>• Permanent</li> <li>• Spool</li> <li>• Temporary</li> </ul> </li> <li>• Number of fallback tables</li> <li>• Collation type</li> <li>• Creation time stamp</li> <li>• The date and time the database was last altered and the name that altered it</li> <li>• Role and profile names</li> <li>• Revision numbers for the UDF library and any XSP libraries by Application Category</li> </ul>      |
| a view or macro              | <ul style="list-style-type: none"> <li>• The text of the view or macro.</li> <li>• Creation time attributes.</li> <li>• User and creator access privileges on the view or macro.</li> </ul>                                                                                                                                                                                                                                                                                                                                                                                                   |
| a stored procedure           | <ul style="list-style-type: none"> <li>• Creation time attributes of the stored procedure.</li> <li>• Parameters of the stored procedure, including parameter name, parameter type, data type, and default format.</li> <li>• User and creator access privileges on the stored procedure.</li> </ul>                                                                                                                                                                                                                                                                                          |
| an external stored procedure | <ul style="list-style-type: none"> <li>• C/C++ source code and object code for the external stored procedure if its language is not Java.</li> <li>• External stored procedure name</li> <li>• External name</li> <li>• Data types of the parameters</li> <li>• Source file language</li> <li>• Data accessing characteristic</li> <li>• Parameter passing convention</li> <li>• Execution protection mode</li> <li>• Character type</li> <li>• Platform type</li> </ul>                                                                                                                      |



| WHEN you create...               | THEN the definition of the object is stored along with the following details...                                                                                                                                                                                                                                                                                                                                                                                                 |
|----------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| a JAR                            | <ul style="list-style-type: none"> <li>• Java object code for the JAR</li> <li>• JAR name</li> <li>• External name</li> <li>• Platform type</li> <li>• Revision number</li> </ul>                                                                                                                                                                                                                                                                                               |
| a Java external stored procedure | <ul style="list-style-type: none"> <li>• Java external stored procedure name</li> <li>• External file reference</li> <li>• Data types of the parameters</li> <li>• Source file language</li> <li>• Data accessing characteristic</li> <li>• Parameter passing convention</li> <li>• Execution protection mode</li> <li>• Character type</li> <li>• Platform type</li> </ul>                                                                                                     |
| a Java user-defined function     | <ul style="list-style-type: none"> <li>• Function call name</li> <li>• Specific name</li> <li>• External name</li> <li>• Data types of the parameters</li> <li>• Function class</li> <li>• Source file language</li> <li>• Data accessing characteristic</li> <li>• Parameter passing convention</li> <li>• Deterministic characteristic</li> <li>• Null-call characteristic</li> <li>• Execution protection mode</li> <li>• Character type</li> <li>• Platform type</li> </ul> |

| WHEN you create...      | THEN the definition of the object is stored along with the following details...                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
|-------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| a trigger               | <ul style="list-style-type: none"> <li>• The IDs of the: <ul style="list-style-type: none"> <li>• Table</li> <li>• Trigger</li> <li>• Database and subject table database</li> <li>• User who created the trigger</li> <li>• User who last updated the trigger</li> </ul> </li> <li>• Timestamp for the last update</li> <li>• Indexes</li> <li>• Trigger name and <ul style="list-style-type: none"> <li>• Whether the trigger is enabled</li> <li>• The event that fires the trigger</li> <li>• The order in which triggers fire</li> </ul> </li> <li>• Default character set</li> <li>• Creation text and time stamp</li> <li>• Overflow text, that is, trigger text that exceeds a specified limit</li> <li>• Fallback tables</li> </ul> |
| a user-defined function | <ul style="list-style-type: none"> <li>• C source code and object code for the user defined function if its language is not Java.</li> <li>• Function call name</li> <li>• Specific name</li> <li>• External name</li> <li>• Data types of the parameters</li> <li>• Function class</li> <li>• Source file language</li> <li>• Data accessing characteristic</li> <li>• Parameter passing convention</li> <li>• Deterministic characteristic</li> <li>• Null-call characteristic</li> <li>• Execution protection mode</li> <li>• Character type</li> <li>• Platform type</li> </ul>                                                                                                                                                          |

| WHEN you create...    | THEN the definition of the object is stored along with the following details...                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|-----------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| a user-defined method | <ul style="list-style-type: none"> <li>• C source code and object code for the user defined method</li> <li>• Function call name</li> <li>• Specific name</li> <li>• External name</li> <li>• Data types of the parameters</li> <li>• Function class</li> <li>• Source file language</li> <li>• Data accessing characteristic</li> <li>• Parameter passing convention</li> <li>• Deterministic characteristic</li> <li>• Null-call characteristic</li> <li>• Execution protection mode</li> <li>• Character type</li> <li>• Platform type</li> </ul>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| a user-defined type   | <p>DBC.UDTInfo - one entry per UDT</p> <ul style="list-style-type: none"> <li>• Type name</li> <li>• Type kind (Distinct or Structured)</li> <li>• Whether the type is instantiable</li> <li>• Default transform group (name)</li> <li>• Ordering form (full ordering or equals only - distinct and structured types are always full)</li> <li>• Ordering category (map or relative - distinct and structured types are always map)</li> <li>• Ordering routine ID</li> <li>• Cast count</li> </ul> <p>DBC.UDTCast - one entry per cast for a UDT</p> <ul style="list-style-type: none"> <li>• Whether cast is implicit assignment</li> <li>• Cast routine ID</li> </ul> <p>DBC.UDFInfo - one entry for the UDT's auto-generated default constructor</p> <ul style="list-style-type: none"> <li>• Entries are the same as for a regular (C/C++) UDF</li> </ul> <p>DBC.UDTTransform - one entry for the UDT's transform</p> <ul style="list-style-type: none"> <li>• Default transform group name</li> <li>• ToSQL routine ID</li> <li>• FromSQL routine ID</li> </ul> |

| WHEN you create... | THEN the definition of the object is stored along with the following details...                                                                                                                                                                                                                                                                                                                                                                                                                    |
|--------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| a user             | <ul style="list-style-type: none"><li>• User name, creator name, and owner name</li><li>• Password string and password change date</li><li>• Space allocation including:<ul style="list-style-type: none"><li>• Permanent</li><li>• Spool</li><li>• Temporary</li></ul></li><li>• Default account, database, collation, character type, and date form</li><li>• Creation timestamp</li><li>• Name and time stamp of the last alteration made to the user</li><li>• Role and profile name</li></ul> |

## Data Dictionary Views

You can examine the information about the system tables in database DBC directly or through a series of views. Typically, you use views to obtain information on the objects in the Data Dictionary rather than querying the actual tables, which can be very large. The database administrator controls who has access to views.

### Users of Data Dictionary Views

Some Data Dictionary views may be restricted to special types of users, while others are accessible by all users. The database administrator controls access to views by granting privileges. The following table defines the information needs of various types of users.

| This type of user...   | Needs to know...                                                                                                                                                                                                                                                                        |
|------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| End                    | <ul style="list-style-type: none"><li>• Objects to which the user has access</li><li>• Types of access available to the user</li><li>• The privileges the user has granted to other users</li></ul>                                                                                     |
| Supervisory            | <ul style="list-style-type: none"><li>• How to create and organize databases</li><li>• How to monitor space usage</li><li>• How to define new users</li><li>• How to allocate access privileges</li><li>• How to create indexes</li><li>• How to perform archiving operations</li></ul> |
| Database administrator | <ul style="list-style-type: none"><li>• Performance</li><li>• Status and statistics</li><li>• Errors</li><li>• Accounting</li></ul>                                                                                                                                                     |

| This type of user...            | Needs to know...                                                                                                                                                                                                          |
|---------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Security administrator          | <ul style="list-style-type: none"> <li>Access logging rules generated by the execution of BEGIN LOGGING statements</li> <li>Results of access checking events, logged as specified by the access logging rules</li> </ul> |
| Operations and Recovery Control | Archive and recovery activities                                                                                                                                                                                           |

## SQL Access to the Data Dictionary

Every time you log on to Teradata Database, perform an SQL query, or type a password, you are using the Data Dictionary.

For security and data integrity reasons, the only SQL DML command you can use on the Data Dictionary is the SELECT statement. You cannot use the INSERT, UPDATE, MERGE, or DELETE SQL statements to alter the Data Dictionary, except for some Data Dictionary tables, such as the AccLogTbl table or the EventLog table.

The Data Dictionary contains Unicode system views. However, to maintain backwards compatibility, Teradata provides compatibility system views to translate the object name to the Kanji/Latin. For more information on Unicode and compatibility system views, see *Data Dictionary* and *Database Administration*.

You can use SELECT to examine any view in the Data Dictionary to which your database administrator has granted you access. For example, if you need to access information in the Personnel database, then you can query the DBC.DatabasesV view as shown:

```
SELECT Databasename,
 Creatorname,
 Ownername,
 Permspace
FROM DBC.DatabasesV
WHERE Databasename='Personnel'
;
```

The query above produces a report like this:

| <u>Databasename</u> | <u>Creatorname</u> | <u>Ownername</u> | <u>Permspace</u> |
|---------------------|--------------------|------------------|------------------|
| Personnel           | Jones              | Jones            | 1,000,000        |

# For More Information

For more information on the topics presented in this chapter, see the following Teradata Database book.

| IF you want to learn more about... | See...                 |
|------------------------------------|------------------------|
| Data Dictionary                    | <i>Data Dictionary</i> |
| Data Dictionary Views              |                        |
| SQL Access to the Data Dictionary  |                        |

## CHAPTER 13 **International Language Support**

---

This chapter describes the capabilities of Teradata Database international language support.

### **Character Set Overview**

A character set (sometimes called a code page) is simply a way of representing characters on a computer. There are many ways to represent characters on a computer, so there are many character sets in use today.

Because different characters are needed for different languages, character sets are often designed to support a particular language. Even for the same language, many different character sets may exist.

When computers or computer applications exchange character data, it is important that they either use the same character set or properly convert the data from one character set to the other during the transfer process. Otherwise, the data received by one machine may no longer have the same meaning as it had before the transfer. The same issue exists for numeric data, but there are fewer ways used to represent numbers, so it is not as big a problem.

A character set has a repertoire of characters it supports, a representation for character strings, and an implied collation based on this representation.

#### **About Repertoire**

Consider English for example. To write English, you need the alphabetic characters, A–Z, the digits, 0–9, and various punctuation characters. Many applications also commonly require the characters a–z, the lower case counterparts of A–Z.

If an application is written in French, you need the alphabetic characters that are required for English, plus accented characters, for example, the é. However, some applications may need accented characters for English as well. The word résumé, borrowed from French, is often displayed in its accented form in English text. Similarly, ö may be used in English text to spell coördinate.

You can see that a repertoire comprises the characters that we need to write a language, and clearly, what we include in our repertoire determines what we can write, and how we must write it. A character string from one character set can only be translated correctly to another character set if every character in the source string exists in the repertoire of the target character set.

## Character Representation

Representing strings of characters is essentially a two-step process:

- Creating a mapping between each character required and an integer.
- Devising an encoding scheme for placing a sequence of numbers into memory.

The simplest systems map the required characters to small integers between 0 and 255, and encode sequences of characters as sequences of bytes with the appropriate numeric values.

Representing characters for repertoires that require more than 256 characters, such as Japanese, Chinese, and Korean, requires more complex schemes.

## External and Internal Character Sets

Client systems communicate with Teradata Database using their own *external* format for numbers and character strings.

Teradata Database converts numbers and strings to its own *internal* format when importing the data, and converts numbers and strings back to the appropriate form for the client when exporting the data.

This allows data to be exchanged between mutually incompatible client data formats. Take for example, channel-attached clients using EBCDIC-based character sets and network-attached clients using ASCII-based character sets. Both clients can access and modify the same data in Teradata Database.

## Character Data Translation

Teradata Database translates the characters:

- Received from a client system into a form suitable for storage and processing on the server.
- Returned to a client into a form suitable for storage, display, printing, and processing on that client.

Thus, the server communicates with each client in the language of that client without the need for additional software. It is essential, for this process to work properly, for the database to be informed of the correct character set used by each client.

## What Teradata Database Supports

Teradata Database supports many external *client* character sets and allows each application to choose the internal *server* character set best suited to each column of character data in Teradata Database. Because of automatic translation, it is normally the repertoire of characters required that determines the server character set to use.

No matter which server character set you chose, communication with the client is always in the client character set (also known as the session charset).



# Teradata Database Character Data Storage

Teradata Database uses internal server character sets to represent user data and data in the Data Dictionary within the system.

## Internal Server Character Sets

Server character sets include:

- LATIN
- UNICODE
- KANJI1
- KANJISJIS
- GRAPHIC

## User Data

User data refers to character data that you store in a character data type column on Teradata Database.

## Object Names in the Data Dictionary

Teradata Database stores a variety of character data in the Data Dictionary, including the names of the following objects:

- |             |                              |                          |
|-------------|------------------------------|--------------------------|
| • Tables    | • Triggers                   | • User-defined functions |
| • Databases | • JAR and ZIP archive files  | • User-defined types     |
| • Users     | • Join indexes               | • User-defined methods   |
| • Columns   | • Hash indexes               | • Authorizations         |
| • Views     | • Stored procedures          |                          |
| • Macros    | • External stored procedures |                          |

## Language Support Modes

During system initialization (sysinit) the database administrator can optimize Teradata Database for one of two language support modes:

- Standard
- Japanese

The language support mode determines the:

- Character set that Teradata Database uses to store system dictionary data.
- Default character set for user data.

| IF you enable this language support mode ... | THEN Teradata Database stores object names using this character set ...                    | AND stores character data other than object names using this character set ... |
|----------------------------------------------|--------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------|
| Standard                                     | UNICODE<br><br>For backward compatibility, object names are processed internally as LATIN  | LATIN.                                                                         |
| Japanese                                     | UNICODE<br><br>For backward compatibility, object names are processed internally as KANJI1 | UNICODE.                                                                       |

## Overriding the Default Character Set for User Data

The language support mode sets the default server character set for a user if the DEFAULT CHARACTER SET clause does not appear in the CREATE USER or MODIFY USER statement.

To override the default character set for a user, you can use the DEFAULT CHARACTER SET clause in a CREATE USER statement. You can also specify a server character set for a character column when you created the table.

## Standard Language Support Mode

Although the Data Dictionary uses UNICODE to store object names in standard language support mode and Japanese language support mode, the permissible range of characters in object names depends on the language support mode.

For user data, the default server character set is LATIN.

### LATIN Character Set

Standard language support provides Teradata Database internal coding for the entire set of printable characters from the ISO 8859-1 (Latin1) and ISO 8859-15 (Latin9) standard, including diacritical marks such as ä, ñ, Ÿ, Œ, and œ, though the Z with caron in Latin9 is not supported. ASCII control characters are also supported for the standard language set.

**Note:** ASCII, as used here, represents the characters that can be stored as the LATIN server character set, referred to as Teradata LATIN. EBCDIC is Teradata Database extended ASCII mapped to the corresponding EBCDIC code points.

### Compatible Languages

The LATIN server character set that Teradata Database uses in Standard language support mode is sufficient for you to use client character sets that support the international languages listed in the following table.

| International Languages That are Compatible with Standard Language Support |          |                                   |                 |
|----------------------------------------------------------------------------|----------|-----------------------------------|-----------------|
| Albanian                                                                   | English  | Germanic                          | Portuguese      |
| Basque                                                                     | Estonian | Greenlandic                       | Rhaeto-Romantic |
| Breton                                                                     | Faroese  | Icelandic                         | Romance         |
| Catalonian                                                                 | Finnish  | Irish Gaelic<br>(new orthography) | Samoan          |
| Celtic                                                                     | French   | Italian                           | Scottish Gaelic |
| Cornish                                                                    | Frisian  | Latin                             | Spanish         |
| Danish                                                                     | Galician | Luxemburgish                      | Swahili         |
| Dutch                                                                      | German   | Norwegian                         | Swedish         |

## Japanese Language Support Mode

Although the Data Dictionary uses UNICODE to store object names in Japanese language support mode and standard language support mode, the permissible range of characters in object names depends on the language support mode.

For user data, the default server character set is UNICODE.

### Advantages of Storing User Data Using UNICODE

Unicode is a 16-bit encoding of virtually all characters in all current languages in the world. Teradata Database UNICODE server character set supports Unicode 4.1 and is designed eventually to store all character data on the server.

UNICODE may be used to store all characters from all single-byte and multibyte client character sets. User data stored as UNICODE can be shared among heterogeneous clients.

### User DBC Default Character Set

With Japanese language mode support, the default character set for user DBC is UNICODE.

## Extended Support

Extended support allows you to customize Teradata Database to provide additional support for local character set usage.

A sufficiently privileged user can create single-byte and multibyte client character sets that support, with certain constraints, any subset of the Unicode repertoire. Moreover, such a user can customize a collation for the entire Unicode repertoire.

Extended support is available on systems that have been enabled with standard language support or Japanese language support.

## For More Information

For more information on the topics presented in this chapter, see *International Character Set Support*.

## CHAPTER 14 Query and Database Analysis Tools

---

Although the Optimizer executes complex strategic and tactical queries in an efficient manner, when you look at a query plan, you may find it difficult to understand how and why the Optimizer chose the plan for a given query.

This chapter discusses:

- Tools found in Teradata Analyst Pack, a set of tools designed to help automate and simplify query plan analysis. These include:
  - Teradata Visual Explain (VE)
  - Teradata System Emulation Tool (SET)
  - Teradata Index Wizard
  - Teradata Statistics Wizard
- Teradata Database Query Analysis Tools (DBQAT), tools designed to improve the overall performance analysis capabilities of Teradata Database. These include:
  - Query Capture Facility
  - Database Query Log
  - Target Level Emulation (TLE)
  - Database Object Use Count

### Teradata Visual Explain

Teradata Visual Explain (VE) is a tool that visually depicts the execution plan of complex SQL requests in a graphical manner.

Teradata VE presents a graphical view of the request broken down into discrete steps showing the flow of data during execution. It has the ability to compare multiple execution plans side by side.

Because comparing optimized queries is easier with Teradata VE, application developers and database administrators can fine-tune the SQL statements so that Teradata Database can access data in the most effective manner.

In order to view an execution plan using Teradata VE, the execution plan information must first be captured in the Query Capture Database (QCD) by means of the Query Capture Facility (QCF).

Teradata VE reads the execution plan, which has been stored in a QCD, and turns it into a series of icons.

## Teradata System Emulation Tool

When Target Level Emulation (TLE) information is stored on a test system, Teradata System Emulation Tool (SET) enables you to generate and examine the query plans using the test system Optimizer as if the plans were processed on the production system.

Using Teradata SET you can:

- Change system configuration details, including DBS Control fields, and table demographics and model the impact of various changes on SQL statement performance.
- Determine the source of various Optimizer-based production problems.
- Provide an environment in which Teradata Index Wizard can produce recommendations for a production system workload.

## Teradata Index Wizard

Teradata Index Wizard analyzes SQL queries and suggests candidate indexes to enhance their performance.

The workload definitions, supporting statistical and demographic data, and index recommendations are stored in various QCD tables.

Using data from a QCD or the Database Query Log (DBQL), Teradata Index Wizard:

- Recommends, using an INITIATE PARTITION ANALYSIS statement, the potential performance benefits from adding a partitioning expression to one or more tables in a given workload.  
The statement does not recommend the complete removal of any defined partitioning expressions. It considers, however, the alteration of an existing partitioning expression if a Partitioned Primary Index (PPI) table is explicitly included in the table\_list.
- Recommends secondary indexes for the tables based on workload details, including data demographics, that are captured using the QCF.
- Enables you to validate index recommendations before implementing the new indexes.
- Enables you to perform what-if analysis on the workload. Teradata Index Wizard allows you to determine whether your recommendations actually improve query performance.
- Interfaces with other Teradata Tools and Utilities, such as Teradata SET to perform offline query analysis by importing the workload of a production system to a test system
- Uses Teradata Visual Explain and Compare tools to provide a comparison of the query plans with and without the index recommendations.

Teradata Index Wizard can be started from Teradata Visual Explain, Teradata SET, and Teradata Statistics Wizard. Teradata Index Wizard can also open these applications to help in your evaluation of recommended indexes.

Demographics

Teradata Index Wizard needs demographic information to perform index analysis and to make recommendations. You can collect the following types of data demographics using SQL:

- Query demographics  
Use the INSERT EXPLAIN statement with the WITH STATISTICS and DEMOGRAPHICS clauses to collect table cardinality and column statistics.
- Table demographics  
Use the COLLECT DEMOGRAPHICS statement to collect the row count and the average row size in each of the subtables in each AMP on the system.

Teradata Statistics Wizard

Teradata Statistics Wizard automates the process of collecting statistics for a particular workload or selecting arbitrary indexes or columns for collection or re-collection. In addition, Teradata Statistics Wizard enables you to validate the proposed statistics on a production system. The validation capability enables you to verify the performance of the proposed statistics before applying the recommendations. The following table contains information about the capabilities of Teradata Statistics Wizard.

| You can...                                                      | For...                                                               |
|-----------------------------------------------------------------|----------------------------------------------------------------------|
| select a workload                                               | analysis and receive recommendations based on the results.           |
| select a database or select several tables, indexes, or columns | analysis and receive recommendations based on the results.           |
| defer                                                           | the schedule for the collection or recollection of statistics.       |
| display and modify statistics                                   | a column or index.                                                   |
| receive recommendations                                         | analysis that is based on table demographics and general heuristics. |

As changes are made within a database, Teradata Statistics Wizard identifies those changes and recommends which tables should have statistics collected, based on age of data and table growth, and the columns/indexes that may benefit from having statistics defined and collected for a specific workload. The administrator is then given the opportunity to accept or reject the recommendations.

# Query Capture Facility

The Query Capture Facility (QCF) allows you to capture the steps of query plan information into well-known user tables called the Query Capture Database (QCD).

The source of the captured data is produced by the Teradata Database Optimizer, which outputs the text of the EXPLAIN request modifier detailing the final stage of optimization.

Applications of QCF include:

- Store all query plans for customer queries. You can then compare and contrast queries as a function of software release, hardware platform, and hardware configuration.
- Provide data so that you can generate your own detailed analyses of captured query steps using standard SQL DML statements and third-party query management tools.
- Provide the input for these utilities: Teradata Index Wizard (which recommends index definitions for improved efficiency), Teradata Statistics Wizard (which allows you to validate proposed statistics), and Teradata VE (which presents a graphical view of the flow of data during query execution).

# Target Level Emulation

Teradata Database supports Target Level Emulation (TLE) both on Teradata Database and in the client as follows.

| Teradata Database supports...        | On the...                                                                                                     |
|--------------------------------------|---------------------------------------------------------------------------------------------------------------|
| Target Level Emulation (TLE)         | Teradata Database.                                                                                            |
| Teradata System Emulation Tool (SET) | client. For information about Teradata SET, see <a href="#">“Teradata System Emulation Tool” on page 150.</a> |

Teradata Database provides the infrastructure for TLE. You can use the standard SQL interface to capture the system configuration details and table demographics on one system and store them on another.

Usually the information is obtained from a production system, then stored on a smaller test or development system. With this capability, the Optimizer can generate access plans similar to those that are generated on a production system. You can use the plans to analyze Optimizer-related production problems. This information can also be used by Teradata SET.



## Database Query Log

The Database Query Log (DBQL) is a Teradata Database feature that provides a series of predefined tables that can store historical records of queries and their duration, performance, and target activity based on rules you specify.

DBQL is flexible enough to log information on the variety of SQL requests that run on Teradata Database, from short transactions to longer-running analysis and mining queries.

You can request that DBQL log particular query information or just a count of qualified queries. You can specify a hierarchy of rules that DBQL applies to sessions. This allows great flexibility in capturing the right amount of logging data for each use of the Teradata Database system. The rules may target:

- Application names, for example FastLoad or MultiLoad
- A user and a specific account
- A user and all accounts
- All users and a specific account
- All users and all accounts

Each rule may specify that the recording criteria be a mix of:

- Summarizing data based on elapsed time, CPU time, or I/O counts as either a series of intervals or a threshold limit
- Reporting details at the request level, including identification, performance, counts, error code, and SQL text
- Additional, optional detailed data, which may include any or all:
  - Objects used in the request
  - Steps and performance data per step
  - Full SQL text
  - Explain text
  - Optimizer query plan information logged as an XML document

You can define rules, for instance, that tell the DBS to log the first 4,000 SQL characters of a query that runs during a session invoked by a specific user under a specific account if the query exceeds a specified time threshold.

To implement DBQL, you use simple SQL requests to control the starting and ending of the logging activity. These rules may be dynamically changed as needed to support management of the Teradata Database system.

## Database Object Use Count

The database administrator and application developer can use Database Object Use Count to capture the number of times an application refers to an object.

Database Object Use Count captures counts for the following:

- Databases
- Tables
- Columns
- Indexes
- Views
- Macros
- Teradata Database stored procedures
- Triggers
- User-defined functions

Object use access information is not counted for EXPLAIN, INSERT EXPLAIN, or DUMP EXPLAIN statements.

Once captured, you can use the information to identify possibly obsolete or unused database objects, particularly those that occupy significant quantities of valuable disk space. Further, Database Object Use Count information can be useful to database query analysis tools like Teradata Index Wizard.

## For More Information

For more information on the topics presented in this chapter, see the following Teradata Tools and Utilities books.

| IF you want to learn more about...                                     | THEN see...                                                                                                                                                                                             |
|------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Teradata Visual Explain                                                | <i>Teradata Visual Explain User Guide</i>                                                                                                                                                               |
| Teradata System Emulation Tool                                         | <i>Teradata System Emulation Tool User Guide</i>                                                                                                                                                        |
| Teradata Index Wizard, including support for Partitioned Primary Index | <ul style="list-style-type: none"><li>• <i>Teradata Index Wizard User Guide</i></li><li>• <i>SQL Request and Transaction Processing</i></li><li>• <i>SQL Data Definition Language</i></li></ul>         |
| Teradata Statistics Wizard                                             | <ul style="list-style-type: none"><li>• <i>SQL Request and Transaction Processing</i></li><li>• <i>Teradata Statistics Wizard User Guide</i></li></ul>                                                  |
| Query Capture Database                                                 | <ul style="list-style-type: none"><li>• <i>SQL Data Definition Language</i></li><li>• <i>SQL Request and Transaction Processing</i></li></ul>                                                           |
| Database Query Log                                                     | <ul style="list-style-type: none"><li>• <i>Database Administration</i></li><li>• <i>Data Dictionary</i></li><li>• <i>Performance Management</i></li><li>• <i>SQL Data Definition Language</i></li></ul> |

| IF you want to learn more about... | THEN see...                                                                                                                                                |
|------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Target Level Emulation             | <ul style="list-style-type: none"><li>• <i>SQL Request and Transaction Processing</i></li><li>• <i>Teradata System Emulation Tool User Guide</i></li></ul> |
| Database Object Use Count          | <i>Database Administration</i>                                                                                                                             |



## CHAPTER 15 Teradata Database Security

This chapter describes the features and methods available to establish and maintain Teradata Database security.

### Security Concepts

Teradata Database security is based on the following concepts.

| Security Element           | Description                                                                                                                                             |
|----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------|
| Database User              | An individual or group of individuals represented by a single user identity.                                                                            |
| Privileges                 | The database privileges explicitly or automatically granted to a user or database.                                                                      |
| Logon                      | The process of submitting user credentials when requesting access to the database.                                                                      |
| Authentication             | The process by which the user identified in the logon is verified.                                                                                      |
| Authorization              | The process that determines the database privileges available to the user.                                                                              |
| Security Mechanism         | A method that provides specific authentication, confidentiality, and integrity services for a database session.                                         |
| Network Traffic Protection | The process for protecting message traffic between Teradata Database and network-attached clients against interception, theft, or other form of attack. |
| Message Integrity          | Checks data sent across the network against what was received to ensure no data was lost or changed.                                                    |
| Access Logs                | Logs that provide the history of users accessing the database and the database objects accessed.                                                        |

For detailed information on these topics, see *Security Administration*.

### Users

Users that access Teradata Database must be defined in the database or a supported directory.

## Permanent Database Users

The CREATE USER statement defines permanent database users. Teradata recommends that the username represent an individual. Each username must be unique in the database.

## Directory-based Users

Directory-based users that access the database must be defined in a supported directory. Creation of a matching database user may or may not be required, depending upon implementation. One or more configuration tasks must be completed before directory users can access the database.

## Proxy Users

Proxy users are end users who access Teradata Database through a middle-tier application set up for trusted sessions. They are authenticated by the application rather than the database. The GRANT CONNECT THROUGH statement assigns role-based database privileges to proxy users. To establish privileges for a particular connection to the database, the application submits the SET QUERY\_BAND statement when it connects the user. The SET QUERY\_BAND statement and the rules for how it is applied to each proxy user must be coded into the application as part of setup for trusted sessions.

# Database Privileges

Users can access only database objects on which they have privileges. The following table lists the types of database privileges and describes how they are acquired by a user.

| Privilege            | Description                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|----------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Implicit (Ownership) | Privileges implicitly granted by the database to the owner of the space in which database objects are created.                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| Automatic            | Privileges automatically provided by the system to: <ul style="list-style-type: none"><li>• The creator of a database, user, or other database object.</li><li>• A newly created user or database.</li></ul>                                                                                                                                                                                                                                                                                                                                                 |
| Inherited            | Privileges that are passed on indirectly to a user based on its relationship to another user or role to which the privileges were granted directly. <ul style="list-style-type: none"><li>• Directory users inherit the privileges of the database users to which they are mapped. Directory users who are members of groups also inherit the privileges defined in external roles to which the groups are mapped.</li><li>• All users inherit the privileges of PUBLIC, the default database user, whether or not they have any other privileges.</li></ul> |

| Privilege        | Description                                                                                                                                                                                                                                    |
|------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Explicit (GRANT) | Privileges granted explicitly to a user or database in one of the following ways: <ul style="list-style-type: none"><li>• GRANT directly to a user or database.</li><li>• GRANT to a role, then GRANT the role to one or more users.</li></ul> |

## Directly Granted Privileges

Privileges can be directly given to users with the GRANT statement. Administrators GRANTing a privilege must have been previously granted the privilege they are granting, as well as the WITH GRANT OPTION privilege on the privilege.

For additional information on how to use SQL statements to GRANT and REVOKE privileges, see *SQL Data Control Language*.

## Roles

Roles can be used to define privileges on database objects for groups of users with similar needs, rather than granting the privileges to individual users. Roles also require less dictionary space than individually granted privileges. Use the CREATE ROLE statement to define each role, then use the GRANT statement to grant roles to users. The CREATE USER statement must also specify the default role for the user. The MODIFY USER statement can be used to assign additional user roles.

A member of a role may access all objects to which a role has privileges. Users can employ the SET ROLE statement to switch from the default to any alternate role of which the user is a member, or use SET ROLE ALL to access all roles.

For more information on use of roles, see *Database Administration*.

### Roles for Proxy Users

Proxy users are users that access the database through a middle-tier application set up to offer trusted sessions. Proxy users are limited to privileges defined in roles that are assigned to them using the GRANT CONNECT THROUGH statement.

For details on using GRANT CONNECT THROUGH, see *Security Administration* and *SQL Data Control Language*.

## External Roles

Use external roles to assign role privileges to directory users. External roles are created exactly like database roles; however, they cannot be granted to directory users because these users do not exist in the database. Instead, directory users must be members of groups that are mapped to external roles in the directory.

Directory users mapped to multiple external roles have access to all of them at logon to the database.

For information on mapping directory users to database objects, see *Security Administration*.

## Profiles

To simplify user management, an administrator can define a profile and assign it to a group of users who share similar values for the following types of parameters:

- Default database assignment
- Spool space capacity
- Temporary space capacity
- Account strings permitted
- Password security attributes

For further information on profiles, see *Database Administration*.

## User Authentication

At logon to the database, users are authenticated, meaning their identity is verified and checked against a list of approved users. Authentication comprises the following elements:

- Authentication method
- Logon formats and controls
- Password formats and controls

### Authentication Method

To identify the method by which they will be authenticated, users must choose a security mechanism as part of the logon string. Security mechanisms also contain configurable properties for customizing user authentication and authorization. If a mechanism is not specified in the logon string, the system will use the default mechanism.

Two categories of authentication are available:

- Teradata Database authentication
- External authentication

#### Teradata Database Authentication

Authentication by Teradata Database requires that the user and its privileges be defined in the database. Teradata Database authentication requires use of the TD2 mechanism, which is the default. Unless another mechanism has been set as the default, TD2 need not be specified at logon.

#### External Authentication

External authentication allows Teradata Database users to be authenticated by an agent running on the same network as Teradata Database and its clients.

External authentication is dependent upon two elements:

- The authenticating agent indicated by the security mechanism specified in the logon.
- The logon type, which determines how user credentials are submitted.



The following table shows the different types of external authentication logons:

| Type                                                 | Description                                                                                                                                   | Requirements                                                                                                                                                                                                                                                                                                                                                                                                          |
|------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Sign-on without resubmitting user credentials</b> |                                                                                                                                               |                                                                                                                                                                                                                                                                                                                                                                                                                       |
| Single Sign-on                                       | Users are authenticated in the client domain. Subsequent logons to Teradata Database do not require them to resubmit a username and password. | <ul style="list-style-type: none"> <li>The username employed to log on to the client domain must match the Teradata Database username.</li> <li>The logon must specify a mechanism that corresponds to the authenticating application (Kerberos, NTLM, or SPNEGO) or it must be the default mechanism.</li> <li>The logon must specify the tdpid of the database, and optional account string information.</li> </ul> |
| <b>Sign-on with user credentials</b>                 |                                                                                                                                               |                                                                                                                                                                                                                                                                                                                                                                                                                       |
| Directory Sign-on                                    | The user is authenticated by the directory.                                                                                                   | <ul style="list-style-type: none"> <li>The user must log on with a directory username.</li> <li>The logon must specify the LDAP mechanism.</li> </ul>                                                                                                                                                                                                                                                                 |
| Sign-on As                                           | The user logs on to Teradata Database with a username and password recognizable by the client domain.                                         | <ul style="list-style-type: none"> <li>The username must be defined in the database.</li> <li>The user must select a mechanism that corresponds to the authenticating application (Kerberos, LDAP, NTLM, or SPNEGO) or it must be the default mechanism.</li> </ul>                                                                                                                                                   |

For complete information on configuring external authentication and the mechanisms that support it, see *Security Administration*.

## Logon Format

Logons to Teradata Database can take the following forms:

- Command line
- Graphical User Interface (GUI)
- Logon from a channel-attached client
- Logon through a middle-tier application

### Command-Line Logon

Users provide the following information when logging on from a network-attached client:

- **.logmech:** Specifies the name of the security mechanism that defines the method by which the user will be authenticated and authorized. If a mechanism is not specified, the logon proceeds using the designated default mechanism.

- **.logdata:** Used only for external authentication. Specifies the username and password, and optionally names an individual username, profile, or role where the logon user has access to multiple user identities, profiles, or roles.
- **.logon:** Must be used for Teradata Database authentication. May be used for external authentication except when the logon requires specification of user=, profile=, or realm= to differentiate among multiple user identities, profiles, or realms.

For either authentication type, .logon specifies the tdpid, username, password, and optional account string information.

**Note:** The format required for the username and password information varies depending on the user is sign-on method.

For more information on command-line logons, see *Security Administration*.

## GUI Logons

Some Teradata Database client applications provide a logon GUI in the form of dialog boxes. The dialog boxes provide fields and buttons that prompt the user to enter the same logon information shown for command-line logons.

For an example of a GUI logon, see *Security Administration*.

## Logons from Channel-Attached Clients

Sessions logged on from channel-attached clients do not support network security features, such as security mechanisms, encryption, or directory management of users. Such logons require submittal of only the username, password, tdpid, and optional account string information using the .logon command.

For information on logons from channel-attached clients, see *Security Administration* and *Teradata Director Program Reference*.

## Logons Through Connection Pools

Some users access the database through a middle-tier application. The application must log on to Teradata Database with a database username. The application then sets up a connection pool that allows individual end users access to the database but does not require that the users formally logon. End user privileges are determined by whether or not the application and its end-users are set up for trusted sessions.

For information, see [“Authorization of Middle-Tier Application Users” on page 164](#).

## Logon Controls

Teradata Database automatically grants permission for all users defined in the database to logon from all connected client systems. But administrators can, for example:

- Modify current or default logon privileges for specific users.
- Give individual users permission to log on to the database only from specific channel or network interfaces.
- Set the maximum number of times a user can submit an unsuccessful logon string.
- Enable authentication of the user by an external application, such as Kerberos or LDAP.

- Restrict access to the database based on the IP address of the machine from which a user logs on.

## Password Format Requirements

Passwords must conform to password format rules, which govern the type and number of characters allowed in the password.

## Password Controls

Teradata Database provides controls to enable the administration of passwords. Administrators can, for example:

- Restrict the content of passwords:
  - Define limits on minimum and maximum password characters.
  - Allow or require that passwords contain upper and lowercase characters, digits, or special characters.
  - Allow or deny use of restricted words.
- Set the number of days for which a password is valid.
- Assign a temporary password.
- Set the lockout time after the user has exceeded the maximum number of logon attempts.
- Define the period during which a user may not reuse a previous password.

# User Authorization

Once users have been authenticated, they are authorized database privileges according to their defined privileges.

## Authorization of Permanent Database Users

Permanent database users are defined in the database with a CREATE USER statement. Once authenticated at logon, permanent users are authorized the following privileges:

- Privileges granted directly to the user with the GRANT statement.
- Privileges indirectly given to the user (automatic, implicit, and inherited privileges).
- Privileges granted to a role that has been granted to the user.

**Note:** Users with more than one role automatically have access to the default role or all their roles, as specified in the CREATE USER statement. The SET ROLE statement allows users to access roles other than their default role.

## Authorization of Directory-Based Users

After being authenticated by the directory, directory-based users are authorized database access privileges according to the following rules:

- If the directory maps users to Teradata Database objects (users, external roles, and profiles), each directory user is authorized the privileges of the objects to which it is mapped.
- If the directory *does not* map users to Teradata Database objects, but the directory username matches a Teradata Database username, the directory user is authorized all the privileges belonging to the matching Teradata Database user.
- If a directory user is neither mapped to any database objects, nor does the directory username match a Teradata Database username, the directory user has no privileges to access the database.

**Note:** One or more setup tasks (depending on implementation) must be completed before a directory user can access the database. For information, see *Security Administration*.

## Authorization of Middle-Tier Application Users

Middle-tier applications may stand between end users and Teradata Database, accepting requests from users, constructing queries from those requests, passing the queries to the database, and then returning results to the users. The middle-tier application logs on to the database, is authenticated as a permanent database user, and establishes a connection pool. The application then authenticates the individual application end users, some of whom may request access to the database through the connection pool.

By default, all end-users accessing the database through a middle-tier application are authorized database privileges and are audited in access logs, based on the single permanent database user identity of the application.

For sites that require end users to be individually identified, authorized, and audited, the middle-tier application can be configured to offer *trusted sessions*. Application end-users that access the database through a trusted session must be set up as *proxy users* and assigned one or more database roles, which determine their privileges in the database. When a proxy user requests database access, the application automatically forwards the user identity and applicable role information to the database.

For further information about the tasks required to set up trusted sessions and proxy users, see *Security Administration*.

## Data Protection

Teradata Database provides the following features to enhance data protection:

- By default, the logon string is encrypted to maintain the confidentiality of the username and password employed to log on to Teradata Database.
- Optional data encryption of messages maintains confidentiality of the data transmitted to and from Teradata Database.
- Automatic integrity checking insures that the data is not corrupted during the encryption/transmission/decryption cycle.

- Optional BAR encryption provides confidentiality of data backups between the BAR server and the storage device.
- Optional SSL/TLS protection for systems using LDAP authentication with simple binding, including:
  - Encryption of the LDAP authentication sequence between Teradata Database and the directory server.
  - Advanced TLS protection, which requires that the database authenticate itself to the directory or that the database and directory mutually authenticate.
- Optional SASL protection for systems using LDAP authentication with Digest-MD5 binding:
  - Protection of the LDAP authentication token exchange sequence between Teradata Database and the directory server.

## Directory Management of Users

Normally, users that log on to Teradata Database have been defined in the database using a CREATE USER request. However, because many potential database users may already be defined in a directory running within the client network, Teradata Database allows for authentication and authorization of users by supported directories. Integration of directory managed users simplifies administration by eliminating the need to create a database instance for every user.

For information on how to set up the database and the directory for integrated user management, see *Security Administration*.

### Supported Directories

Teradata Database is certified for use with the following directories:

- Active Directory
- Active Directory Application Mode (ADAM)
- Novell eDirectory
- Sun Java System Directory Server
- Other LDAPv3-compliant directories may be usable. Contact Teradata Professional Services for integration assistance.

## Database Security Monitoring

To ensure optimal database security, the security administrator should configure the system to audit security events, and then monitor the logs to detect breaches in security and, where necessary, repel security threats. If unauthorized or undesirable activity is detected, the administrator can take remedial actions to address the problem.

## Security Monitoring

Teradata Database provides the capability for two types of user security monitoring:

- All user logon and logoff activity is automatically collected in the Event Log and can be accessed using the DBC.LogOnOffV system view. Listed parameters include:
  - Database username
  - Session number
  - Logon events, including the causes of any unsuccessful logons
- Optional access logging records user attempts to access the database and can be accessed using the DBC.AccessLogV view, including the following access parameters:
  - Type of access
  - Type of request
  - Requesting database username
  - Referenced database object
  - Frequency of access

**Note:** Access log entries are generated each time the database checks a privilege to determine whether or not it will honor a user request.

### Logging of Directory Users

Directory users are logged by directory username rather than by the name of any database user they may be mapped to.

### Logging of Middle-Tier Application Users

Users that access the database through a middle-tier application by means of a trusted session and who are set up as proxy users are logged by their proxy user name. If the middle-tier application and its end users are *not* set up for trusted sessions, all such users will appear in the log as the same username, that is, the name used by the application.

### Enabling and Disabling Access Logging

Use the BEGIN and END LOGGING statements to enable and disable logging and to indicate which access parameters should be logged. Access logging can be set up for almost any database object, for instance, users, databases, or tables.

## Security-related System Views

The Data Dictionary provides a number of security-related system views, including the following:

| View             | Description                                                                                                                                                                                               |
|------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| DBC.AccessLogV   | Each entry indicates a privileges check that has resulted from a user request.                                                                                                                            |
| DBC.AccLogRulesV | Lists the access logging rules contained in each BEGIN and END LOGGING statement. These rules are used by the database to determine the access logging criteria for a particular user or database object. |
| DBC.LogOnOffV    | Lists all logon and logoff activity.                                                                                                                                                                      |
| DBC.LogonRulesV  | Lists the logon rules that result from GRANT and REVOKE LOGON statements. These rules are used by the database to determine whether or not to allow logon privileges to a particular user.                |

For a complete listing of security-related system views, see *Security Administration*.

## Defining a Security Policy

Your security policy should be based on the following considerations:

- Determine your security needs to balance the need for secure data against user needs for quick and efficient data access.
- Review Teradata Database security features to meet your needs.
- Develop a policy that includes both system-enforced and personnel-enforced features.

## Publishing a Security Policy

To ensure administrators and users at your site understand and follow site-specific security procedures, the administrator should create a security handbook. The handbook should summarize how you are using Teradata Database security features for your database and should be published and made available to all users.

A security policy document should include:

- Why security is needed.
- Benefits for both the company and the users of adhering to the security policy.
- A description of the specific implementation of Teradata Database security features at your site.
- Suggested/required security actions for users and administrators to follow.

- Whom to contact when security questions arise.

## For More Information

For more information on the topics presented in this chapter, see the following Teradata Database and Teradata Tools and Utilities books.

| IF you want to learn more about...                                                                                  | See...                                                                                                                      |
|---------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------|
| Security concepts                                                                                                   | <i>Security Administration</i>                                                                                              |
| Users                                                                                                               | <ul style="list-style-type: none"><li>• <i>Security Administration</i></li><li>• <i>SQL Data Control Language</i></li></ul> |
| Database privileges, including information on how to use security-related SQL statements, such as GRANT and REVOKE. |                                                                                                                             |
| User authentication                                                                                                 | <i>Security Administration</i>                                                                                              |
| User authorization, including detailed logon requirements for Teradata Database client applications                 | the user guide for the respective application                                                                               |
| Data protection                                                                                                     | <i>Security Administration</i>                                                                                              |
| Directory management of users                                                                                       |                                                                                                                             |
| Database security monitoring, including security-related system tables and views                                    | <i>Data Dictionary</i>                                                                                                      |
| Defining a security policy                                                                                          | <i>Security Administration</i>                                                                                              |
| Publishing a security policy                                                                                        |                                                                                                                             |



## SECTION 4 **Managing and Monitoring Teradata Database**



## CHAPTER 16 **System Administration**

---

While the system administrator is responsible for creating databases and users, he or she is also responsible for managing the database.

For information on creating database and users, see [Chapter 7: “Database Objects, Databases, and Users.”](#)

### **Session Management**

Users must log on to Teradata Database and establish a session before system administrators can do any system accounting.

#### **Session Requests**

A session is established after the database accepts the username, password, and account number and returns a session number to the process.

Subsequent Teradata SQL requests generated by the user and responses returned from the database are identified by:

- Host id
- Session number
- Request number

The database supplies the identification automatically for its own use. The user is unaware that it exists.

The context for the session also includes a default database name that is often the same as the user name. When the session ends, the system discards the context and accepts no further Teradata SQL statements from the user.

#### **Establishing a Session**

To establish a session, the user logs on to the database.

The procedure varies depending on the client system, the operating system, and whether the user is an application program, or a user in an interactive terminal session using an application program or an interactive user.

## Logon Operands

The logon string can include any of the following operands:

- Optional identifier for the database server, called a tdpid
- User name
- Password
- Optional account number

**Note:** In the Windows environment, the Single Sign-On feature eliminates the need for users to re-submit usernames, passwords, and account ids after they have logged on to Windows using their authorized user names and passwords. For more information about this feature, see [Chapter 15: “Teradata Database Security.”](#)

## Administrative and Maintenance Utilities

A large number of utilities are available to system administrators to perform various administrative and maintenance functions.

The following table lists some of the major Teradata Database utilities.

| Utility                           | Purpose                                                                                                                                                                                                    |
|-----------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Abort Host<br>(aborthost)         | Aborts all outstanding transactions running on a failed host, until the system restarts the host.                                                                                                          |
| AMP Load<br>(ampload)             | Displays the load on all AMP vprocs in a system, including the number of AMP worker tasks (AWTs) available to each AMP vproc, and the number of messages waiting (message queue length) on each AMP vproc. |
| AWT Monitor<br>(awtmon)           | Collects and displays a user-friendly summary of the AMP Worker Task (AWT) in-use count snapshots for the local node or all nodes in Teradata Database.                                                    |
| CheckTable<br>(checktable)        | Checks for inconsistencies between internal data structures such as table headers, row identifiers, and secondary indexes.                                                                                 |
| CNS Run<br>(cnsrun)               | Allows running of database utilities from scripts.                                                                                                                                                         |
| Configuration Utility<br>(config) | Defines AMPs, PEs, and hosts, and describes their interrelationships for Teradata Database.                                                                                                                |
| Control GDO Editor<br>(ctl)       | Displays the fields of the PDE Control Parameters GDO, and allows modification of the settings.                                                                                                            |
| Cufconfig Utility<br>(cufconfig)  | Displays configuration settings for the user-defined function and external stored procedure subsystem, and allows these settings to be modified.                                                           |

| Utility                               | Purpose                                                                                                                                                                                                                                                                                                    |
|---------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Database Initialization Program (DIP) | Executes one or more of the standard DIP scripts packaged with Teradata Database. These scripts create a variety of database objects that can extend the functionality of Teradata Database with additional, optional features.                                                                            |
| DBS Control (dbscontrol)              | Displays the DBS Control Record fields, and allows these settings to be modified.                                                                                                                                                                                                                          |
| Dump Unload/Load (DUL, DULTAPE)       | Saves system dump tables to tape, and restores system dump tables from tape.                                                                                                                                                                                                                               |
| Ferret Utility (ferret)               | Defines the scope of an action, such as a range of tables or selected vprocs, displays the parameters and scope of the action, and performs the action, either moving data to reconfigure data blocks and cylinders, or displaying disk space and cylinder free space percent in use of the defined scope. |
| Filer Utility (filer)                 | Finds and corrects problems within the file system.<br><b>Note:</b> Filer is documented in <i>Support Utilities</i> .                                                                                                                                                                                      |
| Gateway Control (gtwcontrol)          | Modifies default values in the fields of the Gateway Control Globally Distributed Object (GDO).                                                                                                                                                                                                            |
| Gateway Global (gtwglobal)            | Monitors and controls the Teradata Database LAN-connected users and their sessions.                                                                                                                                                                                                                        |
| Lock Display (lokdisp)                | Displays a snapshot capture of all real-time database locks and their associated currently-running sessions.                                                                                                                                                                                               |
| Locking Logger (dumplocklog)          | Logs transaction identifiers, session identifiers, lock object identifiers, and the lock levels associated with currently-executing SQL statements.                                                                                                                                                        |
| Modify MPP List (modmpplist)          | Allows modification of the node list file (mpplist).                                                                                                                                                                                                                                                       |
| Priority Scheduler (schmon)           | Creates, modifies, and monitors Teradata Database process prioritization parameters.<br><br>All processes have an assigned priority based on their Teradata Database session. This priority is used by Priority Scheduler to allocate CPU and I/O resources.                                               |
| Query Configuration (qryconfig)       | Reports the current Teradata Database configuration, including the Node, AMP, and PE identification and status.                                                                                                                                                                                            |
| Query Session (qrysessn)              | Monitors the state of selected Teradata Database sessions on selected logical host IDs.                                                                                                                                                                                                                    |
| Reconfiguration Utility (reconfig)    | Uses the component definitions created by the Configuration Utility to establish an operational Teradata Database.                                                                                                                                                                                         |

| Utility                                               | Purpose                                                                                                                                                                                                                                                                                                  |
|-------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Reconfiguration Estimator<br>(reconfig_estimator)     | Estimates the elapsed time for reconfiguration based upon the number and size of tables on the current system, and provides time estimates for the following phases: <ul style="list-style-type: none"> <li>• Redistribution</li> <li>• Deletion</li> <li>• NUSI building</li> </ul>                     |
| Recovery Manager<br>(rcvmanager)                      | Displays information used to monitor progress of a Teradata Database recovery.                                                                                                                                                                                                                           |
| Resource Check Tools<br>(dbschk, nodecheck, syscheck) | Identifies slow down and potential hangs of Teradata Database, and displays system statistics that help identify the cause of the problem.                                                                                                                                                               |
| Show Locks<br>(showlocks)                             | Displays locks placed by Archive and Recovery and Table Rebuild operations on databases and tables.<br><br>For details Archive and Recovery, see <i>Teradata Archive/Recovery Utility Reference</i> . For details on Table Rebuild, see <i>Utilities Volume 2</i> .                                      |
| System Initializer<br>(sysinit)                       | Initializes Teradata Database. Creates or updates the DBS Control Record and other Globally Distributed Objects (GDOs), initializes or updates configuration maps, and sets hash function values in the DBS Control Record.<br><br><b>Note:</b> Sysinit is documented in <i>Support Utilities</i> .      |
| Table Rebuild<br>(rebuild)                            | Rebuilds tables that Teradata Database cannot automatically recover, including the primary or fallback portions of tables, entire tables, all tables in a database, or all tables in an Access Module Processor (AMP). Table Rebuild can be run interactively or as a background task.                   |
| Tdlocaledef Utility<br>(tdlocaledef)                  | Converts a Specification for Data Formatting file (SDF) into an internal, binary format (a GDO) for use by Teradata Database. The SDF file is a text file that defines how Teradata Database formats numeric, date, time, and currency output.                                                           |
| TDN Statistics<br>(tdnstat)                           | Performs GetStat/ResetStat operations and displays or clears Teradata Database Network Services statistics.                                                                                                                                                                                              |
| TDN Tuner<br>(tdntune)                                | Displays and allows modification of Teradata Network Services tunable parameters. The utility provides a user interface to view, get, or update the Teradata Network Services, which are specific to tunable parameters.                                                                                 |
| Tpareset<br>(tpareset)                                | Resets the PDE and database components of Teradata Database.                                                                                                                                                                                                                                             |
| Two-Phase Commit Console<br>(tpccons)                 | Performs the following two-phase commit (2PC) related functions: <ul style="list-style-type: none"> <li>• Displays a list of coordinators that have in-doubt transactions.</li> <li>• Displays a list of sessions that have in-doubt transactions.</li> <li>• Resolves in-doubt transactions.</li> </ul> |
| Task List<br>(tsklist)                                | Displays information about PDE processes and their tasks.<br><br><b>Note:</b> This utility runs only under Microsoft Windows and Linux.                                                                                                                                                                  |

| Utility                           | Purpose                                                                                                                                                                                                                  |
|-----------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Update DBC<br>(updatedbc)         | Recalculates the PermSpace and SpoolSpace values in the DBASE table for the user DBC, and the MaxPermSpace and MaxSpoolSpace values of the DATABASESPACE table for all databases based on the values in the DBASE table. |
| Update Space<br>(updatespace)     | Recalculates the permanent, temporary, or spool space used by a single database or by all databases in a system.                                                                                                         |
| Verify _pdisks<br>(verify_pdisks) | Verifies that the pdisks on Teradata Database are accessible and are mapped correctly.                                                                                                                                   |
| Vproc Manager<br>(vprocmanager)   | Manages the virtual processors (vprocs). For example, obtains status of specified vprocs, initializes vprocs, forces a vproc to restart, and forces a Teradata Database restart.                                         |

## For More Information

For more information on the topics presented in this chapter, see the following Teradata Database and Teradata Tools and Utilities books.

| IF you want to learn more about...       | See...                                                                                                                                                         |
|------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Roles and profiles                       | <ul style="list-style-type: none"> <li><i>Database Administration</i></li> <li><i>SQL Fundamentals</i></li> <li><i>SQL Data Definition Language</i></li> </ul> |
| Session management                       | <i>Database Administration</i>                                                                                                                                 |
| Administrative and maintenance utilities | <i>Utilities</i>                                                                                                                                               |





## CHAPTER 17 Database Management Tools and Utilities

---

Teradata Database offers a wide variety of utilities, management tools, and peripherals.

Some of these reside on Teradata Database and others are part of the Teradata Tools and Utilities management suite available for installation in client environments.

With database management tools, you can back up and restore important data, save dumps, and investigate and control Teradata Database configuration, user sessions, and various aspects of its operation and performance. Management and analysis tools help keep the database running at optimum performance levels.

### Data Archiving Utilities

Teradata Backup, Archive, and Restore (BAR) supports third party software products that provide data archiving, backup, and restore functions. Teradata utilizes software extensions called TARA (Tiered Archive Restore Architecture) and plug-ins that help connect BAR software to the Teradata Database. The BAR application software offering includes:

- BakBone® NetVault® and NetVault Plug-in

NetVault Teradata plug-in enables NetVault to work with Teradata Database. NetVault allows you to select databases and tables graphically, and specify the types of backups (for example, distributed, online) you want to perform.

It is a flexible, scalable, modular storage management solution designed to backup and restore a Teradata Database of any size. It is also tightly integrated with Teradata making it easy to use.

- Symantec™ NetBackup and NetBackup Extension for Teradata

NetBackup Extension for Teradata is an access module that enables Netbackup to work with Teradata Database. Administrators can schedule automatic, unattended backups for client systems across a network. It supports parallel backups and restores coordinated across multiple hosts in a single Teradata Database. Teradata TARA in the NetBackup framework and is comprised of three components: TARA Server, TARA GUI, and NetBackup Extension for Teradata.

- Tivoli® Storage Manager and Tivoli Storage Manager Teradata Extension  
Tivoli Storage Manager (TSM) Teradata Extension is an access module that enables TSM to work with Teradata Database. Tivoli Storage Manager Teradata Extension manages the I/O interfaces between Teradata ARCMAIN and IBM TSM. Teradata TARA in the TSM framework and is comprised of three components: TARA Server, TARA GUI, and TSM Teradata Extension.

## Teradata Archive/Recovery Utility

Teradata Archive/Recovery utility (ARC), working with BAR application software, writes and reads sequential files on a Teradata client system to archive, restore, recover, and copy Teradata Database table data. Through its associated script language, it also provides an interface between Teradata Backup Application Software solutions and Teradata Database.

Teradata ARC also supports archiving and restoring individual:

- Functions
- Hash indexes
- Join indexes
- Macros
- Stored procedures
- Triggers
- Views

Teradata ARC, working with BAR application software, supports archiving and restoring Teradata Database to any of the following media:

- Tape storage device
- Backup-to-Disk (B2D) storage device

ARC also includes:

- Online archiving of individual tables
- Recovery with rollback and rollforward functions for data tables defined with a journal option. For more information about rollback and rollforward, see [Chapter 11: “Concurrency Control and Transaction Recovery.”](#)

## Teradata Parallel Transporter

Teradata Parallel Transporter (Teradata PT) is an object-oriented client application that provides scalable, high-speed, parallel data extraction, loading, and updating.

Teradata PT uses and expands on the functionality of the traditional Teradata standalone extract and load utilities, that is, FastLoad, MultiLoad, FastExport, and TPump.

Teradata PT supports:

- **Process-specific operators:** Teradata PT jobs are run using *operators*. These are discrete object-oriented modules that perform specific extraction, loading, and updating processes. There are four functional operator types:
  - **Producer operators** that read data from a source and write it to data streams.
  - **Consumer operators** that read from data streams and write it to a data target.
  - **Filter operators** that read data from data streams, perform data filtering operations such selection, validation, and condensing, and then write filtered data to data streams.
  - **Standalone operators** that perform processing that does not involve receiving data from, or sending data to, the data stream.
- **Access modules:** These are software modules that give Teradata PT access to various data stores.
- **A parallel execution structure:** Teradata PT can execute multiple instances of an operator to run multiple and concurrent loads and extracts and perform inline updating of data. Teradata PT maximizes throughput performance through scalability and parallelism.
- **The use of data streams:** Teradata PT distributes data into data streams shared with multiple instances of operators. Data streaming eliminates the need for intermediate data storage (data is not written to disk).
- **A single SQL-like scripting language:** Unlike the standalone extract and load utilities, each of which uses its own scripting language, Teradata PT uses a single script language to specify extraction, loading, and updating operations.
- **A GUI-based Teradata PT Wizard:** The Teradata PT Wizard helps you generate simple Teradata PT job scripts.

## Teradata Parallel Transporter Application Programming Interface

The Teradata Parallel Transporter Application Programming Interface (Teradata PT API) is a set of application programming interfaces used to load into, and extract from, Teradata Database.

Unlike Teradata PT and the standalone utilities, which are script-driven, Teradata PT API is a functional library that is part an application. This allows applications to have more control during load and extract operations. Closer control of the runtime environment simplifies management processes.

## Standalone Data Load and Unload Utilities

### Teradata FastExport

Teradata FastExport extracts large quantities of data in parallel from Teradata Database to a client. The utility is the functional complement of the FastLoad and MultiLoad utilities.

Teradata FastExport can:

- Export tables to client files.
- Export data to an Output Modification (OUTMOD) routine. You can write an OUTMOD routine to select, validate, and preprocess exported data.
- Perform block transfers with multisession parallelism.

## Teradata FastLoad

Teradata FastLoad loads data into unpopulated tables only. Both the client and server environments support Teradata FastLoad.

Teradata FastLoad can:

- Load data into an empty table.  
FastLoad loads data into one table per job. If you want to load data into more than one table, you can submit multiple FastLoad jobs.
- Perform block transfers with multisession parallelism.

## Teradata MultiLoad

Teradata MultiLoad supports bulk inserts, updates, and deletes against initially unpopulated or populated database tables. Both the client and server environments support Teradata MultiLoad.

Teradata MultiLoad can:

- Run against multiple tables.
- Perform block transfers with multisession parallelism.
- Load data from multiple input source files.
- Pack multiple SQL statements and associated data into a request.
- Perform data Upserts.

## Teradata Parallel Data Pump

Teradata TPump uses standard SQL (not block transfers) to maintain data in tables.

TPump also contains a resource governing method whereby you can control the use of system resources by specifying how many inserts and updates occur minute-by-minute. This allows background maintenance for insert, delete, and update operations to take place at any time of day while Teradata Database is in use.

TPump provides the following capabilities:

- Has no limitations on the number of instances running concurrently.
- Uses conventional row hash locking, which provides some amount of concurrent read and write access to the target tables.
- Supports the same restart, portability, and scalability as Teradata MultiLoad.
- Perform data Upserts.

## Access Modules

Access modules are dynamically linked software components that provide block-level I/O interfaces to external data storage devices.

Access modules import data from data sources and return it to a Teradata utility. Access modules are dynamically linked to one or more Teradata utilities by the Teradata Data Connector API.

Access modules provide transparent, uniform access to various data sources, isolating you from, for example, device and data store dependencies.

Teradata Database supports the following access modules, which read data, but do not write it.

- **Named Pipes Access Module**, which enables you to load data into the Teradata Database from a UNIX OS named pipe. A pipe is a type of data buffer that certain operating systems allow applications to use for the storage of data.
- **Teradata WebSphere MQ Access Module**, which enables you to load data from a message queue using IBM's WebSphere MQ (formerly known as MQ Series) message queuing middleware.
- **Teradata Access Module for JMS**, which enables you to load data from a JMS-enabled messaging system using JMS message queuing middleware.
- **Teradata OLE DB Access Module**, which enables you to transfer data between an OLE DB provider and Teradata Database.
- **Custom Access Modules**, which you can use with the DataConnector operator for access to specific systems.

Teradata access modules work on many operating systems with the following Teradata Tools and Utilities:

- BTEQ
- Teradata FastExport
- Teradata FastLoad
- Teradata MultiLoad
- Teradata PT
- Teradata TPump

## Basic Teradata Query

Basic Teradata Query (BTEQ) software is a general-purpose, command-based program that allows users on a workstation to communicate with one or more Teradata Database systems and to format reports for both print and screen output.

Using BTEQ you can submit SQL queries to the Teradata Database. BTEQ formats the results and returns them to the screen, to a file, or to a designated printer.

A BTEQ session provides a quick and easy way to access a Teradata Database. In a BTEQ session, you can:

- Enter Teradata SQL statements to view, add, modify, and delete data.
- Enter BTEQ commands.
- Enter operating system commands.
- Create and use Teradata stored procedures.

## Session and Configuration Management Tools

Database management tools include utilities for investigating active sessions and the state of Teradata Database configuration, such as:

- Query Session
- Query Configuration
- Gateway Global

The following table contains information about the capabilities of each utility.

| This utility...     | Does the following...                                                                                                                                                                                                                                                                                                                                                                                                                                               |
|---------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Query Session       | <ul style="list-style-type: none"><li>• Provides information about active Teradata Database sessions.</li><li>• Monitors the state of all or selected sessions on selected logical host IDs attached to Teradata Database.</li><li>• Provides information about the state of each session including session details for Teradata Index Wizard. For more information about Teradata Index Wizard, see <a href="#">“Teradata Index Wizard” on page 150</a>.</li></ul> |
| Query Configuration | <p>Provides reports on the current Teradata Database configuration, including:</p> <ul style="list-style-type: none"><li>• Node</li><li>• AMP</li><li>• PE identification and status</li></ul>                                                                                                                                                                                                                                                                      |
| Gateway Global      | <p>Allows you to monitor and control the sessions of Teradata Database network-connected users.</p>                                                                                                                                                                                                                                                                                                                                                                 |

# System Resource and Workload Management Tools and Protocols

Teradata Database supports specific tools, protocols, and tool architectures for system resource and workload management. Among these are:

- Write Ahead Logging
- Ferret utility
- Priority Scheduler
- Teradata Active System Management

## Write Ahead Logging

Teradata Database uses a Write Ahead Logging (WAL) protocol. According to this protocol, writes of permanent data are written to a log file that contains the records representing updates. The log file is written to disk at key moments, such as at transaction commit.

Modification to permanent data from different transactions, all written to the WAL log, can also be batched. This achieves a significant reduction in I/O write operations. One I/O operation can represent multiple updates to permanent data.

The WAL Log is conceptually similar to a table, but the log has a simpler structure than a table. Log data is a sequence of WAL records, different from normal row structure and not accessible via SQL.

The WAL Log includes the following:

- Redo Records for updating disk blocks and insuring file system consistency during restarts, based on operations performed in cache during normal operation.
- Transient Journal (TJ) records used for transaction rollback.

WAL protects all permanent tables and all system tables, except Transient Journal (TJ) tables, user journal tables, and restartable spool tables (global temporary tables). Furthermore, WAL allows Teradata Database to be reconstructed from the WAL Log in the event of a system failure.

The file system stages in-place writes through a disk area called the DEPOT, a collection of cylinders. Staging in-place writes through the DEPOT ensures that either the old or the new copy is available after a restart.

## Ferret Utility

The Ferret utility, `ferret`, lets you display and set storage space utilization attributes of the Teradata Database. Ferret dynamically reconfigures the data in the Teradata file system while maintaining data integrity during changes. Ferret works on various data levels as necessary: `vproc`, `table`, `subtable`, `WAL log`, `disk`, and `cylinder`.

The Ferret utility allows display and modification of the WAL Log and its index.

Ferret includes the following:

- The SCANDISK option includes checking the WAL Log by default.
- The SCOPE option can be set to just the WAL Log.
- The SHOWBLOCKS and SHOWSPACE options display log statistics and space.

## Priority Scheduler

Teradata Database Priority Scheduler (also called schmon) is a workload management facility that controls access to resources by the different active jobs in Teradata Database. It allows administrators to define different priorities for different categories of work, and comes with a number of flexible options.

Priority Scheduler is active in all Teradata Database systems.

Priority Scheduler has these capabilities:

- Provides better service for more important work.
- Controls resource sharing among different applications.
- Prevents aggressive queries from over-consuming resources at the expense of other work.
- Automates changes in priority based on query or session CPU usage levels.

## Teradata Active System Management

Teradata Active System Management (Teradata ASM) is a set of products, including system tables and logs, that interact with each other and a common data source. It facilitates automation in the following four key areas of system management:

- Workload management
- Performance tuning
- Capacity planning
- Performance monitoring

With careful planning, Teradata ASM can improve and optimize your workload management and performance. It can also improve response times and ensure more consistent response times for critical work. This can reduce the effort required by DBAs.

### Some Key Teradata ASM Products and Components

The following table describes some of the key products and components of Teradata ASM. For a complete list of products and components, as well as a discussion of how they work together, see *Performance Management*.

| Product / Component                      | Description                                                                                                                                                               |
|------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Teradata Workload Analyzer (Teradata WA) | A product that analyzes collected DBQL data to help group workloads and define rules to manage system performance. See <a href="#">“Database Query Log” on page 153</a> . |



| Product / Component    | Description                                                                                                                                                                                                                                                                                                                     |
|------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Teradata Viewpoint     | A product that enables users to: <ul style="list-style-type: none"> <li>• Create rules for filtering, throttling, and defining classes of queries (workload definitions)</li> <li>• Create events to monitor system resources</li> </ul>                                                                                        |
| Open APIs              | An interface, invoked from any application, that provides an SQL interface to PMPC through User-Defined Functions (UDFs) and external stored procedures.<br>See <a href="#">“Workload Management Application Programming Interface” on page 107</a> .                                                                           |
| Query Bands            | A set of name-value pairs that are defined by the user or middle-tier application. Query Bands allow the user to tag session or transactions with an ID through an SQL interface.<br><br>For specific information on Query Bands, see <i>SQL Data Definition Language, Database Administration and Performance Management</i> . |
| Resource Usage Monitor | Data collection subsystem that includes Teradata ASM-related data collection. See <a href="#">“Resource Usage Monitoring” on page 192</a> .                                                                                                                                                                                     |

## Teradata Viewpoint

Teradata Viewpoint, including the Workload Designer portlet, supports the creation of the following based on business-driven allocations of operating resources:

- Filter rules
- Throttle rules
- Rules that define classes of queries (Workload Definitions [WDs])
- Events to monitor system resources
- States to allow changes to rule values

## Request-Specific Performance Management

Teradata Viewpoint Workload Designer portlet lets users define rules according to which workload is managed. The following table describes the three categories of Teradata Active System Management (ASM) rules.

| Rules                                                | Description                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Filter                                               | <p>Reject unwanted logon and query requests before they are executed.</p> <p>Filters restrict access to specific database objects for some or all types of SQL requests. You can prohibit queries that are estimated to access too many rows, take too long, and perform some types of joins.</p>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| Throttle<br>(also called concurrency rules)          | <p>Enforce session and query concurrency limits on specific objects.</p> <p>When creating throttle rules, you can:</p> <ul style="list-style-type: none"> <li>• Restrict the number of requests simultaneously executed against a database object (such as requests made by a user, or against a table).</li> <li>• Reject requests over the concurrency limit on a state-by-state basis (where <i>state</i> is a complete set of working values for a rule set).</li> <li>• Enforce concurrency limits on FastLoad, MultiLoad, FastExport, and ARC utilities.</li> <li>• Apply them to Priority Scheduler Performance Groups (PGs).</li> </ul>                                                                                                                                                                                                                                                                                                                              |
| Workload<br>(also called Workload Definitions [WDs]) | <p>Specify how Teradata Database should handle queries while they are executing by specifying parameters for up to 36 separate workload definitions.</p> <p>In each workload definition, you can specify:</p> <ul style="list-style-type: none"> <li>• The Include and Exclude conditions or database objects, or both the workload definition and database objects that determine whether a query is assigned to the class.</li> </ul> <p><b>Note:</b> Wildcard characters (such as “*” and “?”) can be used to include all or a group of database objects, and then exclude specific ones.</p> <ul style="list-style-type: none"> <li>• The execution priority (by more or less transparently creating a Priority Scheduler configuration).</li> <li>• The query concurrency limits. Requests over the concurrency limit can be rejected on a state-by-state basis.</li> <li>• The set of conditions that invoke an exception once a query has started running.</li> </ul> |

### **Event-Based Performance Management**

Teradata Viewpoint Workload Designer portlet allows you to specify filter, throttle, and workload rules (WDs) that dynamically adjust their behavior based on system and user-defined events.

An event is any condition or indication that you think is pertinent to workload management.

| Event               | Description                                                                                                                                                                       |
|---------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Health Condition    | Reflects the health of the system, such as a Teradata Database component degrading or failing (a node down, for example), or resources below a threshold for some period of time. |
| Planned Environment | Includes the kinds of work Teradata Database is expected to perform, such as batch and loads or month-end processing, defined as time periods.                                    |

### **Utility Management**

You can manage load utilities (FastLoad, MultiLoad, FastExport) and Teradata Archive/Recovery utility similarly to how you manage SQL requests: by classifying them into workload definitions with throttle limits based on:

- Utility name
- “Who” criteria (such as user, account, or queryband)
- “Where” criteria (the name of the database, table, or view) (not available for Archive/Recovery)

For example, you can specify that:

- User A cannot run more than two FastLoad jobs at the same time
- Only one MultiLoad job can run at a time against Database XYZ

In addition, using Workload Designer, a portlet of Viewpoint, you can define rules that control the number of sessions that Archive/Recovery or a load utility can use. To create session configuration rules, specify at least one of the following criteria:

- Utility name (required)
- Data size
- “Who” criteria (for example, user, account, client address, or query band)

and then specify the number of sessions to be used when the criteria are met.

For example, you can specify using:

- Ten sessions for standalone MultiLoad jobs submitted by Joe
- Four sessions for JDBC FastLoad from the application called WebApp

If you do not create a session configuration rule or the utility currently running does not meet the rule criteria, then Workload Designer automatically uses default session rules to select the number of sessions based on these criteria:

- Utility name
- System configuration (the number of AMPs)
- Optional data size

**Teradata Workload Analyzer**

Teradata Workload Analyzer (WA) is a tool to analyze resource usage patterns of workloads and to correlate these patterns with specific classification and exception criteria. Teradata WA helps:

- DBAs identify classes of queries (workloads).
- Provides recommendations on workload definitions and operating rules to ensure that database performance meets Service Level Goals (SLGs).

Through graphical displays such as pie charts showing CPU and I/O utilization by accounts, applications, users, profiles, or query bands, through histograms showing actual service levels, as well as recommended settings, Teradata WA makes it easier for DBAs to manage distribution of resources effectively.

**Teradata SQL Assistant**

Teradata SQL Assistant provides information discovery capabilities on Windows-based systems. Teradata SQL Assistant retrieves data from any ODBC-compliant database server and allows you to manipulate and store the data on your desktop PC. You can then use this data to produce consolidated results or perform analyses on the data using tools such as Microsoft Excel. The following table contains information about key features of Teradata SQL Assistant.

| This feature... | Allows you to...                                                                                                                                                                                                                                                                               |
|-----------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Queries         | <ul style="list-style-type: none"><li>• Use SQL syntax examples to help compose your SQL statements.</li><li>• Send statements to any ODBC database or the same statement to many different databases.</li><li>• Limit data returned to prevent runaway execution of statements.</li></ul>     |
| Reports         | <ul style="list-style-type: none"><li>• Create reports from any database that provides an ODBC interface.</li><li>• Use an imported file to create many similar reports (query results or answer sets); for example, display the DDL (SQL) that was used to create a list of tables.</li></ul> |

| This feature...                     | Allows you to...                                                                                                                                                                                                                                                                                                                                                            |
|-------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Data manipulation                   | <ul style="list-style-type: none"> <li>• Export data from the database to a file on a PC.</li> <li>• Import data from a PC file directly to the database.</li> <li>• Create a historical record of the submitted SQL with timings and status information, such as success or failure.</li> <li>• Use the Database Explorer Tree to easily view database objects.</li> </ul> |
| Teradata Database stored procedures | Use a procedure builder that gives you a list of valid statements for building the logic of a stored procedure, using Teradata Database syntax.                                                                                                                                                                                                                             |

Teradata SQL Assistant electronically records your SQL activities with data source identification, timings, row counts, and notes. Having this historical data allows you to build a script of the SQL that produced the data. The script is useful for data mining.

## For More Information

For more information on the topics presented in this chapter, see the following Teradata Database and Teradata Tools and Utilities books.

| IF you want to learn more about...                                                                                                                                                                    | See...                                                                                                                                                                                                                                        |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Data Archiving Utilities, including: <ul style="list-style-type: none"> <li>• Archive/Recovery utility</li> <li>• Backup and Restore (BAR), including Backup Application Software products</li> </ul> | <ul style="list-style-type: none"> <li>• <i>Teradata Archive/Recovery Utility Reference</i></li> <li>• <i>Teradata BAR Solutions Guide</i></li> </ul>                                                                                         |
| Teradata Parallel Transporter                                                                                                                                                                         | <ul style="list-style-type: none"> <li>• <i>Teradata Parallel Transporter User Guide</i></li> <li>• <i>Teradata Parallel Transporter Reference</i></li> <li>• <i>Teradata Parallel Transporter Operator Programmer Guide</i></li> </ul>       |
| Teradata Parallel Transporter Application Programming Interface                                                                                                                                       | <i>Teradata Parallel Transporter Application Programming Interface Programmer Guide</i>                                                                                                                                                       |
| Standalone Data Load and Export Utilities                                                                                                                                                             | <ul style="list-style-type: none"> <li>• <i>Teradata FastExport Reference</i></li> <li>• <i>Teradata FastLoad Reference</i></li> <li>• <i>Teradata MultiLoad Reference</i></li> <li>• <i>Teradata Parallel Data Pump Reference</i></li> </ul> |
| Access Modules                                                                                                                                                                                        | <i>Teradata Tools and Utilities Access Module Programmer Guide</i>                                                                                                                                                                            |
| Basic Teradata Query (BTEQ)                                                                                                                                                                           | <i>Basic Teradata Query Reference</i>                                                                                                                                                                                                         |
| Session and Configuration Management                                                                                                                                                                  | <i>Utilities</i>                                                                                                                                                                                                                              |

| IF you want to learn more about...                                                                                                                                                                                                      | See...                                                                                                                                                                                                                                        |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| System Resource and Workload Management, including: <ul style="list-style-type: none"><li>• Write Ahead Logging (WAL)</li><li>• Ferret utility</li><li>• Priority Scheduler</li><li>• Teradata Active System Management (ASM)</li></ul> | <ul style="list-style-type: none"><li>• <i>Support Utilities</i></li><li>• <i>Utilities</i></li><li>• <i>Database Administration</i></li><li>• <i>Performance Management</i></li><li>• <i>Teradata Workload Analyzer User Guide</i></li></ul> |
| Teradata SQL Assistant                                                                                                                                                                                                                  | <ul style="list-style-type: none"><li>• <i>Teradata SQL Assistant for Microsoft Windows User Guide</i></li></ul>                                                                                                                              |

## CHAPTER 18 Aspects of System Monitoring

This chapter discusses various aspects of monitoring Teradata Database, including the monitoring tools used to track system and performance issues.

### Teradata Viewpoint

Teradata Viewpoint enables database and system administrators and business users to monitor and manage Teradata Database systems from anywhere using a standard web browser.

Teradata Viewpoint allows users to view system information, such as query progress, performance data, and system saturation and health through preconfigured portlets displayed from within the Teradata Viewpoint portal. Portlets can also be customized to suit individual user needs. User access to portlets is managed on a per-role basis.

Database administrators can use Teradata Viewpoint to determine system status, trends, and individual query status. By observing trends in system usage, system administrators are better able to plan project implementations, batch jobs, and maintenance to avoid peak periods of use. Business users can use Teradata Viewpoint to quickly access the status of reports and queries and drill down into details. Teradata Viewpoint includes:

| Component               | Description                                                                                                                                                                                                                                                                                                        |
|-------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Viewpoint Server        | The hardware on which the other components are installed. It is an Ethernet-connected server housed in the Teradata rack so it can be managed by server management software.                                                                                                                                       |
| Viewpoint Portal        | A Java-based portal that resides on the Viewpoint Server and is the framework for delivering the Teradata Viewpoint portlets. The portal includes built-in security, user management, and role-based permissions for defining which users can perform specific actions and for granting or limiting system access. |
| Data Collection Service | A Java process that tracks daily activities on Teradata Database and maintains data history for comparison and reporting purposes. Because it maintains a local cache database, users can access the system even during critical events, such as outages.                                                          |
| Portlets                | Preconfigured and customizable content that provides a current view of the workload and throughput of Teradata systems and a historical view of system capacity and use.                                                                                                                                           |

For more information on Teradata Viewpoint, contact your Teradata representative.

## QUERY STATE Command

By definition, Teradata Database is always in one of several states. You can monitor these states using the QUERY STATE command from the DBW Supervisor window.

The following table lists the valid system states.

| System State                                                 | Description                                                                                                                                                                                           |
|--------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Database is not running                                      | Teradata Database has not been started; it cannot be accessed from a client or used for processing.                                                                                                   |
| Database Startup                                             | Teradata Database is undergoing startup processing and is not yet ready to accept requests.                                                                                                           |
| Logons are disabled - Users are logged on                    | No new sessions can log on, but existing sessions are still logged on.                                                                                                                                |
| Logons are disabled - The system is quiescent                | Logons are disabled and no sessions are logged on.                                                                                                                                                    |
| Logons are enabled - Users are logged on                     | New sessions can log on and work is in process.                                                                                                                                                       |
| Logons are enabled - The system is quiescent                 | Logons are enabled, but no sessions are logged on.                                                                                                                                                    |
| Only user DBC Logons are enabled                             | Only new DBC sessions can log on and work is in process.                                                                                                                                              |
| RECONFIG is running                                          | Reconfiguration is being run.                                                                                                                                                                         |
| System is operational without PEs - Sessions are not allowed | Either there are no PEs configured into the system or all PEs are offline/down.                                                                                                                       |
| TABLEINIT is running                                         | Database startup has detected that there are no tables on the system and is running TABLEINIT to create the system tables.<br><br>This usually occurs during the next system restart after a SYSINIT. |

## Resource Usage Monitoring

Teradata Database has facilities that permit you to monitor the use of resources such as:

- AMPs
- AWTs
- CPUs
- BYNET activity
- Disk activity
- Performance Groups



Resource usage data is useful for the following purposes:

- Measuring system performance
- Measuring component performance
- Assisting with on-site job scheduling
- Identifying potential performance impacts
- Planning installation, upgrade, and migration
- Analyzing performance degradation and improvement
- Identifying problems such as bottlenecks, parallel inefficiencies, down components, and congestion

## Resource Usage Data Gathering

Resource usage data gathering is a two-phase process as follows:

- Data collection
- Data logging

Two Teradata Database subsystems work in conjunction with other subsystems to gather resource usage data. These are:

- Parallel Database Extensions (PDE)
- Resource Sampling Subsystem (RSS)

During the data collection phase, both PDE and RSS gather information from the operating system and from Teradata Database. Data collection continues for a user-specified period of time, called the collection interval.

In the data logging phase, which occurs at the end of each logging period and consists of one or more complete collection intervals, the RSS writes all gathered data to resource usage (ResUsage) tables and re-initializes the shared data collection buffers for the next log interval.

## How to Control Collection and Logging of ResUsage Data

Several mechanisms exist within Teradata Database for setting the collection and logging rates of ResUsage data. The control sets allow users to do any of the following:

- Specify data collection rate.
- Specify data logging rate.
- Enable or disable ResUsage data logging on a table-by-table basis.
- Enable or disable summarization of the data.
- Enable or disable active row filtering.

Collection rates control the frequency that resource usage data is made available to applications. Logging rates control the frequency that resource usage data is logged to the ResUsage tables.

You can specify data collection without specifying logging. This capability saves space in system tables while making resource usage data available.

You can use the SET LOGTABLE command to establish the logging of resource usage information. The system inserts data into ResUsage tables every logging period for the tables that have logging enabled. You can use the statistics collected in the ResUsage tables to analyze system bottlenecks, determine excessive swapping, and detect system load imbalances.

## ResUsage Tables and Views

Resource usage data is stored in system tables and views in the DBC database. Macros installed with Teradata Database generate reports that display the data.

You can use SQL to access resource usage data if you have the proper privileges. You can also write your own queries or macros on resource usage data.

## ResUsage Data Categories

Each row of ResUsage data contains two broad categories of information:

- Housekeeping, containing identifying information
- Statistical

Each item of statistical data falls into a defined kind and class. Each kind corresponds to one (or several) different things that may be measured about a resource.

## ResUsage Macros

The facilities for analyzing ResUsage data are provided by means of a set of ResUsage macros tailored to retrieving information from a set of system views designed to collect and present ResUsage information.

## Summary Mode

Summary Mode is one method for improving system performance. It reduces database I/O by consolidating and summarizing data on each node on the system. Because Summary Mode reduces the availability of detail, you can log normally to tables in which greater detail is needed. Activate Summary Mode individually for the tables in which great detail is not needed.

For more information about Summary Mode and those tables supported in Summary Mode, see *Resource Usage Macros and Tables*.

# Performance Monitoring

Several facilities exist for monitoring and controlling system performance.

## Account String Expansion

Account String Expansion (ASE) is a mechanism that enables AMP usage and I/O statistics to be collected. ASE supports performance monitoring for an account string.

The system stores the accumulated statistics for a user/account string pair as a row in DBC.Acctg table in the Data Dictionary. You can use the DBC.AMPUsageV view to access this information.

Each user/account string pair results in a new set of statistics and an additional row. You can use this information in capacity planning or in charge back and accounting software.

At the finest granularity, ASE can generate a summary row for each SQL request. You can also direct ASE to generate a row for each user, each session, or for an aggregation of the daily activity for a user.

ASE permits you to use substitution variables to include date and time information in the account id portion of a user logon string. The system inserts actual values for the variables at Teradata SQL execution time.

## The TDPTMON

Teradata Director Program (TDP) User Transaction Monitor (TDPTMON) is a client routine that enables a system programmer to write code to track TDP elapsed time statistics.

## System Management Facility

The System Management Facility (SMF) is available in the z/OS environment only. This facility collects data about Teradata Database performance, accounting, and usage.

Data is grouped into the following categories:

- Session information
- Security violations
- PE stops

## Workload Management Application Programming Interface

For information on the Workload Management API, including those for Teradata Dynamic Workload Management and Query Banding, see [“Workload Management Application Programming Interface” on page 107](#).

## For More Information

For more information on the topics presented in this chapter, see the following Teradata Database and Teradata Tools and Utilities books.

| IF you want to learn more about... | See...                                       |
|------------------------------------|----------------------------------------------|
| QUERY STATE Command                | “Database Window (xdbw)” in <i>Utilities</i> |
| Resource Usage Monitoring          | <i>Resource Usage Macros and Tables</i>      |

| IF you want to learn more about...                                                                                                                                                                                                       | See...                                                                                                                                                                                 |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Performance Monitoring, including: <ul style="list-style-type: none"><li>• Account String Expansion</li><li>• the TDPTMON</li><li>• System Management Facility</li><li>• Workload Management Application Programming Interface</li></ul> | <ul style="list-style-type: none"><li>• <i>Performance Management</i></li><li>• <i>Database Administration</i></li><li>• <i>Workload Management API: PM/API and Open API</i></li></ul> |

## CHAPTER 19 **Teradata Meta Data Services**

---

Teradata Meta Data Services (MDS) provides a means of storing, administering, and navigating metadata in Teradata. It is the only metadata management system optimized for and integrated with Teradata Database environment.

### **About Metadata**

Metadata is the term applied to the definitions of the data stored in Teradata Database. Simply put, metadata is data about data. In a transaction processing database environment, a Data Dictionary generally satisfies the need for data about data. In the data warehouse environment, the requirements for a more elaborate metadata storage system can exceed the capabilities of the Data Dictionary.

Metadata plays an important role across Teradata Database architecture. In the operational database environment, that role is very formal. All development should use metadata as a standard part of the design and development process. As far as the data warehouse is concerned, metadata is used to locate data. Without it, you cannot not interact with the data in the data warehouse because you have no means of knowing how the tables are structured, what the precise definitions of the data are, or where the data originated.

### **Types of Metadata**

Metadata has been around for as long as there have been programs and data. However, in the world of data warehouses, metadata takes on a new level of importance. Using metadata, you can make the most effective use of Teradata. Metadata allows the decision support system (DSS) analyst to navigate through the possibilities.

The major component of the DSS environment is archival data, that is, data with a timestamp. Because archival data is timestamped, it makes sense to store metadata with the actual occurrences of data, which are time stamped as well.

The following table describes the types of metadata.

| For the...      | The following types of metadata are stored...                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
|-----------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| data model      | <ul style="list-style-type: none"><li>• Description</li><li>• Specification</li><li>• The layout of the physical data model tables</li><li>• Relation between the data model and the data warehouse</li></ul>                                                                                                                                                                                                                                                                                                      |
| data warehouse  | <ul style="list-style-type: none"><li>• Data source (system of record)</li><li>• Definition of the system of record</li><li>• Mapping from system of record to the data warehouse and other places defined in the environment</li><li>• Table structures and attributes</li><li>• Any relationship or artifacts of relationships</li><li>• Transformation of data as it passes into the data warehouse</li><li>• History of extracts</li><li>• Extract logging</li><li>• Common routines for data access</li></ul> |
| columns         | <ul style="list-style-type: none"><li>• Columns in a row</li><li>• Order in which the columns appear</li><li>• Physical structure of the columns</li><li>• Any variable-length columns</li><li>• Any columns with NULL values</li><li>• Unit of measure of any numeric columns</li><li>• Any encoding used</li></ul>                                                                                                                                                                                               |
| database design | <ul style="list-style-type: none"><li>• Description of the layouts used</li><li>• Structure of data as known to the programmers and analysts</li></ul>                                                                                                                                                                                                                                                                                                                                                             |

## Teradata Meta Data Services

Teradata MDS is software that creates a repository in a Teradata Database in which metadata is stored. Teradata MDS also permits the DSS analyst to administer and navigate metadata in the warehouse. The following table describes the benefits of Teradata MDS to several user groups.

| For this type of user... | Teradata MDS...                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
|--------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| application developers   | <ul style="list-style-type: none"> <li>• Provides a persistent store for application metadata so that developers can concentrate on developing application functions.</li> <li>• Allows the developer to manipulate metadata with the same techniques used to manipulate other data.</li> <li>• Provides security (MDS controls the read and write access).</li> <li>• Allows metadata to be shared between applications. This allows integration of tools such as ordered analytical functions and data mining tools.</li> <li>• Allows application data to be modeled around Teradata Database metadata maintained by MDS. MDS maintains the metadata so that the application is kept current with database changes.</li> </ul> |
| database administrator   | <ul style="list-style-type: none"> <li>• Provides a common repository for Teradata components.</li> <li>• Provides a single shared copy of metadata, or a single version of the business. One copy eliminates multiple islands of redundant metadata that can cause confusion and administrative difficulties.</li> <li>• Provides the capabilities to browse through data in the repository and to drill-down to see successive levels of detail.</li> <li>• Shows interrelationships between different data definitions.</li> <li>• Provides impact analysis of proposed changes.</li> </ul>                                                                                                                                    |
| business user            | <ul style="list-style-type: none"> <li>• Provides the foundation for a “warehouse view” of enterprise computing.</li> <li>• Allows business analysts to quickly determine where their data comes from, how it was changed, when it was last updated, and how the answer was determined. This greatly increases the value of the detail data and implicitly the value of the metadata.</li> <li>• Supports third-party tools that can be used to import metadata into MDS for viewing.</li> <li>• Supports a web browser that provides general reporting and search capabilities and shows strategic metadata relationships.</li> </ul>                                                                                            |

## Creating Teradata Meta Data Repository

Teradata MDS repository is a set of tables, views, and macros stored in a Teradata Database. You must use MDS program software to create these tables before metadata can be added, stored, or accessed.

## Connecting to Teradata Meta Data Repository

Each system running a Teradata MDS application must have the following:

- The appropriate Teradata ODBC driver (see *ODBC Driver for Teradata User Guide* for more information).
- An ODBC System Data Source Name (DSN) connection to Teradata Database where the MDS repository resides.

# For More Information

For more information on the topics presented in this chapter, see the following Teradata Meta Data Services and Teradata Tools and Utilities books.

| IF you want to learn more about... | See...                                                                                                                                                                 |
|------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Metadata                           | <ul style="list-style-type: none"><li>• <i>Teradata Meta Data Services Administrator Guide</i></li><li>• <i>Teradata Meta Data Services Programmer Guide</i></li></ul> |
| Types of Metadata                  |                                                                                                                                                                        |
| Teradata MDS                       |                                                                                                                                                                        |



# Glossary

|              |                                                    |
|--------------|----------------------------------------------------|
| <b>1NF</b>   | First Normal Form                                  |
| <b>2NF</b>   | Second Normal Form                                 |
| <b>2PC</b>   | Two-Phase Commit                                   |
| <b>3NF</b>   | Third Normal Form                                  |
| <b>4NF</b>   | Fourth Normal Form                                 |
| <b>5NF</b>   | Fifth Normal Form                                  |
| <b>AMP</b>   | Access Module Processor                            |
| <b>ANSI</b>  | American National Standards Institute              |
| <b>API</b>   | Application Programming Interface                  |
| <b>ARC</b>   | Teradata Archive/Recovery Utility                  |
| <b>ASCII</b> | American Standard Code for Information Interchange |
| <b>ASE</b>   | Account String Expansion                           |
| <b>AWS</b>   | Administration Workstation                         |
| <b>AWT</b>   | AMP Worker Task                                    |
| <b>BCNF</b>  | Boyce-Codd Normal Form                             |
| <b>BTEQ</b>  | Basic Teradata Query                               |
| <b>BYNET</b> | Banyan Network (high-speed interconnect)           |
| <b>CAS</b>   | Channel-Attached System                            |
| <b>CICS</b>  | Customer Information Control System                |
| <b>CLiV2</b> | Call-Level Interface, Version 2                    |
| <b>CNS</b>   | Console Subsystem                                  |
| <b>DB2</b>   | DATABASE 2                                         |
| <b>DBC</b>   | Database Computer                                  |
| <b>DBQAT</b> | Database Query Analysis Tools                      |
| <b>DBQL</b>  | Database Query Log                                 |
| <b>DBS</b>   | Database System or Database Software               |

|                    |                                                     |
|--------------------|-----------------------------------------------------|
| <b>DDE</b>         | Dynamic Data Exchange                               |
| <b>DDL</b>         | Data Definition Language                            |
| <b>DIP</b>         | Database Initialization Program                     |
| <b>DML</b>         | Data Manipulation Language                          |
| <b>DNS</b>         | Domain Name Source                                  |
| <b>DSS</b>         | Decision Support System                             |
| <b>EBCDIC</b>      | Extended Binary Coded Decimal Interchange Code      |
| <b>FIPS</b>        | Federal Information Processing Standards            |
| <b>GDO</b>         | Globally Distributed Object                         |
| <b>HI</b>          | Hash Index                                          |
| <b>IBM</b>         | International Business Machines Corporation         |
| <b>ID</b>          | Identification                                      |
| <b>IMS</b>         | Information Management System                       |
| <b>I/O</b>         | Input/Output                                        |
| <b>ISV</b>         | Independent Software Vender                         |
| <b>JBOD</b>        | Just a Bunch Of Disks                               |
| <b>JDBC</b>        | Java Database Connectivity                          |
| <b>JI</b>          | Join Index                                          |
| <b>LAN</b>         | Local Area Network                                  |
| <b>LUN</b>         | Logical Unit                                        |
| <b>MDS</b>         | Meta Data Services                                  |
| <b>MIPS</b>        | Millions of Instructions Per Second                 |
| <b>MLPPI</b>       | Multilevel Partitioned Primary Index                |
| <b>MOSI</b>        | Micro Operating System Interface                    |
| <b>MPP</b>         | Massively Parallel Processing                       |
| <b>MTDP</b>        | Micro Teradata Director Program                     |
| <b>NAS</b>         | Network-Attached System                             |
| <b>NoPI Tables</b> | Tables that are defined with no primary index (PI). |
| <b>NPPI</b>        | Nonpartitioned Primary Index                        |

|                     |                                                   |
|---------------------|---------------------------------------------------|
| <b>NUPI</b>         | Nonunique Primary Index                           |
| <b>NUSI</b>         | Nonunique Secondary Index                         |
| <b>ODBC</b>         | Open Database Connectivity                        |
| <b>OS/VS</b>        | Operating System/Virtual Storage                  |
| <b>PDE</b>          | Parallel Database Extensions                      |
| <b>PE</b>           | Parsing Engine                                    |
| <b>PI</b>           | Primary Index                                     |
| <b>PL/I</b>         | Programming Language 1                            |
| <b>PJ/NF</b>        | Projection-Join Normal Form                       |
| <b>PP2</b>          | Preprocessor2                                     |
| <b>PPI</b>          | Partitioned Primary Index                         |
| <b>PUT</b>          | Parallel Upgrade Tool                             |
| <b>QCD</b>          | Query Capture Database                            |
| <b>QCF</b>          | Query Capture Facility                            |
| <b>RAID</b>         | Redundant Array of Independent Disks              |
| <b>RCT</b>          | Resource Check Tools                              |
| <b>RI</b>           | Referential Integrity                             |
| <b>SIA</b>          | Shared Information Architecture                   |
| <b>SMP</b>          | Symmetric Multiprocessing                         |
| <b>SNMP</b>         | Simple Network Management Protocol                |
| <b>SR</b>           | Single Request                                    |
| <b>SSO</b>          | Single Sign On                                    |
| <b>TCP/IP</b>       | Transmission Control Protocol/Internet Protocol   |
| <b>TDGSS</b>        | Teradata Generic Security Services                |
| <b>TDP</b>          | Teradata Director Program                         |
| <b>TDSP</b>         | Teradata Stored Procedures                        |
| <b>Teradata ASM</b> | Teradata Active System Management                 |
| <b>TPA</b>          | Trusted Parallel Application                      |
| <b>TS/API</b>       | Transparency Series/Application Program Interface |

**UPI** Unique Primary Index

**USI** Unique Secondary Index

**vproc** Virtual Processor

**z/OS** IBM System z Operating System

## Symbols

.NET Data Provider 29  
.NET Data Provider for Teradata 66

## Numerics

1NF, first normal form 121  
2NF, second normal form 121  
2PC 133  
2PL, phases of 126  
3NF, third normal form 121

## A

Access level lock 129  
Access logs 157, 165  
Access Module for JMS 31  
Access Module Processor 45  
Access Module Processor. See AMP  
Access modules 32, 181  
Account String Expansion 194  
Account String Expansion. See ASE  
Accounting  
    ASE 194  
    DBC.AMPUsage table 195  
Active data warehouse 19  
Active Directory 165  
Active System Management. See Teradata Active System Management  
ADAM 165  
Administrative utilities 172  
Aggregate functions 95  
Aggregate join indexes 116  
AMP  
    clusters 46, 57  
    down AMP journal 59  
    down AMP recovery 132  
    functions 45  
    hashing and 120  
    operation 48  
    SELECT statement processing 50  
    step processing 50  
    vproc migration 55  
Analytical functions. See Ordered analytical functions  
ANSI mode  
    transactions 126  
ANSI mode transactions 126

ANSI-compliant data types  
    examples of 93  
API  
    Workload Management 195  
    workload management 107  
API for Network-Attached Clients 29  
APIs 29  
APIs. See also Open APIs  
Application development  
    client applications 99  
    embedded SQL applications 99  
    macros 100  
    platforms 100  
    Preprocessor2 100  
    SQL stored procedures 102  
Application development languages  
    C 100  
    COBOL 100  
    PL/I 100  
ARC. See Teradata Archive/Recovery  
Archive utilities  
    Teradata Archive/Recovery 60, 178  
Archive/Recovery. See Teradata Archive/Recovery  
ASE, accounting 194  
Attachment Methods  
    .NET Provider for Teradata 66  
    CLIv2 68  
    ODBC 66  
    OLE DB Provider for Teradata 66  
Attachment methods  
    channel 25  
    CLIv2 66, 67  
    network 25  
    Teradata Database  
        JDBC 66  
Authentication 157  
    external 160  
    user 160  
Authorization 157  
    user 163  
Automatic privileges 158  
AWS  
    function 51  
    platform 51

**B**

Backup/Archive/Recovery. See BAR

BAR

- encryption 165

- NetBackup 177

- NetBackup Extension for Teradata 34

- NetVault 177

- Tivoli Storage Manager Teradata Extension 34, 178

Basic Teradata Query. See BTEQ

Battery backup 61

BEGIN LOGGING 166

Boardless BYNET 42

BTEQ 30, 182

BYNET

- boardless 42

- function 42

- inter-network communication 41

- multiple 61

**C**

C Preprocessor 30

C, application development language 100

Call-Level Interface version 2. See CLIV2

Channel-attached clients 162

Channel-attached systems

- mainframe 25

- multiple connections 61

- TDP 68

Character set support. See International language support

Child table 122

CICS 29

Client platforms 28

Cliques

- definition 43

- disk arrays 43

- hardware fault tolerance 62

- vproc migration 62

CLIV2 29, 30, 66, 67, 68

Clusters

- AMP 46

- fault tolerance 46

COBOL Preprocessor 30

COBOL, application development language 100

Columns

- attributes 37

- identity 120

Command-line logon 161

Communications interfaces

- TDP 68

Concurrency 125

Concurrency control

- 2PC and 133

- locks 127

- transactions 125

Confidentiality

- definition 164

Connection pools 162

Constraints

- normal forms and 120

- table 38

Constructor methods 81

CREATE PROCEDURE statement 77

CREATE USER 158

Credentials 157, 160

Cursors 96

- definition 96

- Preprocessor2 97

- SQL statements related to 96

- stored procedures 97

Customer Information Control System. See CICS

**D**

Data 135

Data attributes, summary of 94

Data Control Language. See DCL

Data Definition Language. See DDL

Data Dictionary

- compatibility system views 141

- DBC.AMPUsage table 195

- SQL statements and 141

- Unicode system views 141

- views 140

Data distribution

- hashing 120

- indexes and 111

Data encryption 164

Data export utilities

- Teradata FastExport 179

Data load utilities

- controlling the number of sessions 187

- Teradata FastLoad 180

- Teradata MultiLoad 180

- Teradata TPump 180

- throttling 187

Data management

- active sessions 182

- archive utilities 178

- Backup Application Software 177

Data Manipulation Language. See DML

Data protection 164

Data type attributes 94

Data type phrase 93

Data types 93

- data type phrase and 93

- UDT 82

- Data warehouse
  - active data warehouse 19
  - definition 19
- Database
  - management tools 33, 191
  - space allocation 82
- Database access, monitoring 165
- Database level lock 128
- Database Management and Query Analysis Tools 32, 33
- Database object use count 153
- Database recovery 132
- Database security 157
- Database window
  - server software 27
- Databases
  - database object use count 153
  - definition 82
- DBQAT
  - database object use count 153
- DCL, examples of 88
- DDL
  - CREATE PROCEDURE 77
  - REPLACE PROCEDURE 77
- DDL, examples of 87
- Deadlocks
  - resolution 130
  - transaction rollback 130
- Derived tables 74
- Directory Management of Users
  - supported directories 165
- Directory sign-on 161
- Directory-based users 158
- Disk arrays
  - LUNs 42
  - RAID 42
  - RAID1 61
- Dispatcher
  - function 47
  - operation 48
- Distinct UDTs 82
- DML, examples of 89
- Down AMP
  - journal 59
  - recovery 132

## E

- Embedded SQL applications 99
- END LOGGING 166
- Error logging tables 74
- Event log 166
- Event-based performance management 187
- Exclusive HUT lock 131
- Exclusive level lock 129

- EXPLAIN Request Modifier
  - definition 104
  - use 104
- Explicit privileges 159
- Extended language support. See International language support
- External authentication 160
- External roles 159
- External stored procedures 76, 77
  - C and C++ 76
  - CREATE PROCEDURE 77
  - Java 76
  - usage 76

## F

- Fallback table 56
- FastExport 179
  - controlling the number of sessions 187
  - using throttle limits 187
- FastExport. See Teradata FastExport
- FastLoad 180
  - controlling the number of sessions 187
  - using throttle limits 187
- FastLoad. See Teradata FastLoad
- Fault tolerance
  - AMP clusters 57
  - clusters 46
  - fallback tables 56
  - hardware 61
  - software 55
  - Table Rebuild utility 60
  - Teradata Archive/Recovery 60
  - vproc migration 55
- Ferret utility 183
- Filters 185
- Foreign key 112
  - referential integrity and 122
- Formats, logon 161
- Full table scans, strengths and weaknesses 119
- Functions
  - aggregate 95
  - definition 95
  - ordered analytical 96
  - scalar 95

## G

- Gateway Global utility 182
- Generator, function 46
- Global temporary tables 73
- Global temporary trace tables 73
- GRANT 163
- Group Read HUT lock 131
- GUI logon 162

**H**

Hardware fault tolerance  
 battery backup 61  
 cliques 62  
 hot swap 61  
 multiple BYNETS 61  
 multiple channel and network connections 61  
 redundant power supplies and fans 61  
 server isolation 61

Hash indexes 117  
 strengths and weaknesses 119

Hashing  
 data distributing 120  
 primary index 120  
 secondary index 120

Host Utility Consoles. *See* HUTCNS

Host Utility lock. *See* HUT lock

Hot standby nodes  
 definition 44  
 function 44

Hot swap  
 components 61  
 definition 61

HUT lock  
 characteristics of 131  
 Exclusive 131  
 Group Read 131  
 Read 131  
 Teradata Archive/Recovery and 130  
 Write 131

HUTCNS 29

**I**

IBM IMS 30

Identity column  
 column attribute 120  
 unique row number generator 120

Implicit privileges 158

IMS. *See* IBM IMS

Index Wizard 150

Indexes  
 hash 117  
 join 115  
 secondary 114  
 SQL statements and 117  
 strengths and weakness 118  
 types of 111  
 uses 111

Inherited privileges 158

Instance methods 81

Integrity  
 definition 164

International character set support. *See* International

language support

International language support  
 character data translation 144  
 character sets, internal 145  
 compatible languages 146  
 extended support 147  
 external character sets 144  
 internal character sets 144  
 Japanese mode 145  
 language support modes 145  
 LATIN character set 146  
 object names 145  
 session charset 144  
 standard mode 145  
 standard support mode 146  
 user data 145  
 user data, default character set 146

**J**

JDBC 29, 66

JMS Access Module 181

Join indexes  
 aggregate 116  
 covering 115  
 multitable 116  
 multitable, partially covering 115  
 partially covering 115  
 single table 115  
 sparse 116  
 strengths and weaknesses 119

Joins  
 SELECT statement and 92

Journals  
 down AMP 59  
 permanent 59  
 transient 59

**K**

Key  
 foreign 112, 122  
 parent 122  
 primary 112, 122

**L**

Load utilities. *See* Data load utilities

Locks 127  
 access 129  
 database level lock 128  
 deadlocks 130  
 exclusive 129  
 HUT 130  
 levels 128



- read 129
- row hash level lock 128
- table level locks 128
- write 129
- Logging database activity 165
- Logical Units. *See* LUNs
- Logon
  - command-line 161
  - controls 161
  - formats 161
  - GUI 162
  - logon string operands 172
  - sessions 171
- Logon controls 162
- Logons 161
- LUNs
  - RAID 42

**M**

- Macros
  - definition 77, 100
  - multi-user 77
  - processing 77
  - resource usage 194
  - single-user 77
  - SQL statements and 100, 101
  - SQL statements related to 77
  - use 101
- Maintenance utilities 172
- Maintenance utilities. *See also* Utilities
- Management software
  - server software 28
- Massively Parallel Processing. *See* MPP
- MDS 198
- Mechanism, security 157, 161
- Message integrity 157, 164
- Meta Data Services 33
- Meta Data Services. *See* MDS
- Metadata
  - definition 197
  - types of 197
- Methods
  - constructor 81
  - instance 81
- Micro Teradata Director Program 67
- Middle-tier application 162
- Middle-tier applications 164
- Monitoring database activity 166
- MOSI 68
- MPP
  - architecture 41
  - hardware platform 41
  - workstation connections 51

- MTDP
  - functions 67
- MTDP. *See* Micro Teradata Director Program
- Multilevel partitioned primary index 114
- MultiLoad 180
  - controlling the number of sessions 187
  - using throttle limits 187
- MultiLoad. *See* Teradata MultiLoad
- Multitable join indexes 116
  - strengths and weaknesses 119
- Multitable join indexes, partially covering 115

## N

- Named Pipes Access Module 31, 181
- NetBackup 177
- NetBackup Extension for Teradata 34
- NetVault 177
- Network-attached systems
  - LAN 25
  - multiple connections 61
- Nonpartitioned primary index 113
  - partitioned primary index and 114
- Nonunique primary index. *See* NUPI
- Nonunique secondary index 114
- Nonunique secondary index. *See* NUSI
- Normal forms
  - 1NF 121
  - 2NF 121
  - 3NF 121
  - definition 121
  - first 121
  - second 121
  - third 121
- Normalization
  - normal forms 120
  - purpose 120
- Novell eDirectory 165
- NUPI, strengths and weaknesses 118
- NUSI, strengths and weaknesses 119

## O

- ODBC 29, 66
- OLE DB Access Module 31, 181
- OLE DB Provider 29
- OLE DB Provider for Teradata 66
- Open APIs 185
- Optimizer
  - function 46
  - SQL request implementation 48
- Ordered analytical functions 96

**P**

Parallel Data Extensions. See PDE  
 Parallel Database Extensions 50  
 Parallel Upgrade Tool. See PUT  
 Parent key 122  
 Parent table 122  
 Parser  
   function 46  
   PE element 46  
   request processing 47  
 Parsing Engine 46  
 Parsing Engine. See PE  
 Partitioned primary index 113  
   nonpartitioned primary index and 114  
 Password  
   controls 163  
   format 163  
 PDE  
   function 50  
   MPP system enabling 50  
   server software 28  
   vprocs 50  
 PE  
   dispatcher 47  
   function 46  
   generator 46  
   migration 43  
   optimizer 46  
   parser 46  
   request processing 47  
   SELECT statement processing 50  
   session control 47  
   vproc migration 55  
 Performance management 185, 187  
 Performance monitoring 194  
 Performance monitoring. See System performance monitoring  
 Permanent database users 158  
 Permanent journals 59  
 Permanent tables 73  
 PL/I  
   application development language 100  
 PL/I Preprocessor 30  
 PM/API 107  
 Preprocessor2  
   application development 100  
   C 30  
   COBOL 30  
   cursors 97  
   PL/I 30  
 Primary index  
   hashing 120  
   multilevel partitioned 114

  nonpartitioned 113  
   partitioned 113  
   primary key 112  
   secondary index and 115  
 Primary key 112  
   first normal form 121  
   primary index 112  
   referential integrity and 122  
   second normal form 121  
   third normal form 121  
 Priority Scheduler 184  
 Privileges 157, 158  
 Processor node 41  
 Profiles 160  
   security 160  
 Proxy user 158, 164  
   roles for 159  
 PUT  
   installation and 28  
   operational modes 28

**Q**

QCD 149, 150, 152  
   Teradata Visual Explain 149  
 QCF 149, 150, 152  
 Queries  
   strategic 20  
   tactical 20  
   Teradata SQL Assistant 188  
 Query Analysis Tools. See Database Management and Query Analysis Tools  
 Query Bands 185  
 Query Capture Database 152  
 Query Capture Database. See QCD  
 Query Capture Facility 152  
 Query Capture Facility. See QCF  
 Query Configuration utility 182  
 Query Director 33  
 Query plan  
   capturing steps of 152  
 Query Scheduler 33  
 Query Session utility 182  
 QUERY STATE command 192  
 Queue tables 73

**R**

RAID  
   LUNs 42  
   RAID1 61  
   storage technology 42  
 Read HUT lock 131  
 Read level lock 129

- Recovery
    - database 132
    - definition 131
    - down AMP 132
    - single transaction 132
    - transaction 132
  - Recursive query 94
  - Referenced table (parent) 122
  - Referencing table (child) 122
  - Referential integrity
    - benefits of 123
    - referenced tables 122
    - referencing tables 122
    - terminology 122
  - Referential integrity terminology
    - child table 122
    - foreign key 122
    - parent key 122
    - parent table 122
    - primary key 122
  - Relational database
    - definition 37
    - relational model and 37
    - set theory and 37
    - set theory terminology 37
  - Relational model
    - relational databases and 37
    - set theory and 37
  - REPLACE PROCEDURE statement 77
  - Request processing 47
  - Resource Usage Monitor 185
  - Resource Usage. See ResUsage
  - Restarts, system 132
  - ResUsage 192
    - categories of data 194
    - collection rate control 193
    - macros 194
    - monitoring 192
  - Roles 159
    - for directory users (See External roles) 159
    - security 159
  - Row hash level lock 128
  - Rows
    - definition of 38
    - row hash lock 128
- S**
- SASL protection 165
  - Scalar functions 95
  - Secondary index 114
    - hashing 120
    - nonunique 114
    - primary index and 115
    - subtables 114
    - unique 114
  - Security
    - concepts 157
    - threats 165
  - Security policy 167
  - SELECT statement
    - cursor declaration 96
    - joins and 92
    - options 92
    - processing 49
    - set operators 92
  - Session charset 144
  - Session control
    - function 47
  - Session management 171
  - Sessions
    - establishing 171
    - logon 171
  - Set operators
    - SELECT statement and 92
  - Set theory
    - relational databases and 37
    - relational model and 37
  - Set theory terminology
    - relation 37
    - tuple 37
  - Sign-on As 161
  - Single sign-on 161
  - Single transaction recovery 132
  - Single-table join indexes 115
    - strengths and weaknesses 119
  - SMP
    - architecture 41
    - boardless BYNET 42
    - hardware platform 41
    - workstation connections 51
  - Space allocation 82
    - databases 82
  - Sparse join indexes 116
    - strengths and weaknesses 119
  - SQL
    - aggregate function 95
    - cursors 96
    - data control language statements 88
    - data definition language statements 87
    - Data Dictionary statements 141
    - data manipulation language statements 89
    - data types 93
    - embedded 99
    - EXPLAIN 104
    - indexes 117
    - macros 100, 101
    - non-ANSI compliant development 26

- ordered analytical function 96
- relational databases and 87
- scalar functions 95
- SELECT statement processing 49
- statement execution 91
- statement punctuation 91
- statement syntax 90
- statements related to macros 77
- statements, types of 87
- stored procedure statements 103
- stored procedures 75, 102
- Teradata SQL 26
- transaction statements 126, 127
- UDFs 79
- SQL Assistant 33, 188
- SQL Stored procedures
  - SQL statements and 102
- SSL protection 165
- Statistics Wizard 151
- Storage management utilities 34
- Stored procedures
  - benefits 75
  - cursors 97
  - definition 75
  - elements 76
  - external 77
  - SQL statements and 103
  - use 75
- Strategic queries 20
- Structured UDTs 82
- Subtables, secondary index and 114
- Sun Java System Directory Server 165
- Symmetric Multi-Processing. See SMP
- System administration
  - performance monitoring 194
  - session management 171
  - space allocation 82
  - utilities 172
- System console
  - function 51
  - platform 51
- System Emulation Tool 150
- System Emulation Tool. See Teradata SET
- System Management Facility 195
- System monitoring
  - resource usage 192
  - system status 192
  - Teradata Database Window 192
- System performance monitoring
  - performance monitoring 194
  - system management facility 195
  - TDPTMON 195
- System resource management 183
  - Priority Scheduler 184

- Teradata ASM 184
- Teradata WA 188
- Viewpoint 185
- System restarts 132
- System status
  - configuration 192
  - states 192
- System views, security-related 167

## T

- Table
  - derived 74
  - error logging 74
  - global temporary 73
  - global temporary trace 73
  - permanent 73
  - queue 73
  - types 73, 74
  - volatile 74
- Table level lock 128
- Table Rebuild utility 60
- Tables
  - child 122
  - constraints 38
  - DBC.AMPUsage 195
  - definition of 38
  - fallback 56
  - locks 128
  - parent 122
  - referenced table (parent) 122
  - referencing (child) 122
  - relations 37
- Tactical queries 20
- Target Level Emulation. See TLE
- TDP 30
  - channel-attached systems 68
  - definition 68
  - functions 68
- tdpid 162
- TDPTMON 195
- Teradata
  - interface to 182
- Teradata Access Modules 31
- Teradata active data warehouse
  - active access 21
  - active availability 22
  - active enterprise integration 21
  - active events 21
  - active load 20
  - active workload management 21
- Teradata Active System Management 32, 184
- Teradata Administrator 33
- Teradata Analyst Pack 32

- Teradata Index Wizard 149
- Teradata SET 149
- Teradata Statistics Wizard 149
- Teradata Visual Explain 149
- Teradata architecture
  - BYNET 41
  - cliques 43
  - disk arrays 42
  - hot standby nodes 44
  - MPP 41
  - processor node 41
  - SMP 41
  - vprocs 44
  - workstations 51
- Teradata Archive/Recovery 34
  - controlling the number of sessions 187
  - fault tolerance 60
  - HUT locks 130
  - use 178
  - using throttle limits 187
- Teradata ASM 21, 184
- Teradata Call-Level Interface, Version 2. See CLIV2
- Teradata Data Connector 29, 32
- Teradata data types
  - examples of 93
- Teradata Database
  - ANSI SQL 26
  - ANSI transaction semantics 126
  - architecture
    - hardware 41
    - software 41
  - as single data store 26
  - attachment methods 65
    - .NET Provider for Teradata 66
    - CLIV2 66, 67, 68
    - JDBC 66
    - network 65
    - ODBC 66
    - OLE DB Provider for Teradata 66
  - capabilities 27
  - character support 26
  - CLIV2 68
  - database window 27
  - designing 25
  - management software 28
  - methods of attachment 25
  - PDE 28
  - PUT installation software 28
  - referential integrity 122
  - status 192
  - Teradata Gateway 27
  - Teradata mode transactions 127
  - Teradata SQL 26
  - Teradata transaction semantics 126
  - third-party software 107
  - Teradata Database Query Analysis Tools. See DBQAT
  - Teradata Database Window 51
  - Teradata Director Program User Transaction Monitor. See TDPTMON
  - Teradata Director Program. See TDP
  - Teradata FastExport 30, 179
  - Teradata FastLoad 31, 180
  - Teradata file system
    - function 51
  - Teradata Gateway
    - server software 27
  - Teradata Index Wizard 32
    - demographics 151
    - Teradata Visual Explain and 150
    - use 150
  - Teradata mode
    - transactions 127
  - Teradata MultiLoad 31, 180
  - Teradata Parallel Transporter Application Programming Interface. See TPTAPI
  - Teradata Parallel Transporter. See Teradata PT
  - Teradata PT 31
  - Teradata PTAPI 31
  - Teradata Query Scheduler 33
  - Teradata recursive query 94
  - Teradata security 160
    - external authentication 160
    - logons 161
    - policy, defining 167
    - policy, publishing 167
    - roles 159
    - user authentication 160
    - user authorization 163
  - Teradata SET 33
    - client support for 150
    - TLE and 150
    - use 150
  - Teradata solution
    - operational intelligence 20
    - strategic intelligence 20
  - Teradata SQL Assistant 33, 188
  - Teradata Statistics Wizard 33
    - statistics collection 151
  - Teradata System Emulation Tool. See Teradata SET
  - Teradata Tools and Utilities
    - .NET Data Provider 29
    - API for network-attached clients 29
    - BTEQ 30
    - C Preprocessor 30
    - channel-attached clients 29
    - CICS 29
    - client platforms 28
    - CLIV2 29, 30

- COBOL Preprocessor 30
- Database Management and Query Analysis Tools 32, 33
- defined 28
- Host Utility Console 29
- IBM IMS 30
- installation guides for 29
- JDBC 29
- Load and export utilities 30
- ODBC 29
- OLE DB Provider 29
- PL/I Preprocessor 30
- storage management utilities 34
- TDP 30
- Teradata Access Modules 31
- Teradata Active System Management 32
- Teradata Administrator 33
- Teradata Archive/Recovery 34, 60
- Teradata Data Connector 29
- Teradata FastExport 30
- Teradata FastLoad 31
- Teradata Index Wizard 32
- Teradata MultiLoad 31
- Teradata Parallel Transporter 31
- Teradata Parallel Transporter Application Programming Interface 31
- Teradata Query Scheduler 33
- Teradata SQL Assistant 33
- Teradata Statistics Wizard 33
- Teradata System Emulation Tool 33
- Teradata TPump 31
- Teradata Visual Explain 33
- Teradata Workload Analyzer 32
- Teradata TPump 31, 180
- Teradata VE 33, 149
  - QCD 149
  - Teradata Index Wizard and 150
- Teradata Viewpoint 33, 191
- Teradata Visual Explain. See Teradata VE
- Teradata WA 32, 184
- Teradata Workload Analyzer. See Teradata WA
- Third-party software
  - Teradata Database, compatible with 107
  - TS/API products 107
- Throttles 185
- Tivoli Storage Manager Teradata Extension 34, 178
- TLE
  - Teradata SET and 152
  - use 152
- TLS protection 165
- TPump 180
- TPump. See Teradata TPump
- Transaction
  - consistency 133
- Transaction recovery, definition 131

- Transactions
  - 2PL 125
  - ANSI mode 126
  - ANSI mode, rollback 126
  - control using 2PL 125
  - deadlock resolution 130
  - definition 125
  - recovery 132
  - semantics 126
  - serializability 125
  - SQL statements and 126, 127
  - Teradata mode 127
  - Teradata mode, rollback 127
- Transient journals 59
- Triggers
  - definition 78
  - firing 78
  - types of 78
  - use 79
- TS/API
  - third-party product 107
- Two-phase commit protocol 133
- Two-Phase Locking protocol 125
- Two-Phase Locking. See 2PL

## U

- UDFs
  - aggregate function 80
  - scalar function 80
  - SQL 79
  - table function 80
  - types 80
  - usage 80
- UDM
  - constructor methods 81
  - instance methods 81
- UDT
  - examples of 93
- UDT data types 82
  - distinct 82
  - dynamic 82
  - structured 82
- Unique primary index. See UPI
- Unique secondary index 114
- Unique secondary index. See USI
- Unload utilities. See Data export utilities
- UPI, strengths and weaknesses 118
- User authentication 160
- User authorization 163
- User-defined functions. See UDFs
- User-Defined Method. see UDM
- Users 157
  - definition 82

- space allocation 82
- types 157
- USI, strengths and weaknesses 118
- Utilities
  - administrative 172
  - Backup Application Software 177
  - maintenance 172
  - Teradata Archive/Recovery 178
  - Teradata FastExport 179
  - Teradata FastLoad 180
  - Teradata MultiLoad 180
  - Teradata TPump 180
- Utility
  - Ferret 183
  - Gateway Global 182
  - management 187
  - Priority Scheduler 184
  - Query Configuration 182
  - Query Session 182
  - session configuration rules 187
- Utility management 187

## V

- Viewpoint 32, 33, 191
- Views
  - base tables 74
  - benefits 75
  - DBC.AMPUsage 195
  - definition 74
  - restrictions 75
  - users 140
- Virtual processors 44
- Virtual processors. See Vprocs
- Visual Explain. See Teradata VE
- Volatile tables 74
- Vproc migration
  - cliques 62
  - hardware fault tolerance 62
  - software fault tolerance 55
- Vprocs
  - definition 44
  - function 44
  - maximum per system 45
  - PDE 50
  - vproc migration 55
- Vprocs, types of 44

## W

- WAL 183
- Web-based tools 33, 191
- WebSphere MQ Access Module 32, 181
- Workload Analyzer. See Teradata WA
- Workload definitions 185

- Workload management 183, 185
  - Archive/Recovery utility 187
  - data load utilities 187
  - Priority Scheduler 184
  - Teradata ASM 184
  - Teradata WA 188
  - Write Ahead Logging 183
- Workload Management API 195
- Workload management API 107
- Write Ahead Logging. See WAL
- Write HUT lock 131
- Write level lock 129
- Write-Ahead Logging 183

